

**TECH
VAULT**

MATERIALIZED VIEWS

AMR ELHELW



\d orders

Table "public.orders"				
Column	Type	Collation	Nullable	Default
order_id	integer			
customer_id	integer			
amount	numeric			

SELECT count(*) FROM orders;

count

20000000
(1 row)

SELECT * FROM orders LIMIT 10;

order_id	customer_id	amount
1	2	188.12
2	2	95.80
3	2	252.20
4	3	383.40
5	3	297.55
6	1	172.34
7	2	180.04
8	3	270.27
9	2	291.12
10	3	435.66

(10 rows)

```
SELECT customer_id, count(1), avg(amount) FROM orders GROUP BY customer_id;
```

customer_id	count	avg
1	6671419	250.0531631216687185
2	6665436	249.9877122486811065
3	6663145	250.0006444809470603

(3 rows)

```
CREATE VIEW v_cust AS SELECT customer_id, count(1), avg(amount) FROM orders GROUP BY customer_id;  
CREATE VIEW
```

```
SELECT * FROM v_cust;
```

customer_id	count	avg
1	6671419	250.0531631216687185
2	6665436	249.9877122486811065
3	6663145	250.0006444809470603

(3 rows)

```
EXPLAIN ANALYZE SELECT * FROM v_cust;
```

QUERY PLAN

```
Finalize GroupAggregate (cost=254943.62..254944.41 rows=3 width=44) (actual time=948.322..949.728 rows=3 loops=1)
  Group Key: customer_id
    -> Gather Merge (cost=254943.62..254944.32 rows=6 width=44) (actual time=948.313..949.718 rows=9 loops=1)
      Workers Planned: 2
      Workers Launched: 2
      -> Sort (cost=253943.59..253943.60 rows=3 width=44) (actual time=933.816..933.816 rows=3 loops=3)
        Sort Key: customer_id
        Sort Method: quicksort Memory: 25kB
        Worker 0: Sort Method: quicksort Memory: 25kB
        Worker 1: Sort Method: quicksort Memory: 25kB
        -> Partial HashAggregate (cost=253943.53..253943.57 rows=3 width=44) (actual time=933.784..933.785 rows=3 loops=3)
          Group Key: customer_id
          Batches: 1 Memory Usage: 24kB
          Worker 0: Batches: 1 Memory Usage: 24kB
          Worker 1: Batches: 1 Memory Usage: 24kB
          -> Parallel Seq Scan on orders (cost=0.00..191443.02 rows=8333402 width=10) (actual time=0.079..283.788 rows=6666667 loops=3)

Planning Time: 0.210 ms
Execution Time: 949.820 ms
(18 rows)
```

```
CREATE MATERIALIZED VIEW mv_cust AS SELECT customer_id, count(1), avg(amount) FROM orders GROUP BY customer_id;
SELECT 3
```

\d

List of relations			
Schema	Name	Type	Owner
public	mv_cust	materialized view	postgres
public	orders	table	postgres
public	v_cust	view	postgres

(3 rows)

```
SELECT * FROM pg_matviews WHERE matviewname = 'mv_cust';
```

schemaname	matviewname	matviewowner	tablespace	hasindexes	ispopulated	definition
public	mv_cust	postgres		f	t	SELECT customer_id, + count(1) AS count, + avg(amount) AS avg + FROM orders + GROUP BY customer_id;

(1 row)

```
SELECT * FROM mv_cust;
```

customer_id	count	avg
1	6671419	250.0531631216687185
2	6665436	249.9877122486811065
3	6663145	250.0006444809470603

(3 rows)

```
EXPLAIN ANALYZE SELECT * FROM v_cust;
```

QUERY PLAN

```
-----
Finalize GroupAggregate (cost=254943.62..254944.41 rows=3 width=44) (actual time=948.322..949.728 rows=3 loops=1)
  Group Key: customer_id
    -> Gather Merge (cost=254943.62..254944.32 rows=6 width=44) (actual time=948.313..949.718 rows=9 loops=1)
        Workers Planned: 2
        Workers Launched: 2
        -> Sort (cost=253943.59..253943.60 rows=3 width=44) (actual time=933.816..933.816 rows=3 loops=3)
            Sort Key: customer_id
            Sort Method: quicksort Memory: 25kB
            Worker 0: Sort Method: quicksort Memory: 25kB
            Worker 1: Sort Method: quicksort Memory: 25kB
            -> Partial HashAggregate (cost=253943.53..253943.57 rows=3 width=44) (actual time=933.784..933.785 rows=3 loops=3)
                Group Key: customer_id
                Batches: 1 Memory Usage: 24kB
                Worker 0: Batches: 1 Memory Usage: 24kB
                Worker 1: Batches: 1 Memory Usage: 24kB
                -> Parallel Seq Scan on orders (cost=0.00..191443.02 rows=8333402 width=10) (actual time=0.079..283.788 rows=6666667 loops=3)
Planning Time: 0.210 ms
Execution Time: 949.820 ms
(18 rows)
```

```
EXPLAIN ANALYZE SELECT * FROM mv_cust;
```

QUERY PLAN

```
-----
Seq Scan on mv_cust (cost=0.00..21.30 rows=1130 width=44) (actual time=0.032..0.036 rows=3 loops=1)
Planning Time: 0.158 ms
Execution Time: 0.069 ms
(3 rows)
```

SELECT * FROM v_cust;

customer_id	count	avg
1	6671419	250.0531631216687185
2	6665436	249.9877122486811065
3	6663145	250.0006444809470603

(3 rows)

SELECT * FROM mv_cust;

customer_id	count	avg
1	6671419	250.0531631216687185
2	6665436	249.9877122486811065
3	6663145	250.0006444809470603

(3 rows)

INSERT INTO orders VALUES (20000001, 4, 500);
INSERT 0 1

SELECT customer_id, count(1), avg(amount) FROM orders GROUP BY customer_id;

customer_id	count	avg
1	6671419	250.0531631216687185
2	6665436	249.9877122486811065
3	6663145	250.0006444809470603
4	1	500.0000000000000000

(4 rows)

SELECT * FROM v_cust;

customer_id	count	avg
1	6671419	250.0531631216687185
2	6665436	249.9877122486811065
3	6663145	250.0006444809470603
4	1	500.0000000000000000

(4 rows)

SELECT * FROM mv_cust;

customer_id	count	avg
1	6671419	250.0531631216687185
2	6665436	249.9877122486811065
3	6663145	250.0006444809470603

(3 rows)

REFRESH MATERIALIZED VIEW mv_cust;
REFRESH MATERIALIZED VIEW

SELECT * FROM mv_cust;

customer_id	count	avg
1	6671419	250.0531631216687185
2	6665436	249.9877122486811065
3	6663145	250.0006444809470603
4	1	500.0000000000000000

(4 rows)

Refresh Frequency

- **Auto refresh** - every time the data changes
 - Materialized view always up-to-date
 - Most expensive
- **Fixed frequency** - fixed time interval
 - Simple to implement
 - Tricky to figure out the right frequency
- **Variable Frequency** - after certain number of data writes
 - Better at adjusting frequency
 - More complex to implement

	View	Materialized Views
View result	Computed each time the view is used	Pre-computed and stored in the database. Refreshed automatically or on-demand.
Data freshness	Always up-to-date	May not be up-to-date. Depends on refresh frequency.
Speed to retrieve view results	Slow	Fast
Extra Storage	No	Yes

When to use Materialized Views

- Complex query with small result set
- Data is read only or very few updates
- Don't care if results are slightly out of date