


**TECH  
VAULT**

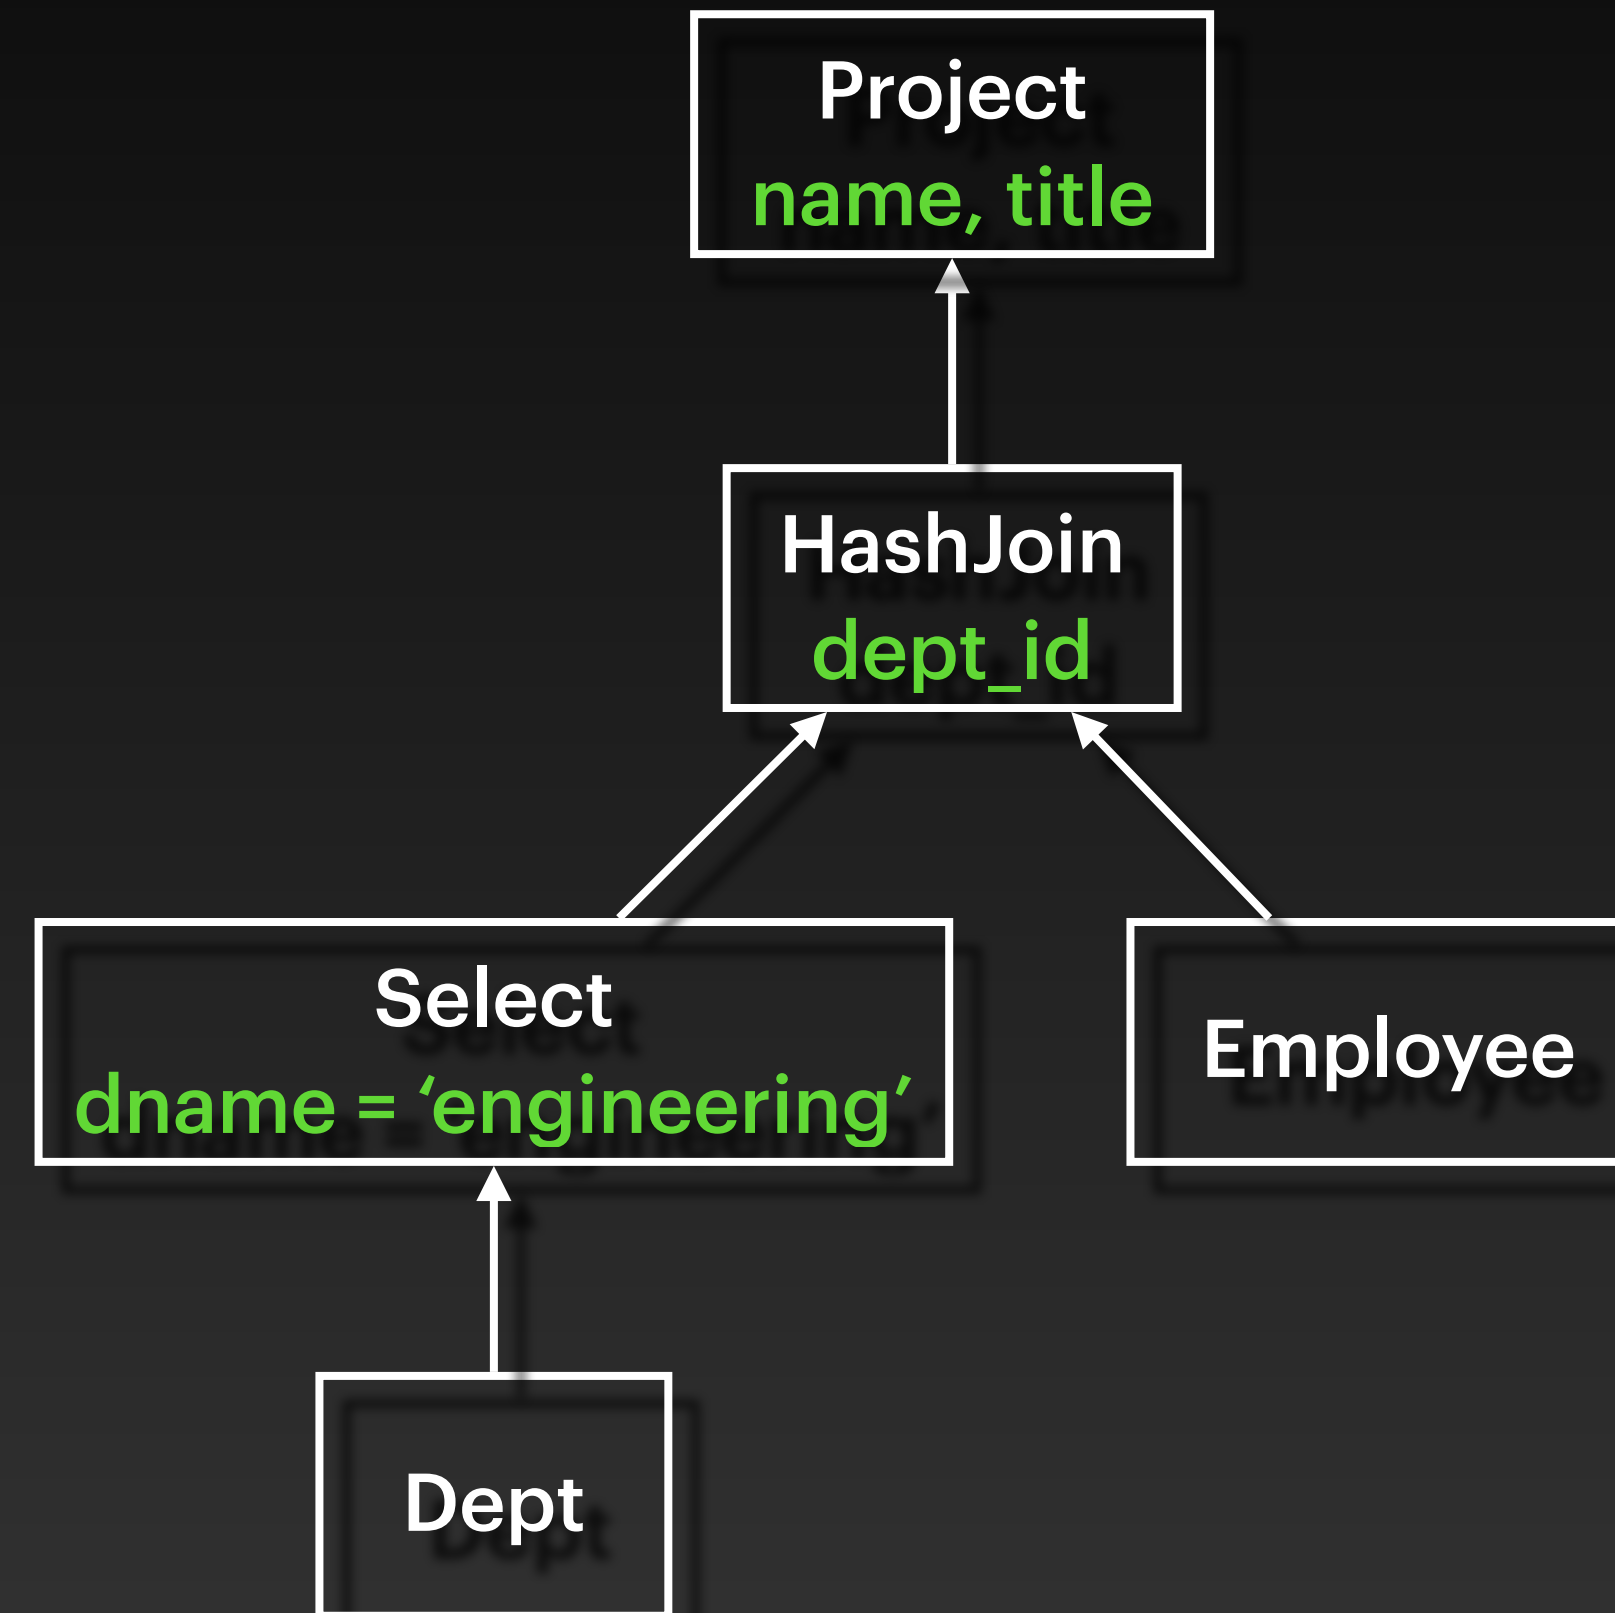
# UNDERSTANDING “EXPLAIN” IN POSTGRES

AMR ELHELW



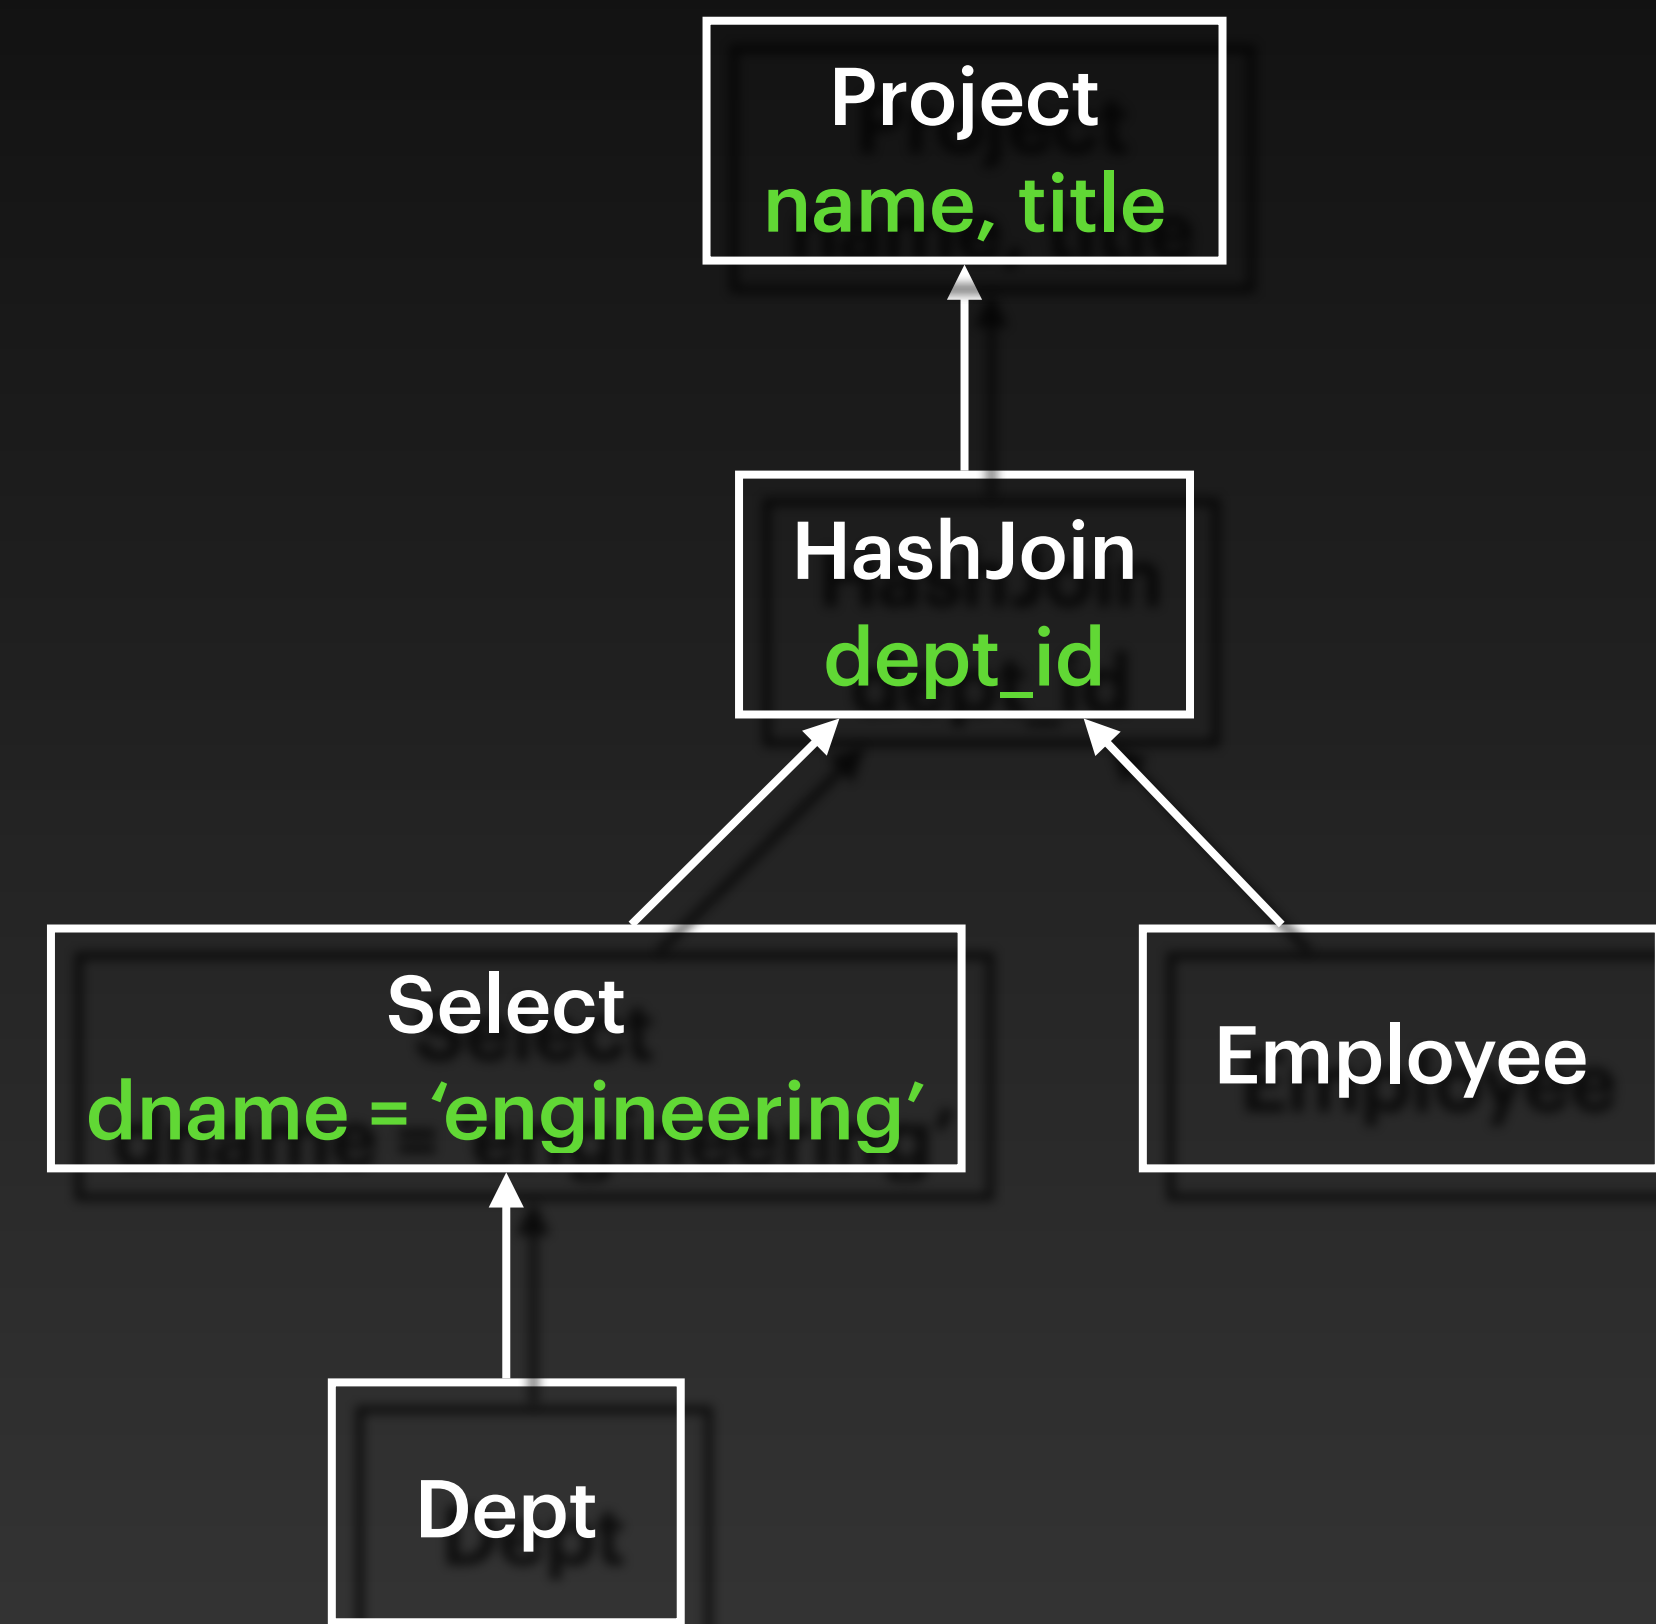


$\Pi_{\text{name, title}}(\text{Employee} \bowtie_{\text{dept\_id}} \sigma_{\text{dname} = \text{'engineering'}}(\text{Dept}))$



# Node Types

- Leaf (no inputs)
- Unary (1 input)
- Binary (2 inputs)
- N-ary ( $N$  inputs,  $N > 2$ )



```
EXPLAIN SELECT * FROM employee;  
          QUERY PLAN
```

```
-----  
Seq Scan on employee (cost=0.00..82.00 rows=5000 width=21)  
(1 row)
```

```
EXPLAIN ANALYZE SELECT * FROM employee;
```

```
          QUERY PLAN
```

```
-----  
Seq Scan on employee (cost=0.00..82.00 rows=5000 width=21) (actual time=0.032..0.488 rows=5000 loops=1)  
Planning Time: 0.195 ms  
Execution Time: 0.749 ms  
(3 rows)
```

```
EXPLAIN SELECT id FROM employee;  
          QUERY PLAN
```

```
-----  
Seq Scan on employee (cost=0.00..82.00 rows=5000 width=4)  
(1 row)
```

```
EXPLAIN SELECT name FROM employee;  
          QUERY PLAN
```

```
-----  
Seq Scan on employee (cost=0.00..82.00 rows=5000 width=13)  
(1 row)
```

```
EXPLAIN SELECT * FROM employee;  
      QUERY PLAN
```

```
-----  
Seq Scan on employee  (cost=0.00..82.00 rows=5000 width=21)  
(1 row)
```

```
EXPLAIN SELECT * FROM employee ORDER BY name;  
      QUERY PLAN
```


```
-----  
Sort  (cost=389.19..401.69 rows=5000 width=21)  
  Sort Key: name  
    -> Seq Scan on employee  (cost=0.00..82.00 rows=5000 width=21)  
(3 rows)
```

```
EXPLAIN ANALYZE SELECT * FROM employee ORDER BY name;  
      QUERY PLAN
```

```
-----  
Sort  (cost=389.19..401.69 rows=5000 width=21) (actual time=2.947..3.522 rows=5000 loops=1)  
  Sort Key: name  
  Sort Method: quicksort  Memory: 427kB  
    -> Seq Scan on employee  (cost=0.00..82.00 rows=5000 width=21) (actual time=0.016..0.459 rows=5000 loops=1)  
Planning Time: 0.105 ms  
Execution Time: 3.888 ms  
(6 rows)
```

```
EXPLAIN SELECT * FROM employee ORDER BY id;  
      QUERY PLAN
```

```
-----  
Index Scan using emp_id on employee  (cost=0.28..174.28 rows=5000 width=21)  
(1 row)
```



```
EXPLAIN SELECT e.id, e.name FROM employee e JOIN dept d ON e.dept_id = d.id WHERE d.name='Engineering' ORDER BY e.name;
      QUERY PLAN
```

```
-----
Sort  (cost=114.47..115.09 rows=250 width=17)
  Sort Key: e.name
  -> Hash Join  (cost=1.26..104.51 rows=250 width=17)
    Hash Cond: (e.dept_id = d.id)
    -> Seq Scan on employee e  (cost=0.00..82.00 rows=5000 width=21)
    -> Hash  (cost=1.25..1.25 rows=1 width=4)
      -> Seq Scan on dept d  (cost=0.00..1.25 rows=1 width=4)
        Filter: ((name)::text = 'Engineering'::text)
```

(8 rows)

```
EXPLAIN ANALYZE SELECT e.id, e.name FROM employee e JOIN dept d ON e.dept_id = d.id WHERE d.name='Engineering' ORDER BY e.name;
      QUERY PLAN
```

```
-----
Sort  (cost=114.47..115.09 rows=250 width=17) (actual time=1.389..1.411 rows=276 loops=1)
  Sort Key: e.name
  Sort Method: quicksort  Memory: 37kB
  -> Hash Join  (cost=1.26..104.51 rows=250 width=17) (actual time=0.085..1.244 rows=276 loops=1)
    Hash Cond: (e.dept_id = d.id)
    -> Seq Scan on employee e  (cost=0.00..82.00 rows=5000 width=21) (actual time=0.028..0.512 rows=5000 loops=1)
    -> Hash  (cost=1.25..1.25 rows=1 width=4) (actual time=0.036..0.037 rows=1 loops=1)
      Buckets: 1024  Batches: 1  Memory Usage: 9kB
      -> Seq Scan on dept d  (cost=0.00..1.25 rows=1 width=4) (actual time=0.021..0.023 rows=1 loops=1)
        Filter: ((name)::text = 'Engineering'::text)
        Rows Removed by Filter: 19
```

Planning Time: 0.315 ms

Execution Time: 1.486 ms

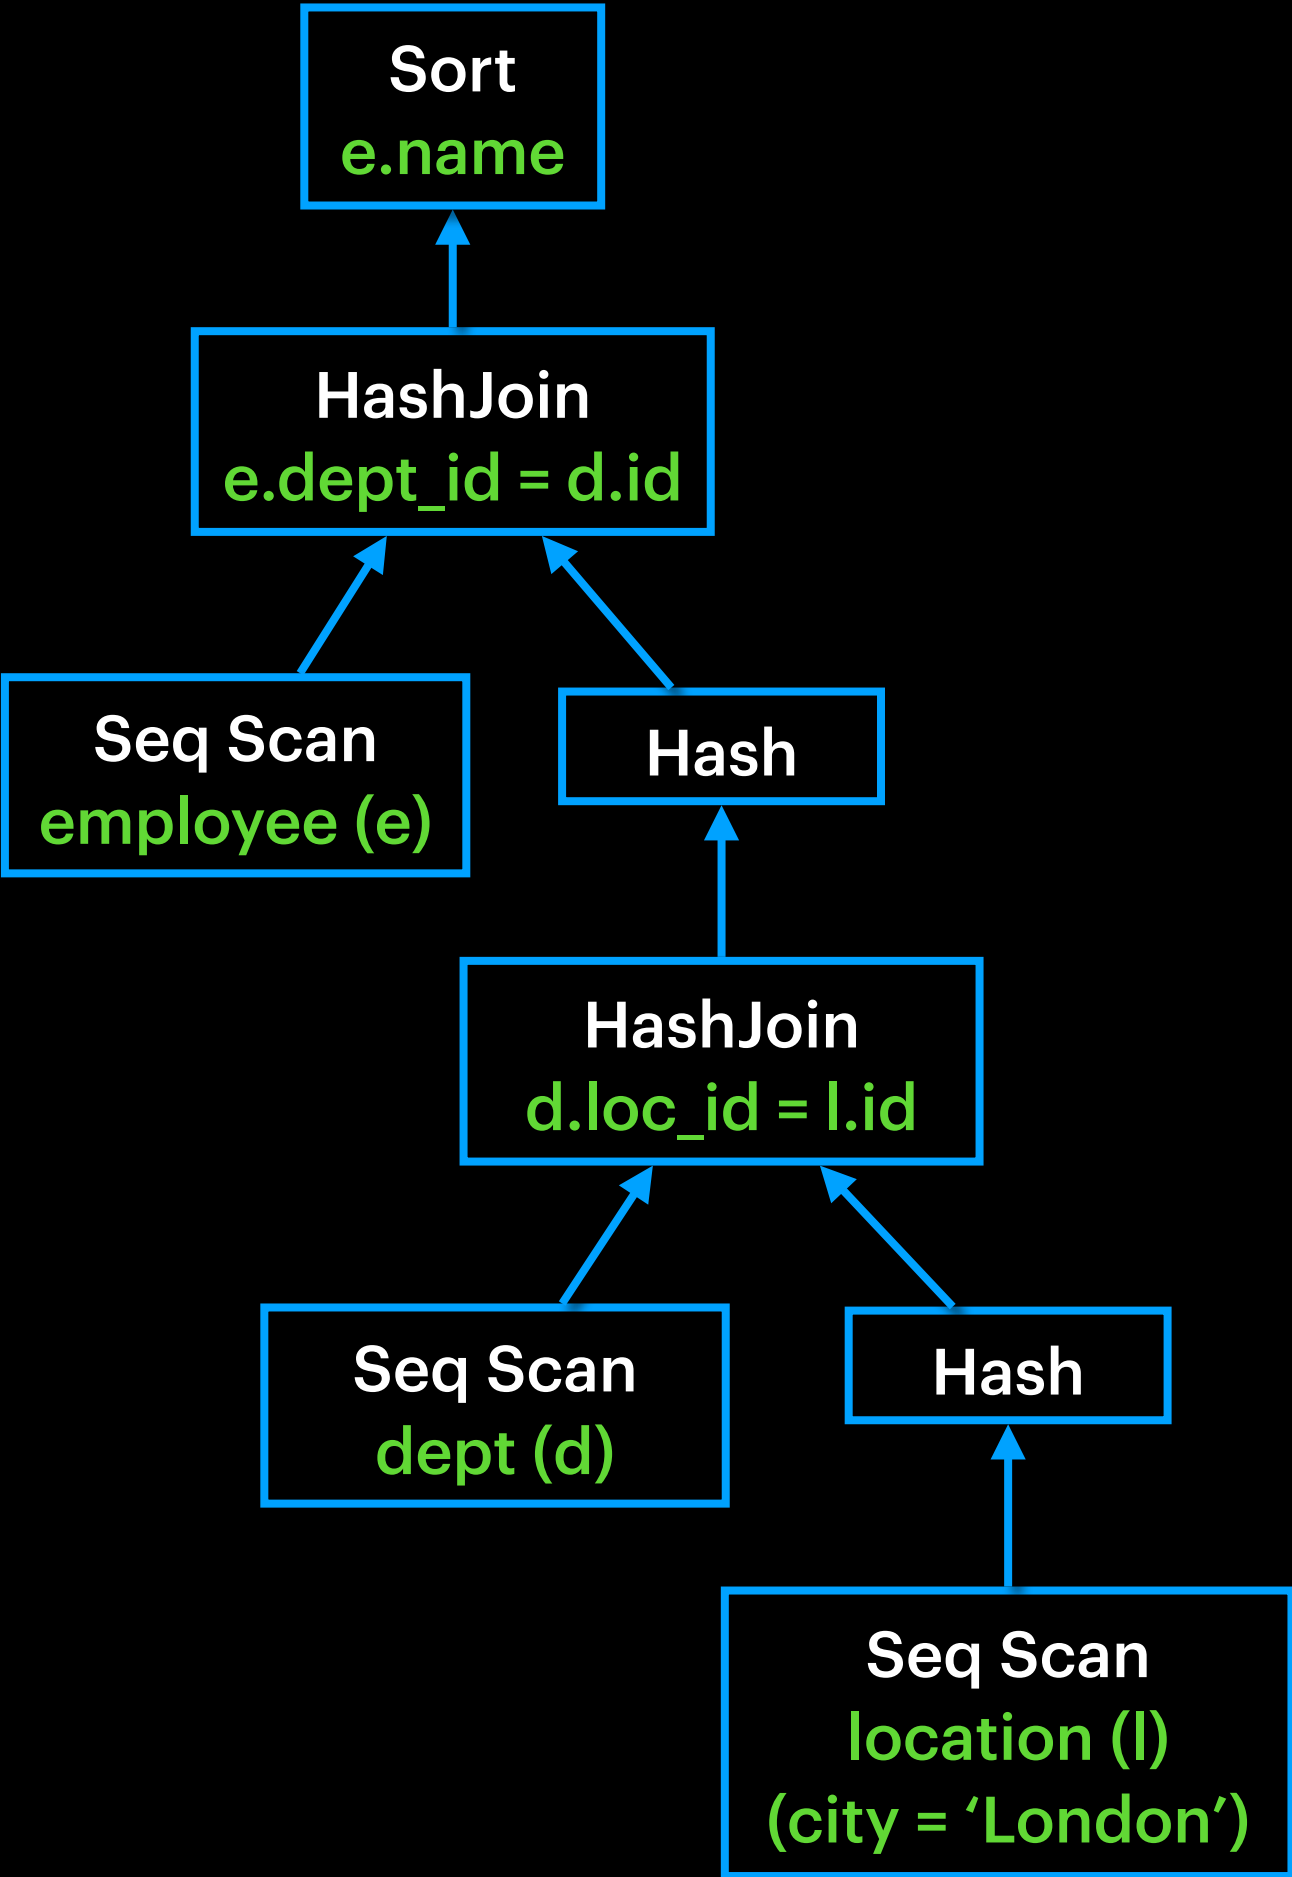
(13 rows)

```
EXPLAIN SELECT e.id, e.name FROM employee e JOIN dept d ON e.dept_id = d.id JOIN location l ON d.loc_id = l.id WHERE l.city='London' ORDER BY e.name;
```

QUERY PLAN

```
Sort (cost=232.95..237.95 rows=2000 width=17)
  Sort Key: e.name
    -> Hash Join (cost=2.54..123.29 rows=2000 width=17)
      Hash Cond: (e.dept_id = d.id)
        -> Seq Scan on employee e (cost=0.00..82.00 rows=5000 width=21)
        -> Hash (cost=2.44..2.44 rows=8 width=4)
          -> Hash Join (cost=1.09..2.44 rows=8 width=4)
            Hash Cond: (d.loc_id = l.id)
              -> Seq Scan on dept d (cost=0.00..1.20 rows=20 width=8)
              -> Hash (cost=1.06..1.06 rows=2 width=4)
                -> Seq Scan on location l (cost=0.00..1.06 rows=2 width=4)
                  Filter: ((city)::text = 'London'::text)
```

(12 rows)





<https://www.pgexplain.dev/>

Plan

QUERY PLAN

Sort (cost=232.95..237.95 rows=2000 width=17)

Sort Key: e.name

-> Hash Join (cost=2.5

Hash Cond: (e.dept

-> Seq Scan on emp

-> Hash (cost=2.44

-> Hash Join (c

Hash Cond: (

-> Seq Scan

-> Hash (co

-> Seq S

Filter:

Plan Raw Query Stats

Execution time: N/A Planning time: N/A Triggers: N/A

Settings

time rows estimation cost

- #1 Sort
- #2 Hash Join
- #3 Seq Scan
- #4 Hash
- #5 Hash Join
- #6 Seq Scan
- #7 Hash
- #8 Seq Scan

#3 Seq Scan  
on employee as e

#1 Sort  
by e.name

#2 Hash Join  
on e.dept\_id = d.id

#4 Hash

#5 Hash Join  
on d.loc\_id = l.id

#6 Seq Scan  
on dept as d

#7 Hash

#8 Seq Scan  
on location as l