

**TECH  
VAULT**

# **CARDINALITY ESTIMATION**

## **PART (3)**



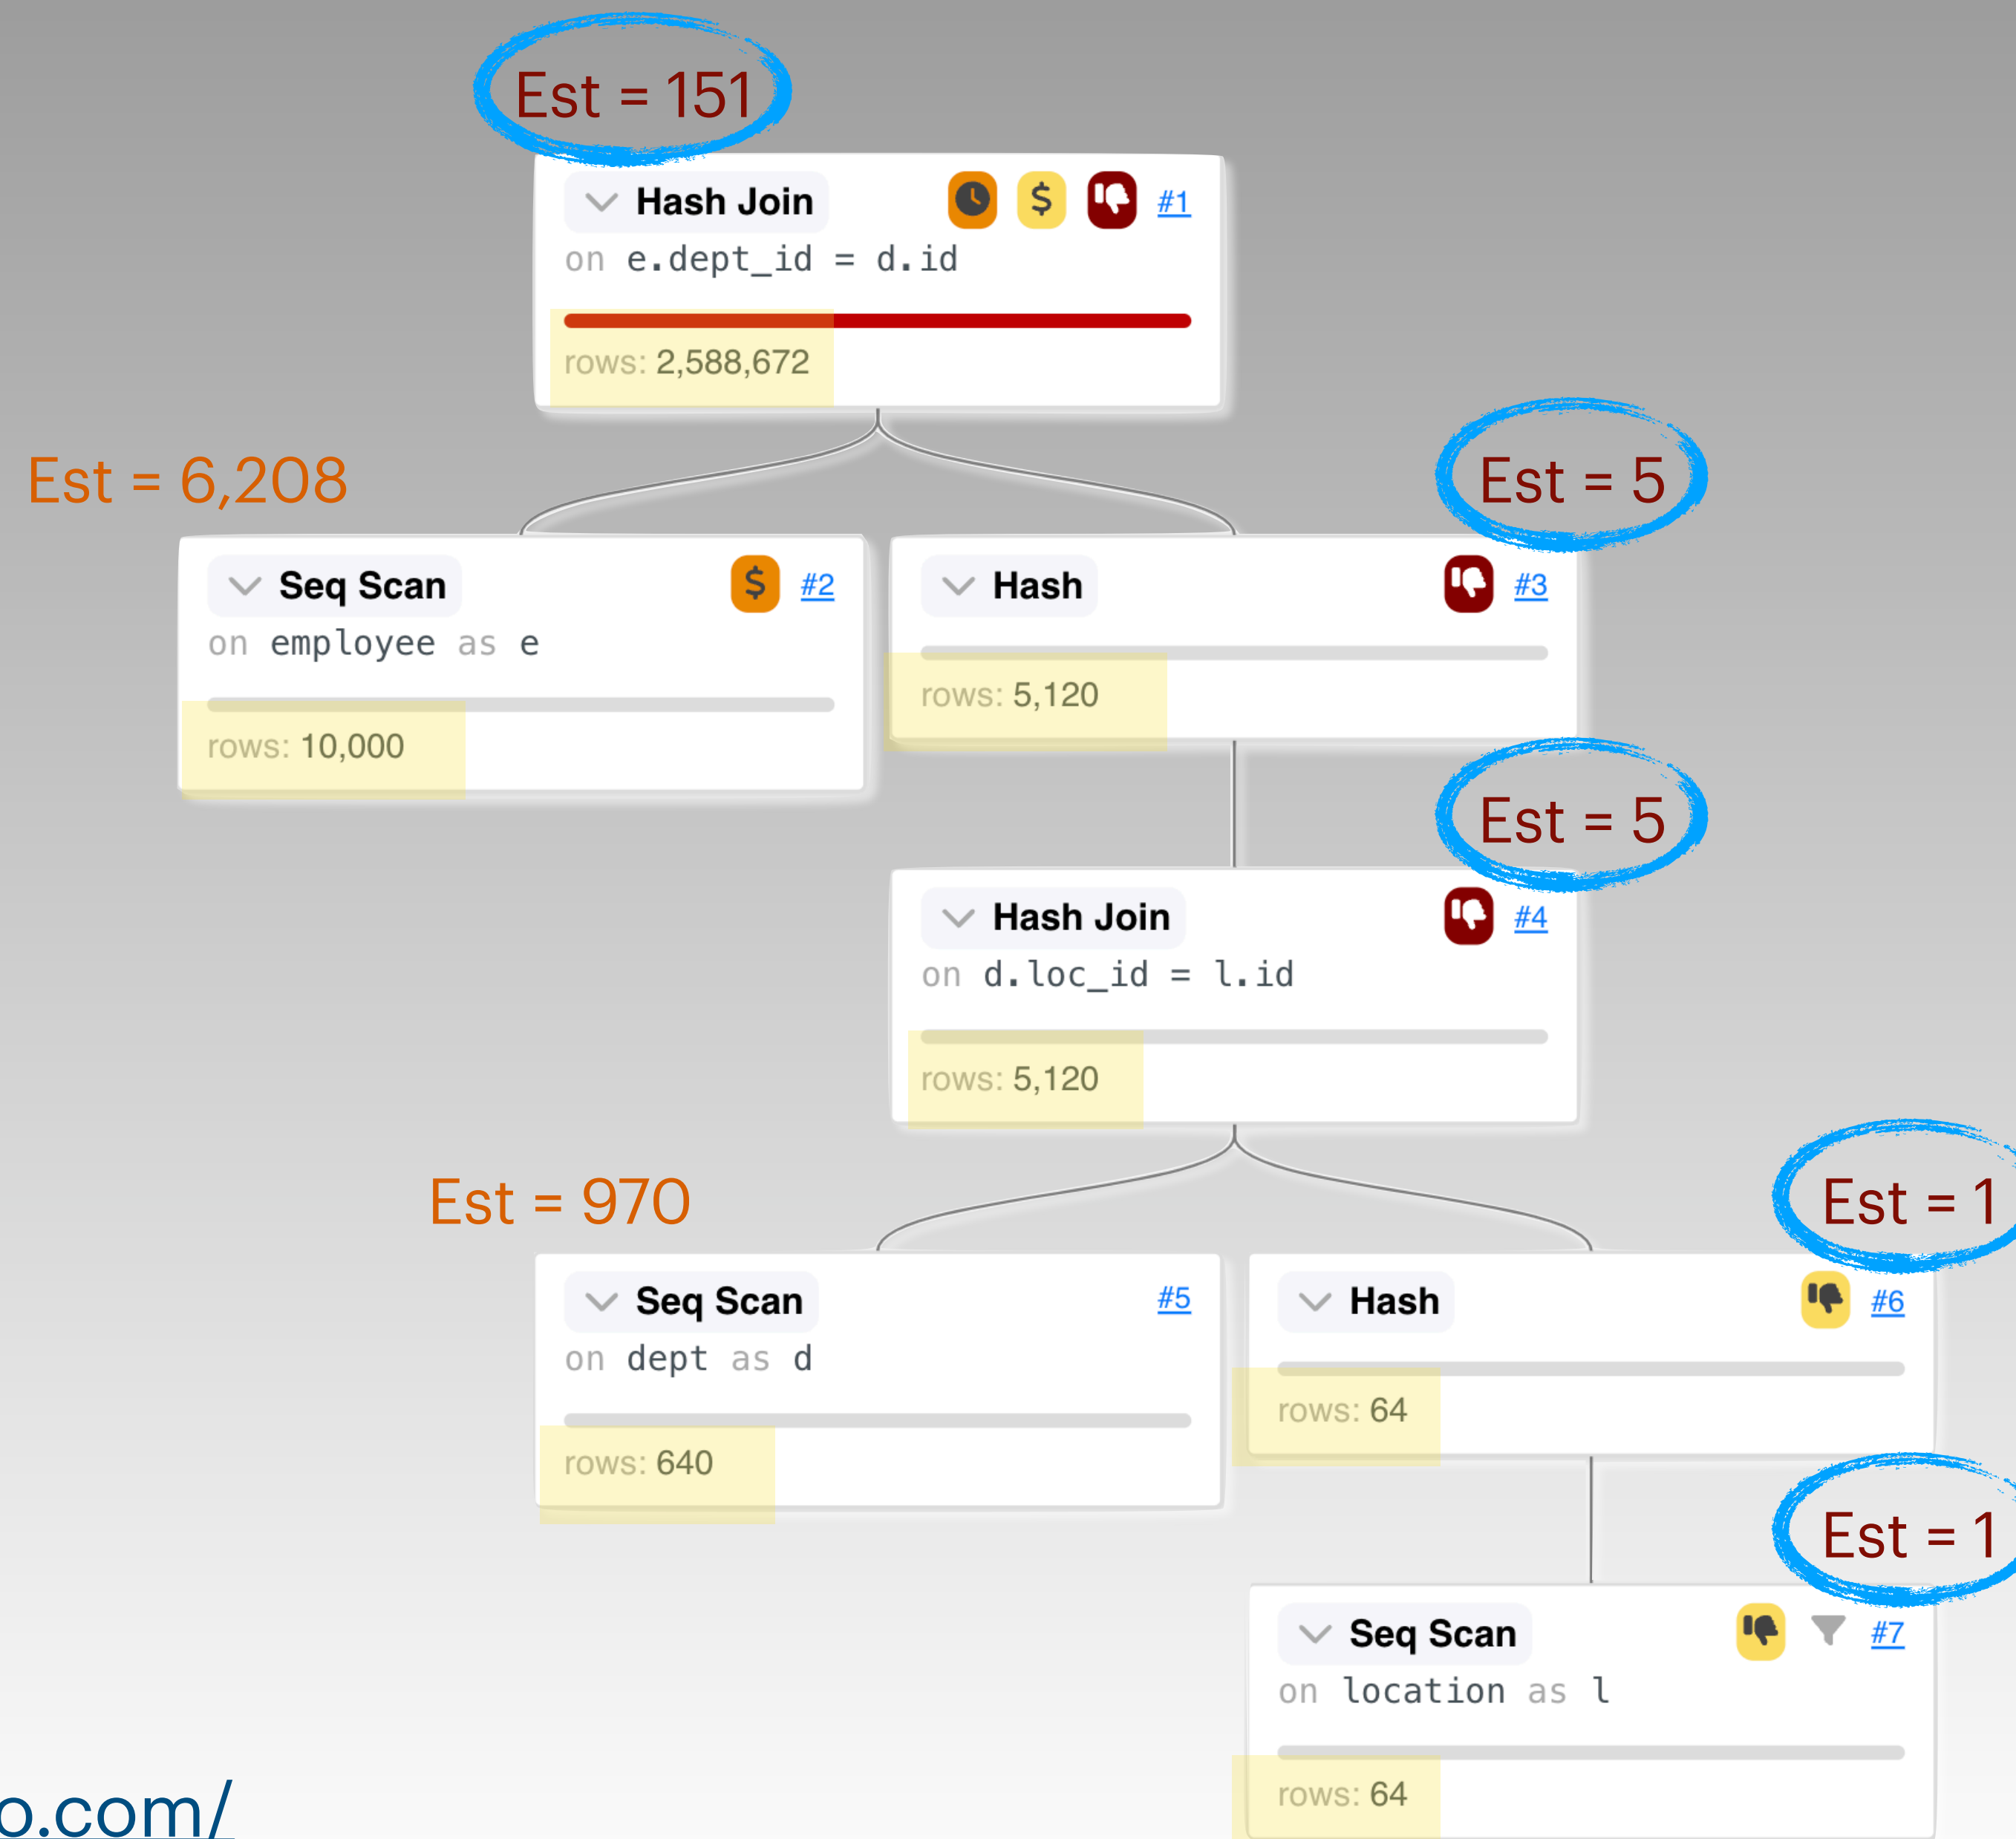
**AMR ELHELW**



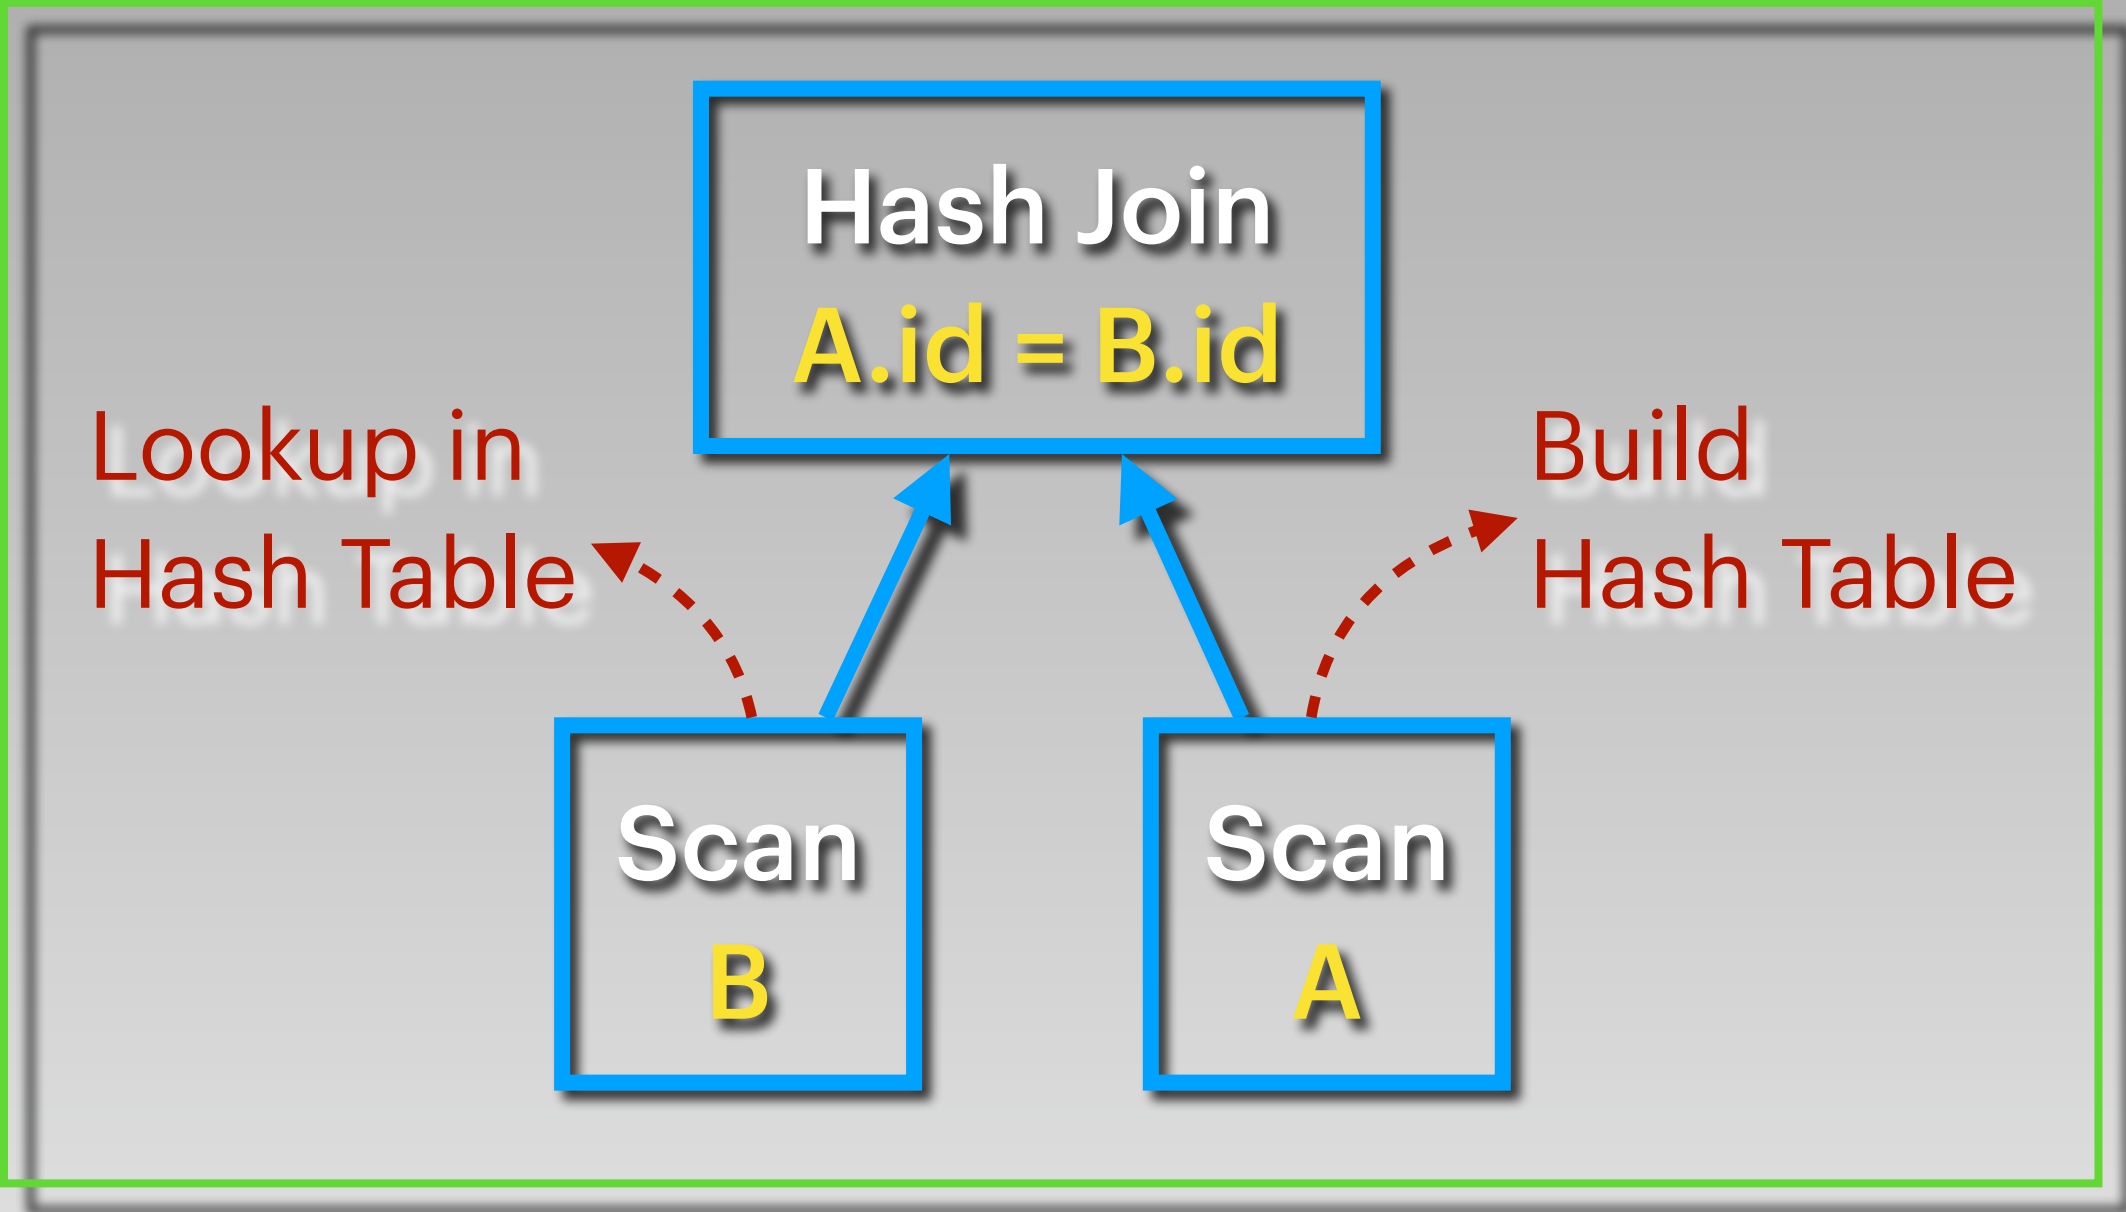

```
EXPLAIN ANALYZE SELECT e.name, d.name FROM employee e JOIN dept d ON e.dept_id = d.id JOIN location l ON d.loc_id = l.id
WHERE country = 'UK' AND city = 'London';

QUERY PLAN
-----
Hash Join (cost=41.71..192.62 rows=151 width=96) (actual time=2.082..130.412 rows=2588672 loops=1)
  Hash Cond: (e.dept_id = d.id)
    -> Seq Scan on employee e (cost=0.00..126.08 rows=6208 width=52) (actual time=0.035..0.603 rows=10000 loops=1)
    -> Hash (cost=41.65..41.65 rows=5 width=52) (actual time=2.059..2.061 rows=5120 loops=1)
      Buckets: 8192 (originally 1024) Batches: 1 (originally 1) Memory Usage: 288kB
      -> Hash Join (cost=18.26..41.65 rows=5 width=52) (actual time=0.103..1.087 rows=5120 loops=1)
        Hash Cond: (d.loc_id = l.id)
          -> Seq Scan on dept d (cost=0.00..19.70 rows=970 width=56) (actual time=0.010..0.089 rows=640 loops=1)
          -> Hash (cost=18.25..18.25 rows=1 width=4) (actual time=0.079..0.079 rows=64 loops=1)
            Buckets: 1024 Batches: 1 Memory Usage: 11kB
            -> Seq Scan on location l (cost=0.00..18.25 rows=1 width=4) (actual time=0.023..0.059 rows=64 loops=1)
              Filter: (((country)::text = 'UK'::text) AND ((city)::text = 'London'::text))
              Rows Removed by Filter: 96

Planning Time: 0.451 ms
Execution Time: 192.303 ms
(15 rows)
```



```
SELECT ... FROM A JOIN B ON A.id = B.id;
```



	Estimated no. of rows	Actual no. of rows
Table A	100	200,000
Table B	5,000	1,000

```
SELECT relname, reltuples, relpages FROM pg_class WHERE relname IN ('location','dept','employee');
```

relname	reltuples	relpages
employee	-1	0
dept	-1	0
location	-1	0

(3 rows)

```
SELECT tablename, attname, null_frac, n_distinct FROM pg_stats WHERE tablename IN ('location','dept','employee');
```

tablename	attname	null_frac	n_distinct
-----------	---------	-----------	------------

(0 rows)

```
ANALYZE;
```

```
SELECT relname, reltuples, relpages FROM pg_class WHERE relname IN ('location','dept','employee');
```

relname	reltuples	relpages
dept	640	4
employee	10000	64
location	160	2

(3 rows)

```
EXPLAIN ANALYZE SELECT e.name, d.name FROM employee e JOIN dept d ON e.dept_id = d.id JOIN location l ON d.loc_id = l.id
WHERE country = 'UK' AND city = 'London';
```

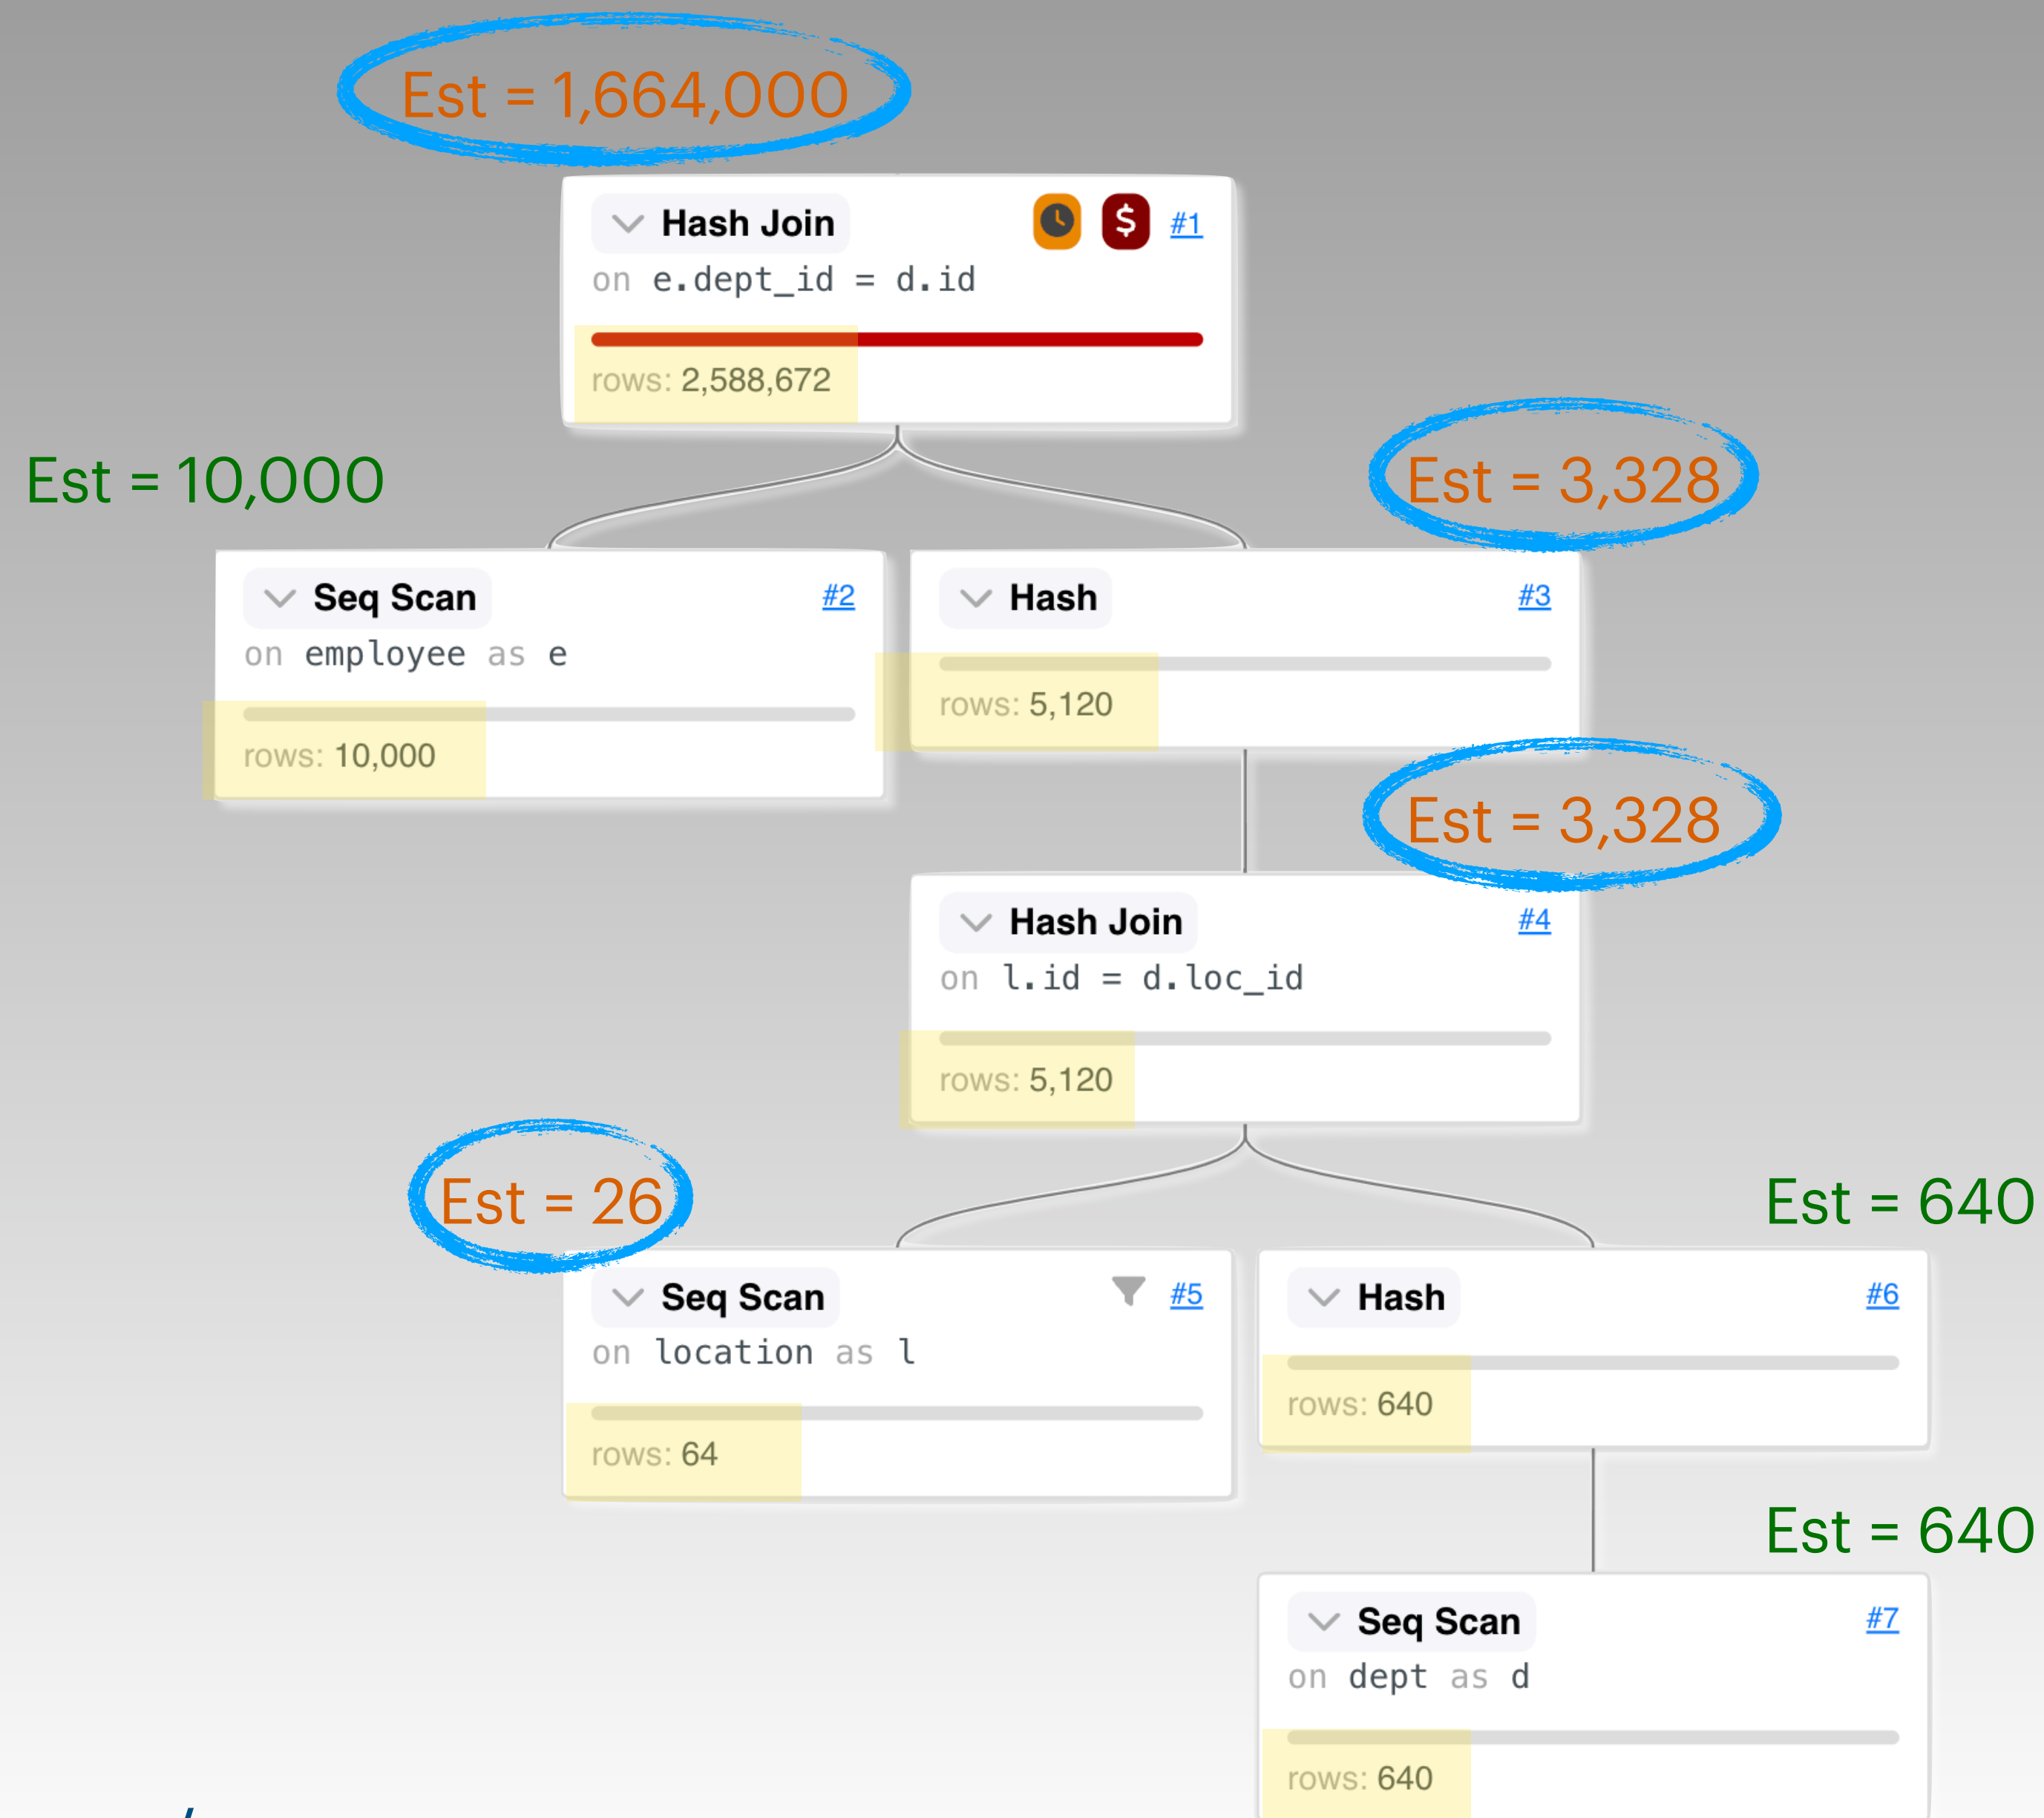
#### QUERY PLAN

```
Hash Join (cost=103.98..19007.99 rows=1664000 width=20) (actual time=2.646..125.921 rows=2588672 loops=1)
  Hash Cond: (e.dept_id = d.id)
    -> Seq Scan on employee e (cost=0.00..164.00 rows=10000 width=17) (actual time=0.036..0.529 rows=10000 loops=1)
    -> Hash (cost=62.38..62.38 rows=3328 width=11) (actual time=2.595..2.597 rows=5120 loops=1)
      Buckets: 8192 (originally 4096) Batches: 1 (originally 1) Memory Usage: 288kB
      -> Hash Join (cost=18.40..62.38 rows=3328 width=11) (actual time=0.301..1.436 rows=5120 loops=1)
        Hash Cond: (l.id = d.loc_id)
          -> Seq Scan on location l (cost=0.00..4.40 rows=26 width=4) (actual time=0.021..0.055 rows=64 loops=1)
            Filter: (((country)::text = 'UK'::text) AND ((city)::text = 'London'::text))
            Rows Removed by Filter: 96
          -> Hash (cost=10.40..10.40 rows=640 width=15) (actual time=0.268..0.269 rows=640 loops=1)
            Buckets: 1024 Batches: 1 Memory Usage: 39kB
            -> Seq Scan on dept d (cost=0.00..10.40 rows=640 width=15) (actual time=0.013..0.133 rows=640 loops=1)
```

Planning Time: 0.498 ms

Execution Time: 186.140 ms

(15 rows)





Customer

Name	City	Country
...	London	UK
...	Madrid	Spain
...	New York	US
...	Seattle	US
...	Paris	France

Rows in 'customer' table = 1,000

No. of distinct values in 'city' column = 5

No. of distinct values in 'country' column = 4

```
SELECT ... FROM customer WHERE city = 'Seattle' AND country = 'US';
```

Assuming  
Independence

Selectivity = 1/5

Selectivity = 1/4

Combined Selectivity = 1/5 \* 1/4 = 1/20

Estimated Rows = 1000 \* 1/20 = 50

No. of distinct (city, country) combinations = 5

Selectivity = 1/5

Estimated Rows = 1000 \* 1/5 = 200



```
EXPLAIN ANALYZE SELECT e.name, d.name FROM employee e JOIN dept d ON e.dept_id = d.id JOIN location l ON d.loc_id = l.id
WHERE country = 'UK' AND city = 'London';
```

#### QUERY PLAN

```
Hash Join (cost=103.98..19007.99 rows=1664000 width=20) (actual time=2.646..125.921 rows=2588672 loops=1)
  Hash Cond: (e.dept_id = d.id)
    -> Seq Scan on employee e (cost=0.00..164.00 rows=10000 width=17) (actual time=0.036..0.529 rows=10000 loops=1)
    -> Hash (cost=62.38..62.38 rows=3328 width=11) (actual time=2.595..2.597 rows=5120 loops=1)
        Buckets: 8192 (originally 4096) Batches: 1 (originally 1) Memory Usage: 288kB
        -> Hash Join (cost=18.40..62.38 rows=3328 width=11) (actual time=0.301..1.436 rows=5120 loops=1)
            Hash Cond: (l.id = d.loc_id)
            -> Seq Scan on location l (cost=0.00..4.40 rows=26 width=4) (actual time=0.021..0.055 rows=64 loops=1)
                Filter: (((country)::text = 'UK'::text) AND ((city)::text = 'London'::text))
                Rows Removed by Filter: 96
            -> Hash (cost=10.40..10.40 rows=640 width=15) (actual time=0.268..0.269 rows=640 loops=1)
                Buckets: 1024 Batches: 1 Memory Usage: 39kB
                -> Seq Scan on dept d (cost=0.00..10.40 rows=640 width=15) (actual time=0.013..0.133 rows=640 loops=1)
```

Planning Time: 0.498 ms

Execution Time: 186.140 ms

(15 rows)

```
CREATE STATISTICS stats_loc ON country, city FROM location;
CREATE STATISTICS

ANALYZE location;
ANALYZE

SELECT m.* FROM pg_statistic_ext JOIN pg_statistic_ext_data ON (oid = stxoid), pg_mcv_list_items(stxdmcv) m WHERE stxname =
'stats_loc';
```

index	values	nulls	frequency	base_frequency
0	{London,UK}	{f,f}	0.4	0.160000000000000003
1	{Madrid,Spain}	{f,f}	0.2	0.040000000000000001
2	{"New York",US}	{f,f}	0.2	0.040000000000000001
3	{Paris,France}	{f,f}	0.2	0.040000000000000001

(4 rows)

```
EXPLAIN ANALYZE SELECT e.name, d.name FROM employee e JOIN dept d ON e.dept_id = d.id JOIN location l ON d.loc_id = l.id
WHERE country = 'UK' AND city = 'London';
```

#### QUERY PLAN

```
Hash Join (cost=222.64..46496.64 rows=4096000 width=20) (actual time=2.252..137.630 rows=2588672 loops=1)
  Hash Cond: (e.dept_id = d.id)
    -> Seq Scan on employee e (cost=0.00..164.00 rows=10000 width=17) (actual time=0.015..0.538 rows=10000 loops=1)
    -> Hash (cost=120.24..120.24 rows=8192 width=11) (actual time=2.225..2.226 rows=5120 loops=1)
      Buckets: 8192 Batches: 1 Memory Usage: 288kB
      -> Hash Join (cost=18.40..120.24 rows=8192 width=11) (actual time=0.283..1.251 rows=5120 loops=1)
        Hash Cond: (l.id = d.loc_id)
          -> Seq Scan on location l (cost=0.00..4.40 rows=64 width=4) (actual time=0.015..0.048 rows=64 loops=1)
            Filter: (((country)::text = 'UK'::text) AND ((city)::text = 'London'::text))
            Rows Removed by Filter: 96
          -> Hash (cost=10.40..10.40 rows=640 width=15) (actual time=0.257..0.258 rows=640 loops=1)
            Buckets: 1024 Batches: 1 Memory Usage: 39kB
            -> Seq Scan on dept d (cost=0.00..10.40 rows=640 width=15) (actual time=0.009..0.126 rows=640 loops=1)
```

Planning Time: 0.547 ms

Execution Time: 209.576 ms

(15 rows)



```
SELECT ... FROM T WHERE col1 + 6 * col2 < 120;
```

```
SELECT ... FROM T WHERE value_str LIKE '%st1%';
```


```
SELECT ... FROM T WHERE func(col1) = 0;
```

One solution is “Sampling”

```
... WHERE col1 + 6 * col2 < 120;
```

## Apply filter

# T



# 100,000 Rows

## Random Sample (e.g. 1%)

# T\_sample

1,000 Rows

176 rows satisfy filter

Selectivity = 0.176