

CONCURRENCY CONTROL

Two-Phase Locking
Part (2)



AMR ELHELW

T1

T2

S-Lock (A)

Read (A)

X-Lock (B)

S-Lock (B)

Read (B)

X-Lock (A)

Time

T1

T2

S-Lock (A)

Read (A)



X-Lock (B)

Write (B)

...

S-Lock (B)

Read (B)



X-Lock (A)

Write (A)

...

A: S-Locked by T1

B: S-Locked by T2

Deadlock!

Amr Elhelw's

TECH

VAULT

DeadLock Detection

- Uses “waits-for” graph
 - Each node is a transaction
 - Edge from T_i to T_j if T_i is waiting for a lock held by T_j .
 - Cycle = deadlock!!

T1

T2

S-Lock (A)

Read (A)

X-Lock (B)

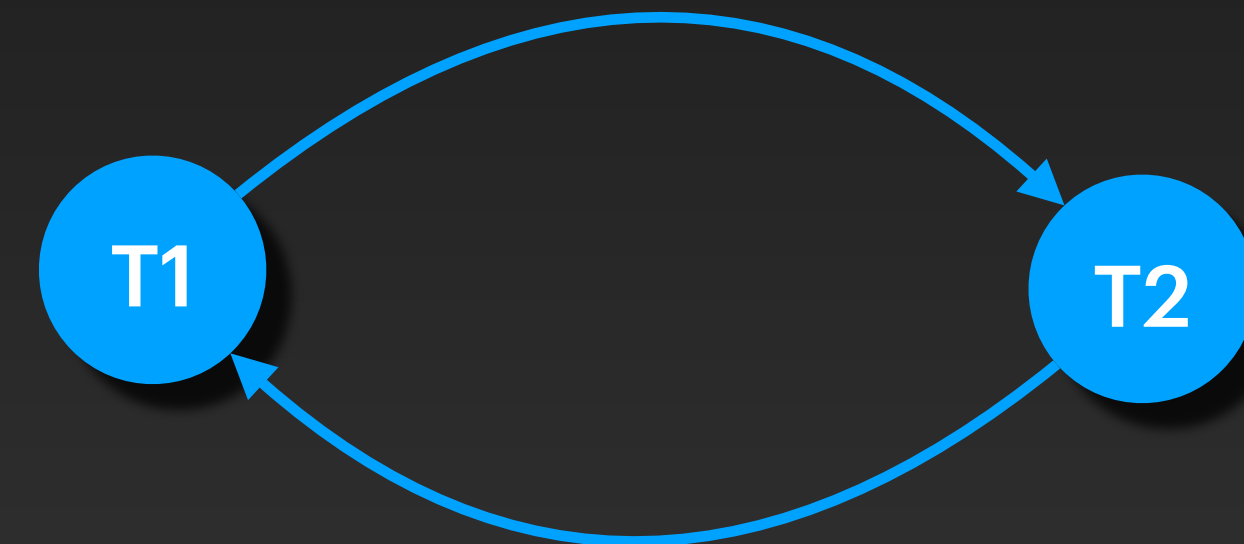
S-Lock (B)

Read (B)

X-Lock (A)

Time

“Waits-for” Graph



Deadlock!!

Victim Selection

Possible criteria

- Newest (by timestamp)
- Least number of writes done
- Most number of locked items
- Least number of rollbacks (avoid starvation)

• ...

DeadLock Prevention

- **Method 1:** Use **Conservative 2PL**
 - Transaction request all locks at start.
 - If any lock is unavailable, don't acquire any locks, wait, then retry.

DeadLock Prevention

- **Method 2: Timestamp-based**
 - Each transaction has a timestamp *ts*
 - *T1* is “older” than *T2* if $ts(T1) < ts(T2)$
 - If a transaction requests a lock held by another transaction —> Either *block* or *abort* one of the two, based on the timestamps

DeadLock Prevention

- **Method 2: Timestamp-based**

- T_{hld} holds lock, T_{req} requests lock

- **Wait-Die**

- If T_{req} older than $T_{hld} \rightarrow T_{req}$ waits, else abort T_{req}

- **Wound-Wait**

- If T_{req} older than $T_{hld} \rightarrow$ abort T_{hld} , else T_{req} waits

Lock Hierarchy

Database (not very common)

↪ Table (very common)

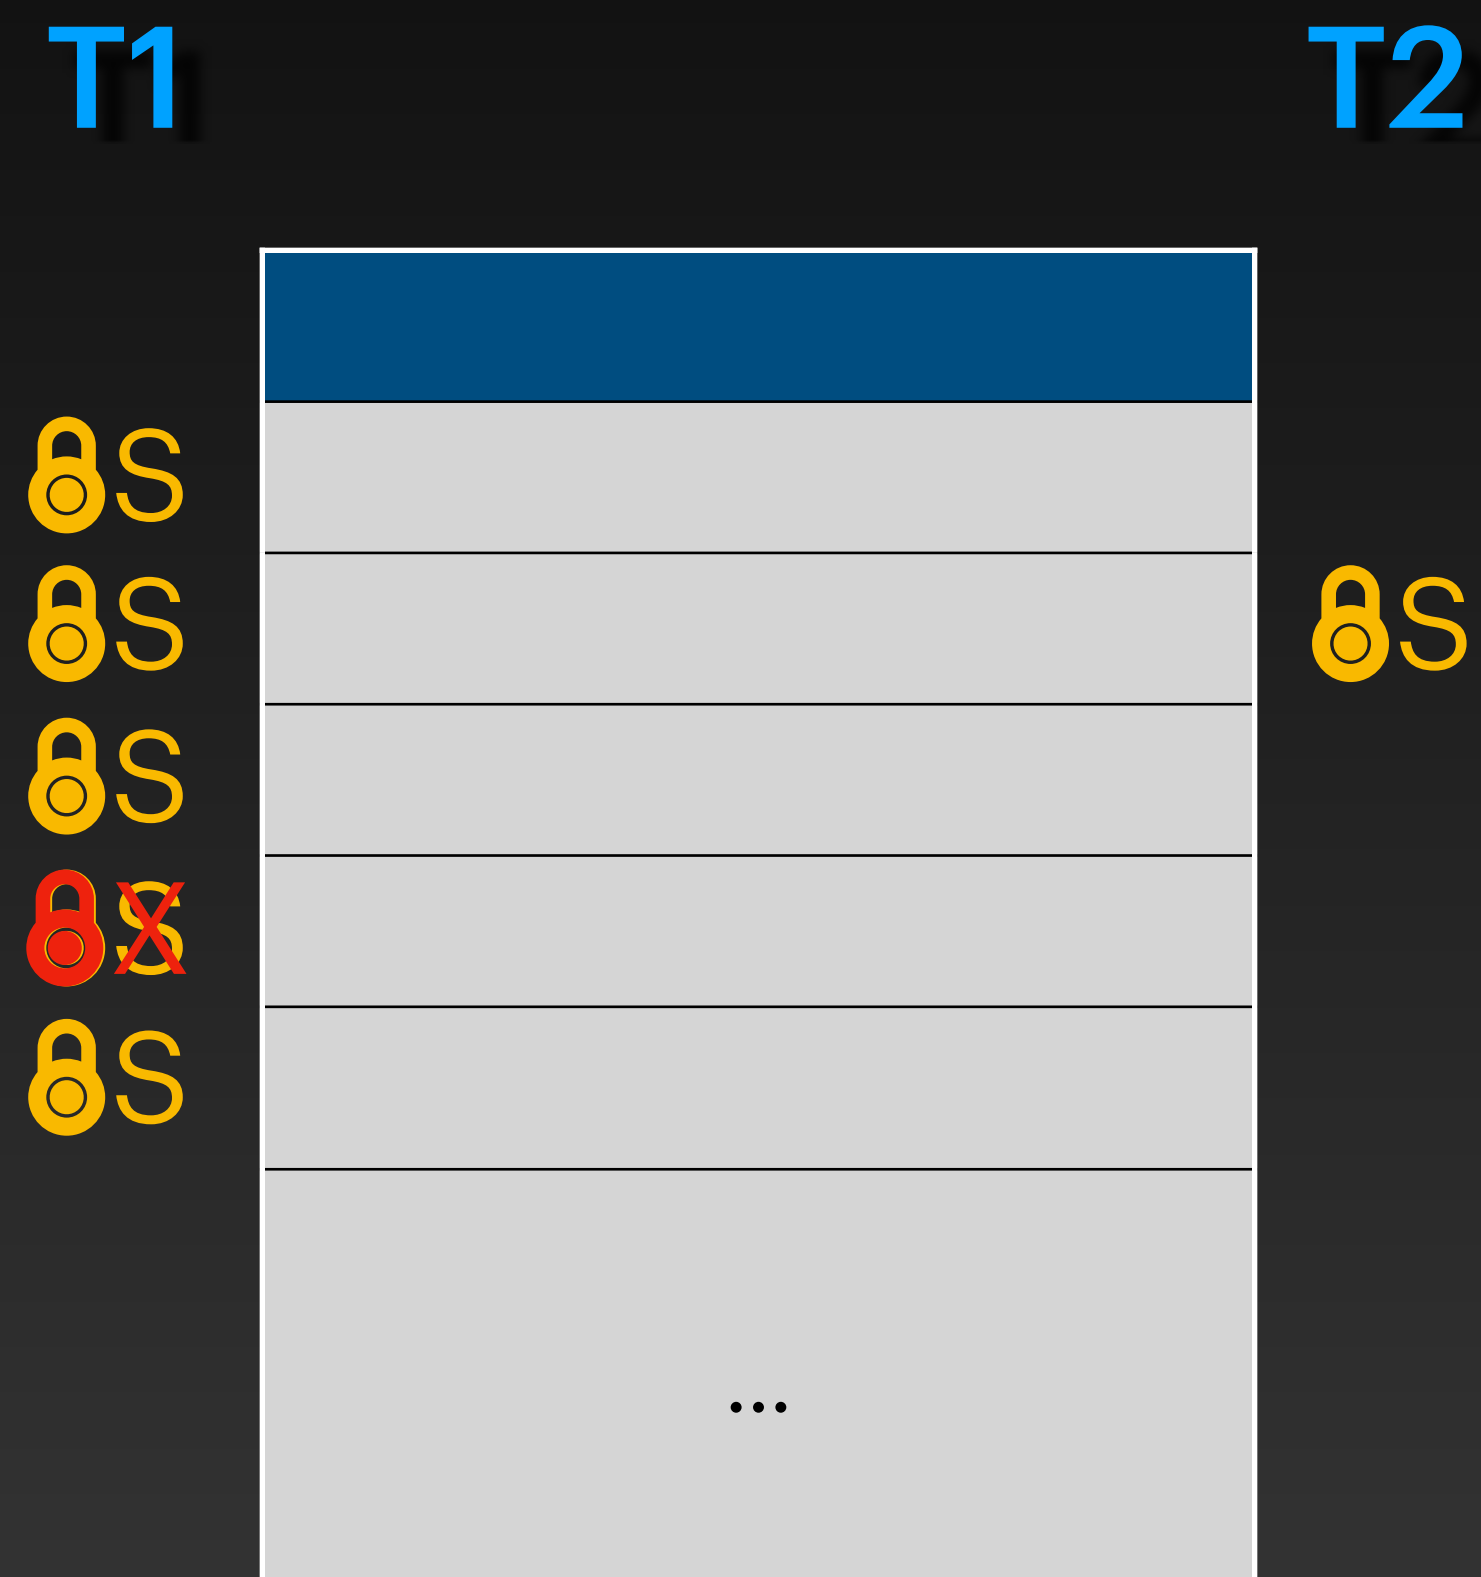
↪ Page (common)

↪ Row (very common)

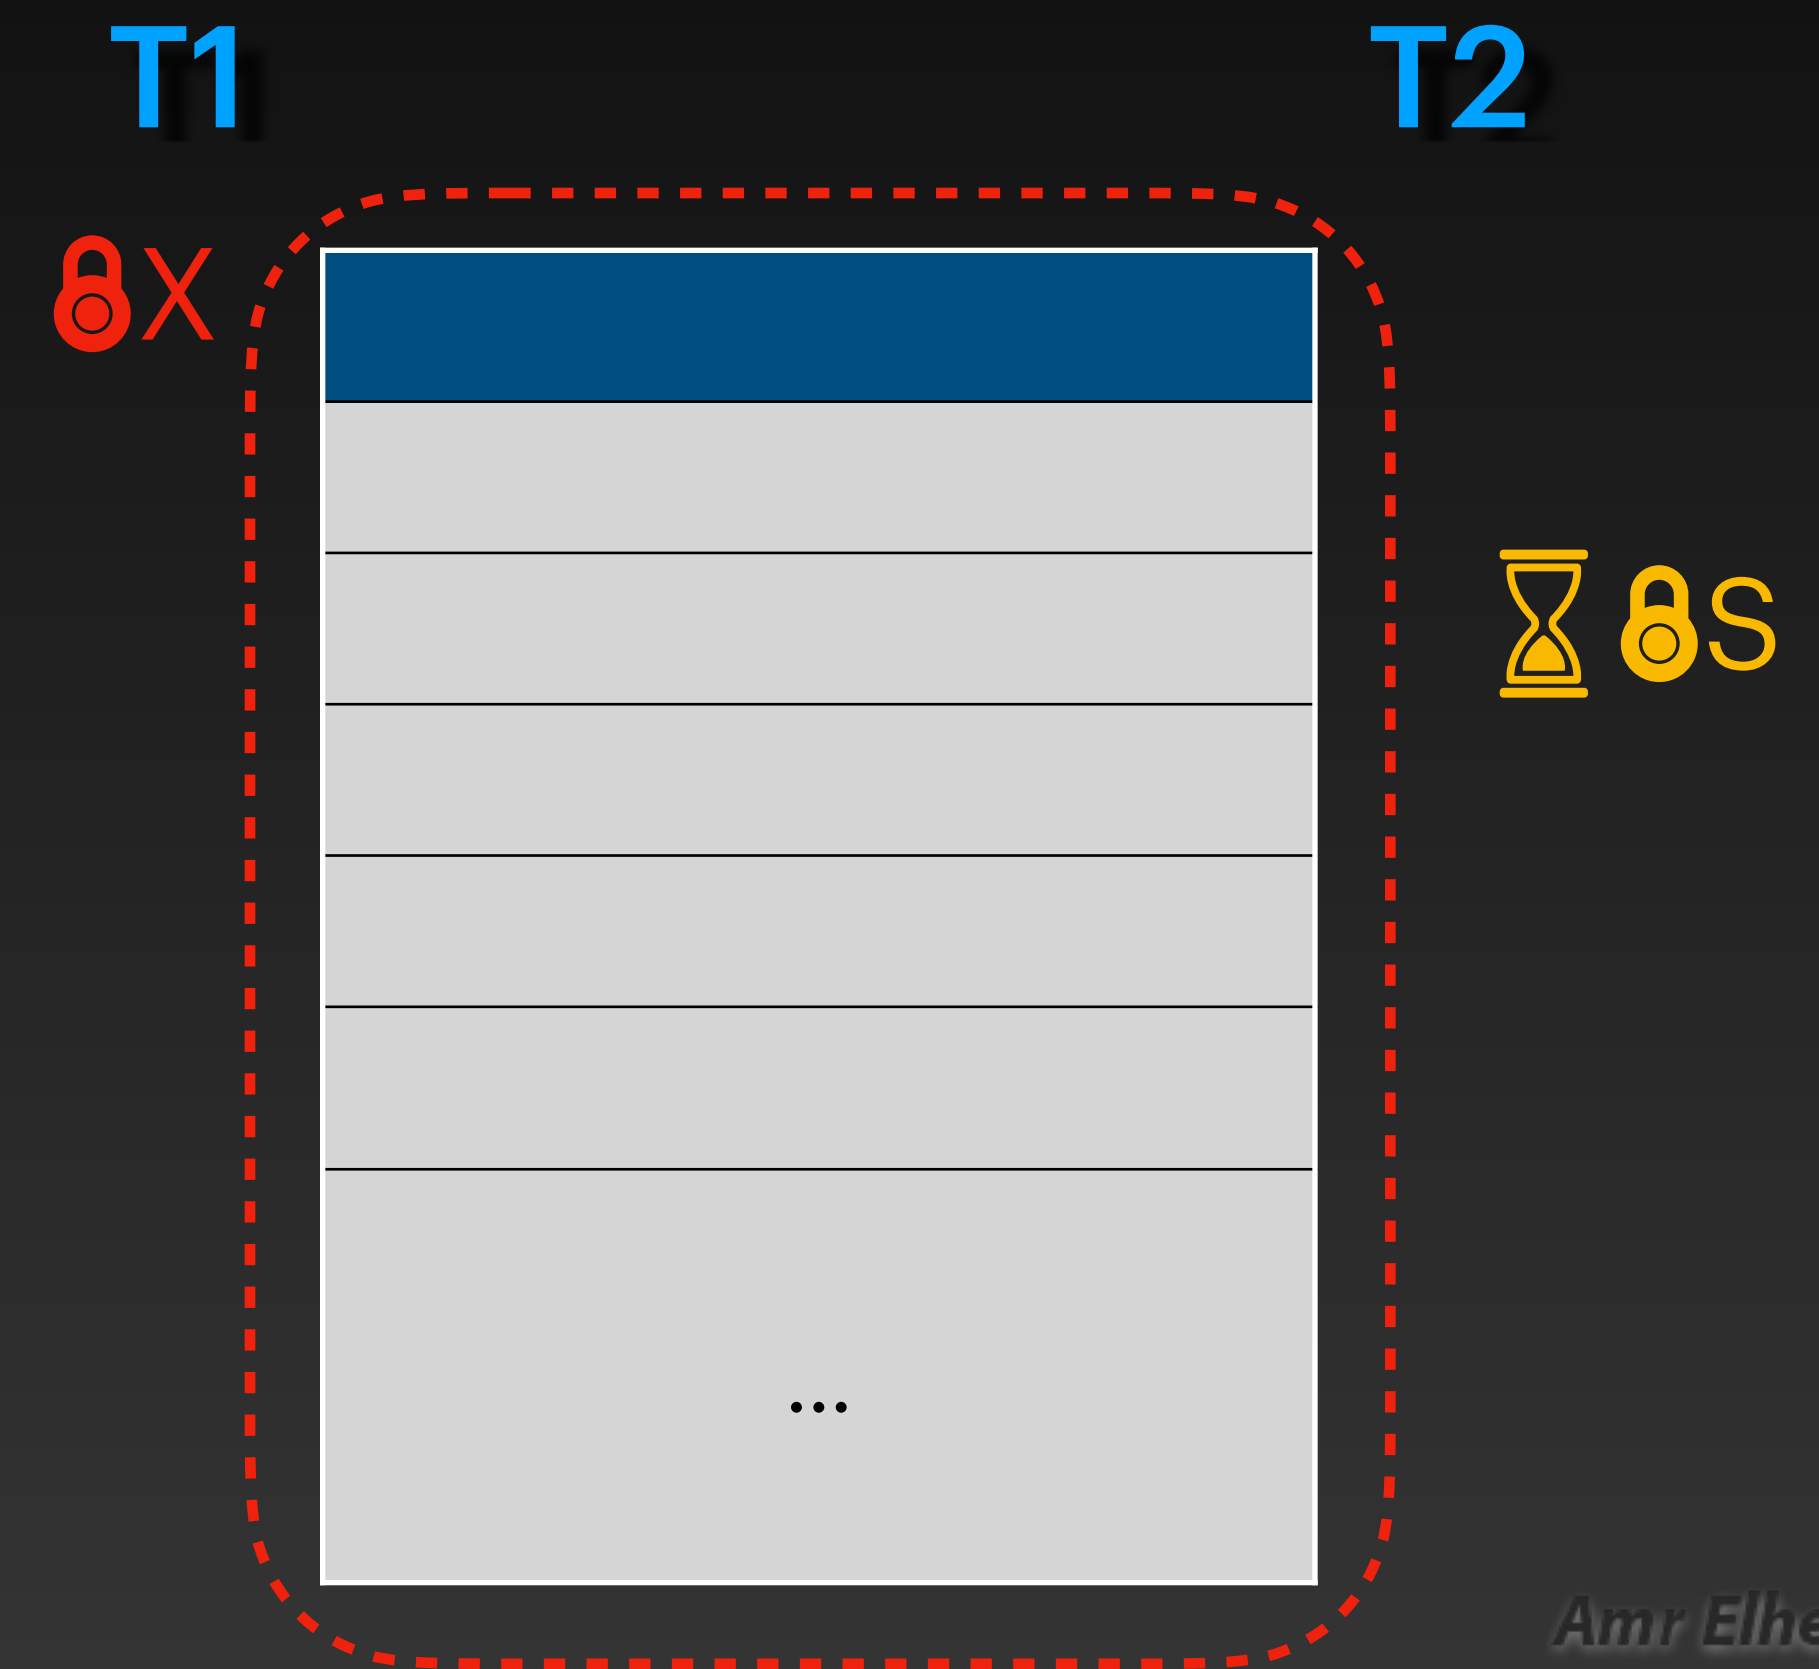
↪ Attribute (rare)

T1: scan all table and then update one row

T2: read a single row



Too many locks - expensive



Bad concurrency

Intent Lock

- A lock acquired on a higher level object to allow the transaction to lock a lower level object
- The DBMS (**lock manager**) internally manages all intent locks. The user does not have direct control over them

Intent Lock Types

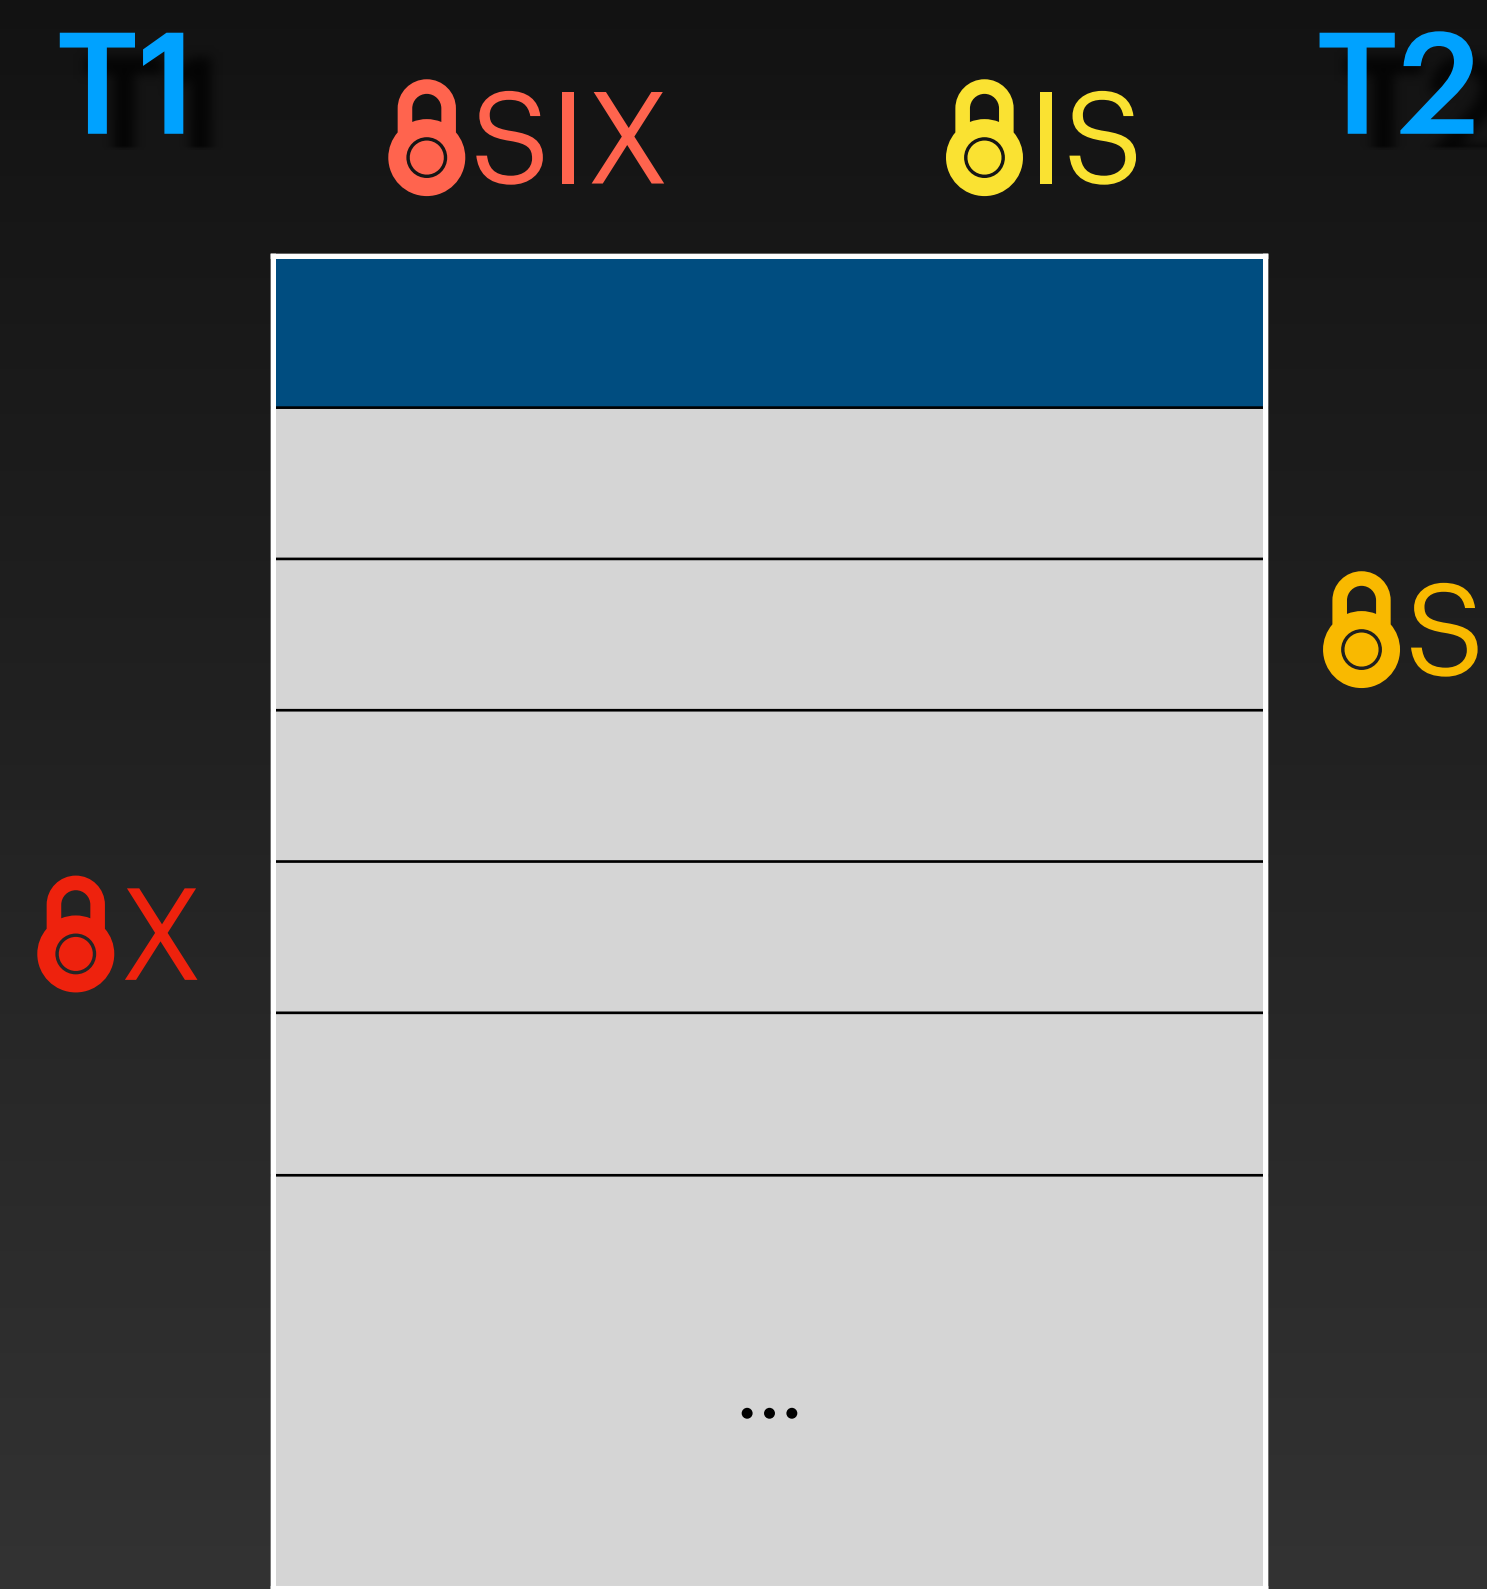
- Intent **S**hared lock (**IS**)
 - Intent to get **S-lock(s)** at a lower level
- Intent **eX**clusive lock (**IX**)
 - Intent to get **X-lock(s)** at a lower level
- **S**hared + Intent **eX**clusive lock (**SIX**)
 - **S-lock** at higher level (and all its descendants) + intent to get **X-lock(s)** at a lower level

Compatibility Matrix

		T2 Requests				
		X	S	IS	IX	SIX
T1 Holds	X	✗	✗	✗	✗	✗
	S	✗	✓	✓	✗	✗
	IS	✗	✓	✓	✓	✓
	IX	✗	✗	✓	✓	✗
	SIX	✗	✗	✓	✗	✗

T1: scan all table and then update one row

T2: read a single row



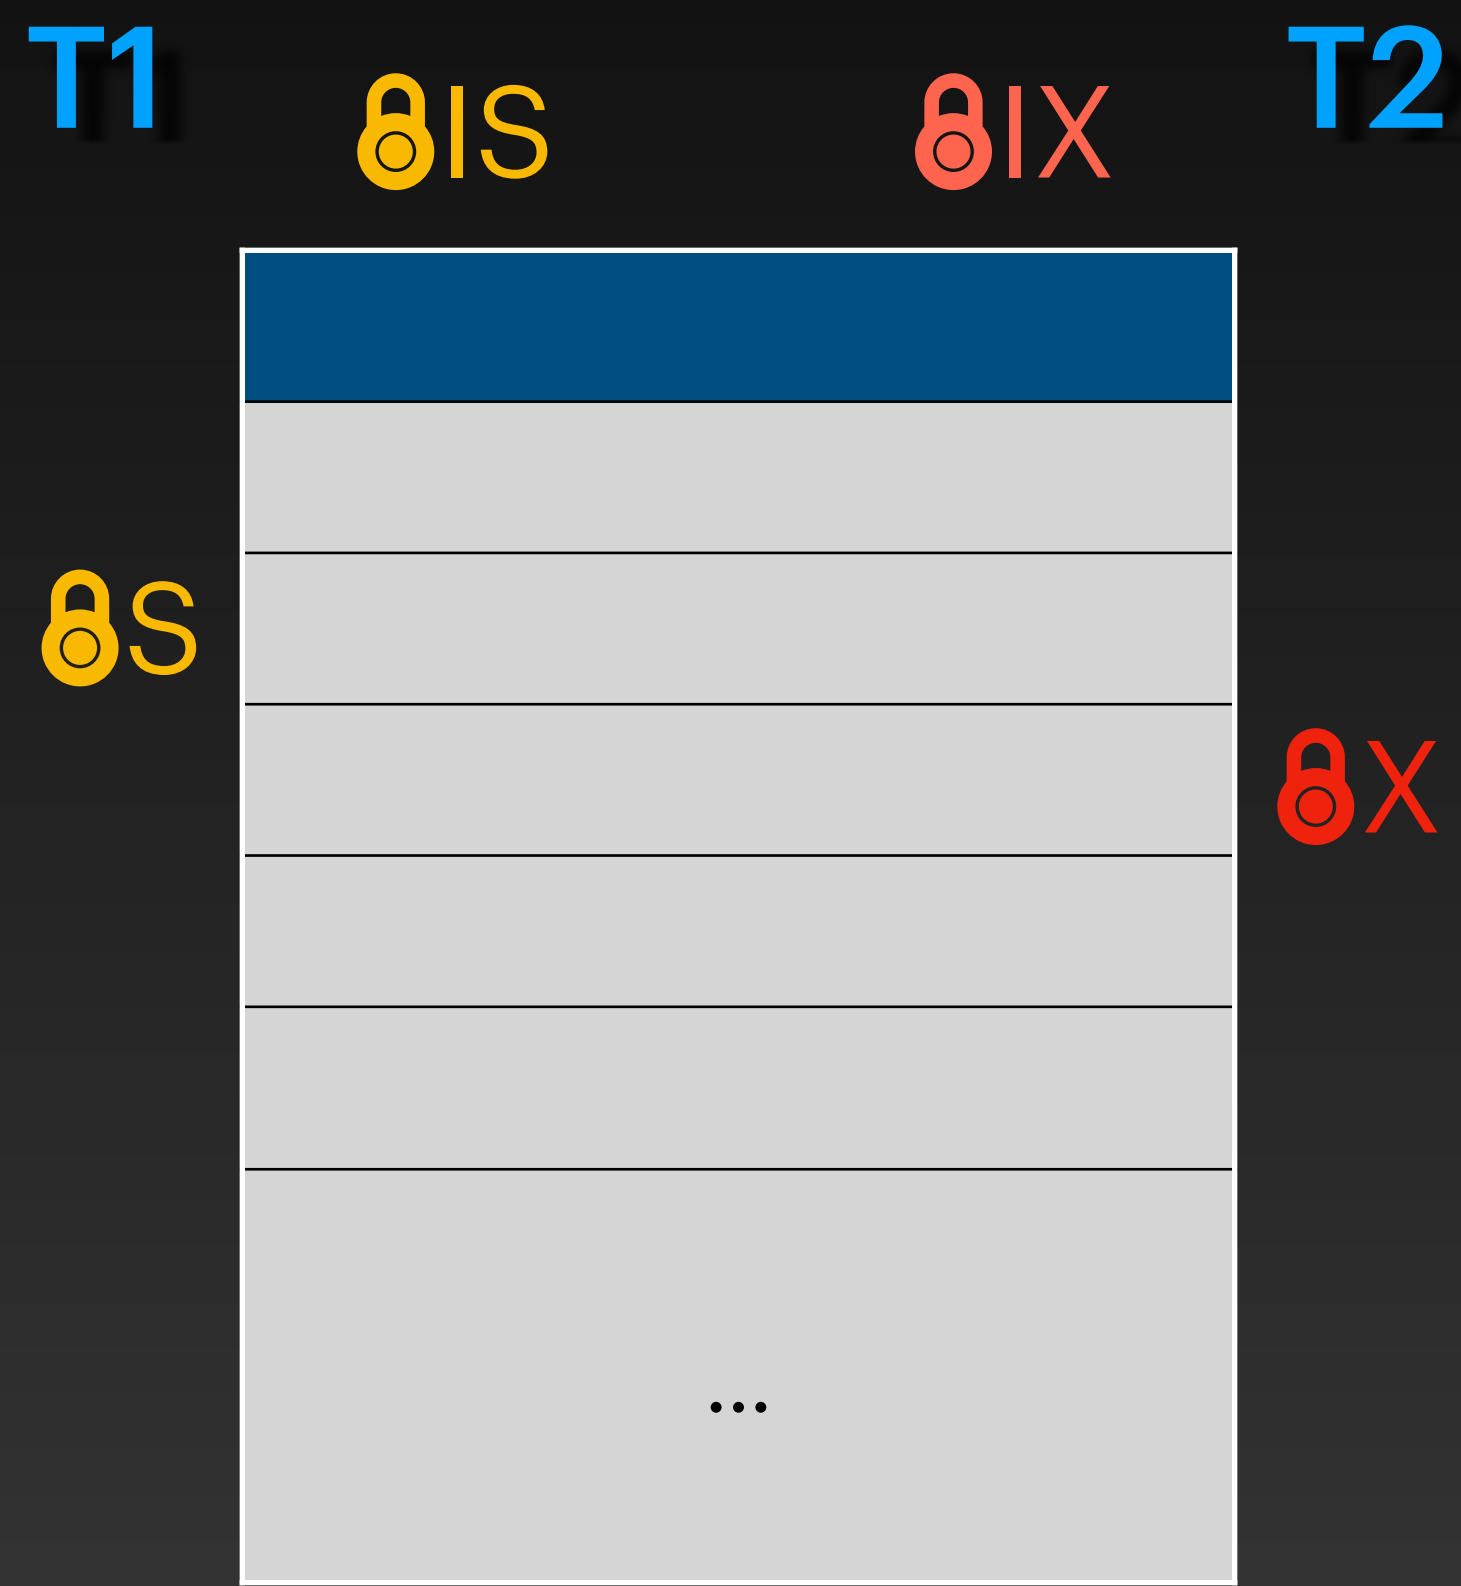
Transaction requests

	X	S	IS	IX	SIX
X	✗	✗	✗	✗	✗
S	✗	✓	✓	✗	✗
IS	✗	✓	✓	✓	✓
IX	✗	✗	✓	✓	✗
SIX	✗	✗	✓	✗	✗

Transaction holds

Amr Elhelw's
TECH
VAULT

- T1: read a single row
- T2: update a single (different) row



Transaction requests

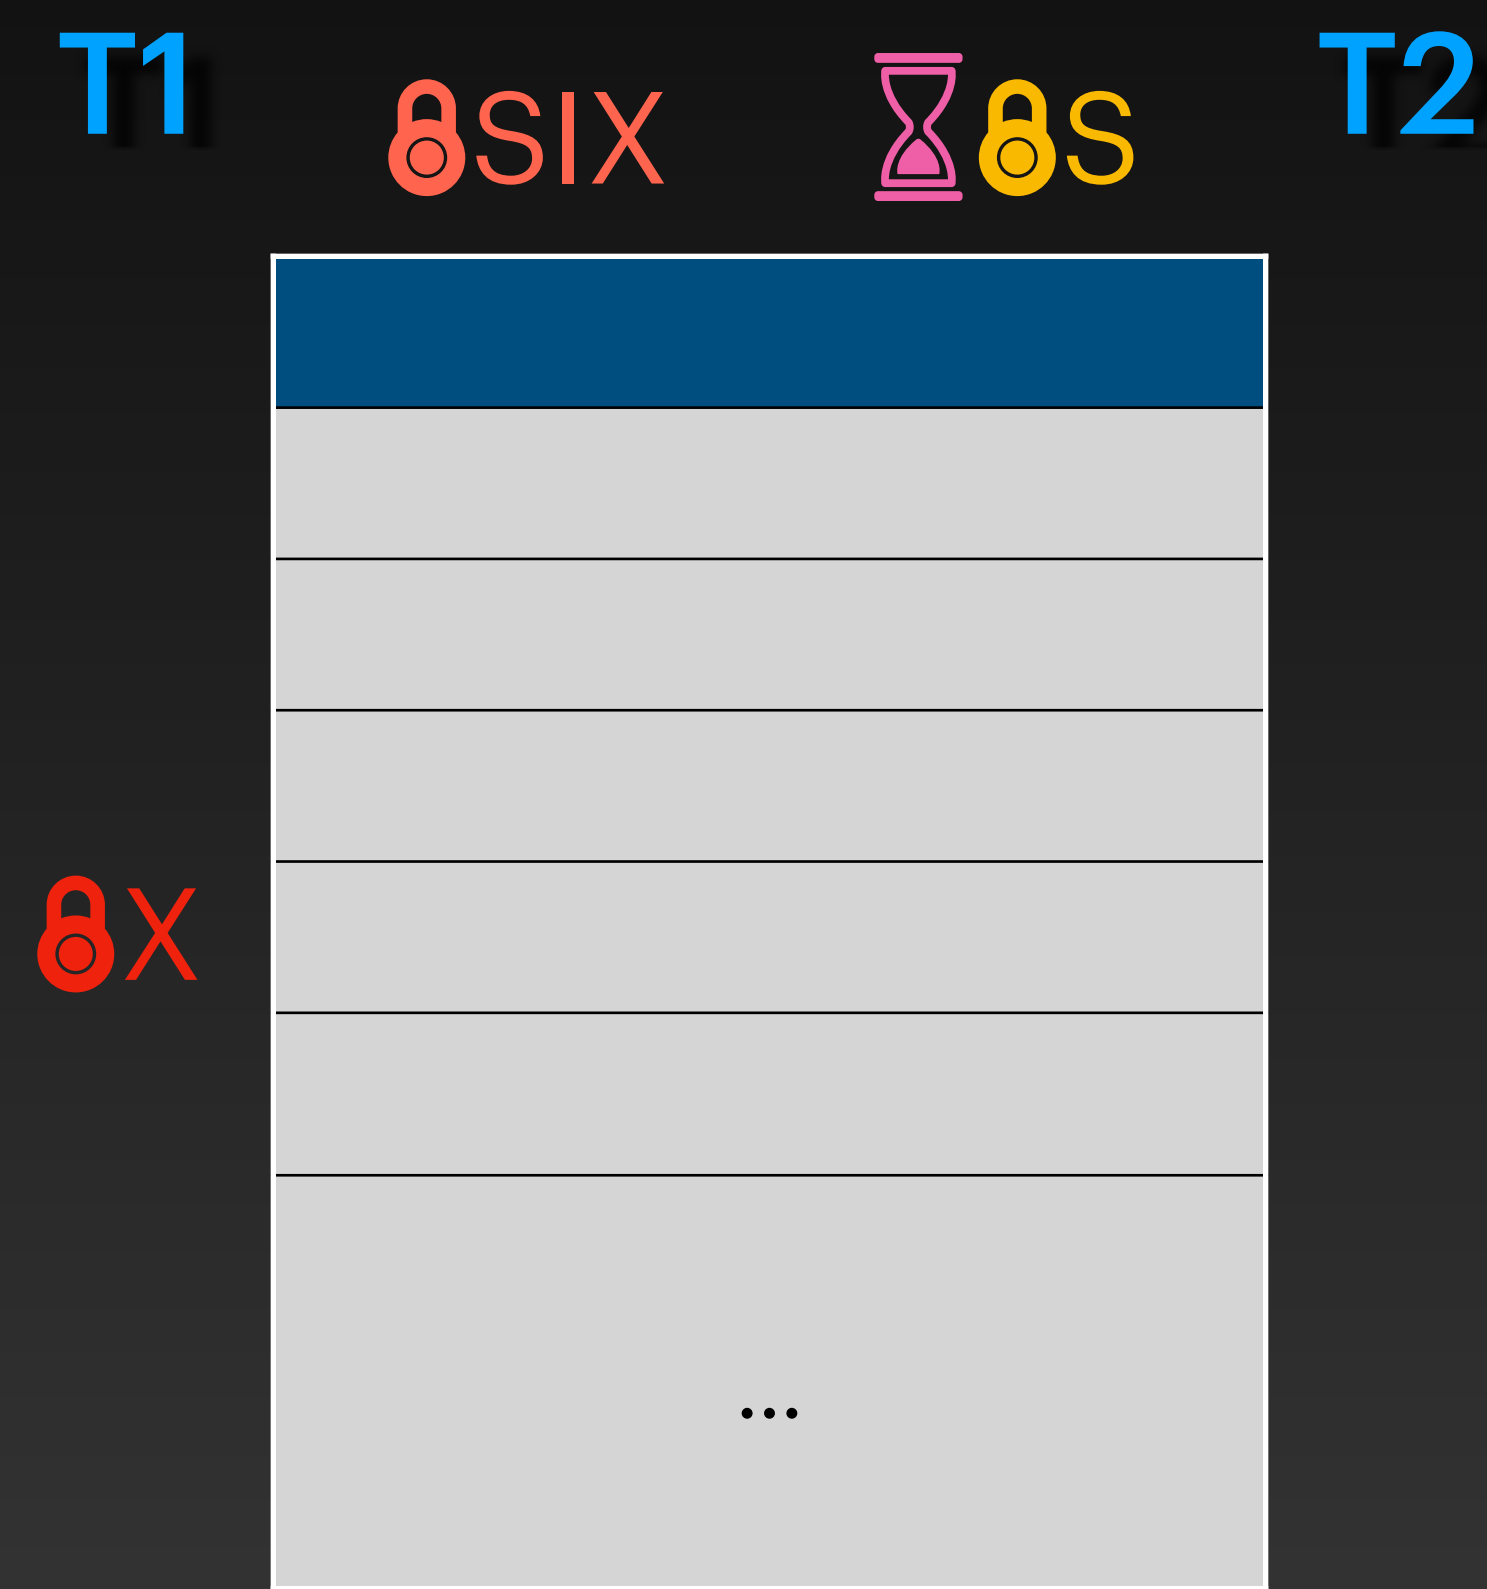
	X	S	IS	IX	SIX
X	✗	✗	✗	✗	✗
S	✗	✓	✓	✗	✗
IS	✗	✓	✓	✓	✓
IX	✗	✗	✓	✓	✗
SIX	✗	✗	✓	✗	✗

Transaction holds

Amr Elhelw's
TECH
VAULT

T1: scan all table and then update one row

T2: scan all table



Transaction requests

	X	S	IS	IX	SIX
X	✗	✗	✗	✗	✗
S	✗	✓	✓	✗	✗
IS	✗	✓	✓	✓	✓
IX	✗	✗	✓	✗	✗
SIX	✗	✗	✓	✗	✗

Transaction holds

Amr Elhelw's
TECH
VAULT

MySQL

- Table-level locking

```
LOCK TABLE customers READ;
```

```
LOCK TABLES customers READ, items WRITE;
```

```
UNLOCK TABLES;
```

- Row-level locking

```
SELECT * FROM users WHERE NAME = 'Jones' FOR SHARE;
```

```
SELECT * FROM t WHERE i = 2 FOR UPDATE;
```

PostgreSQL

- Table-level locking

`LOCK TABLE customers IN <lockmode> MODE;`

<lockmode> can be:

ACCESS SHARE | ROW SHARE | ROW EXCLUSIVE | SHARE UPDATE EXCLUSIVE | SHARE |
SHARE ROW EXCLUSIVE | EXCLUSIVE | ACCESS EXCLUSIVE

- Row-level locking

`SELECT ... FOR <mode>;`

<mode> can be:

UPDATE | NO KEY UPDATE | SHARE | KEY SHARE

- More info and compatibility matrices

<https://www.postgresql.org/docs/current/explicit-locking.html>