

CONCURRENCY CONTROL

Multi Version
Concurrency Control


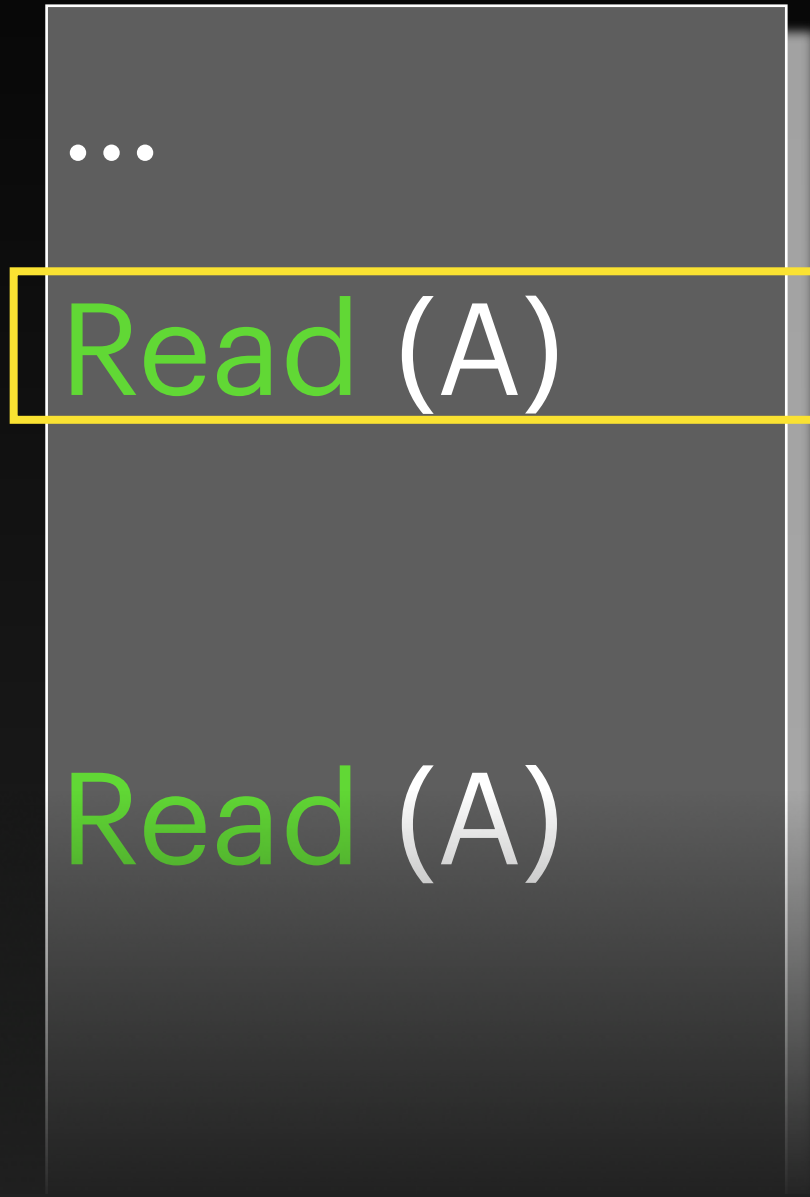


AMR ELHELW

T1

T2

Time



A = 100

T1

T2

...

Read (A)

...

Write (A)

Read (A)

A = 500

Time

T1

T2

...

Read (A)

...

Write (A)

Read (A)

A = 500



T1

T2

...

Read (A)

Read (A)

...

Write (A)

$A_0 = 100$

$A_1 = 500$



Time

MVCC with Timestamps

- Each transaction is assigned a timestamp ts
- Timestamps are monotonically increasing
- If $ts(T_i) < ts(T_j)$ then we want a schedule equivalent to the serial schedule where T_i runs before T_j

MVCC with Timestamps

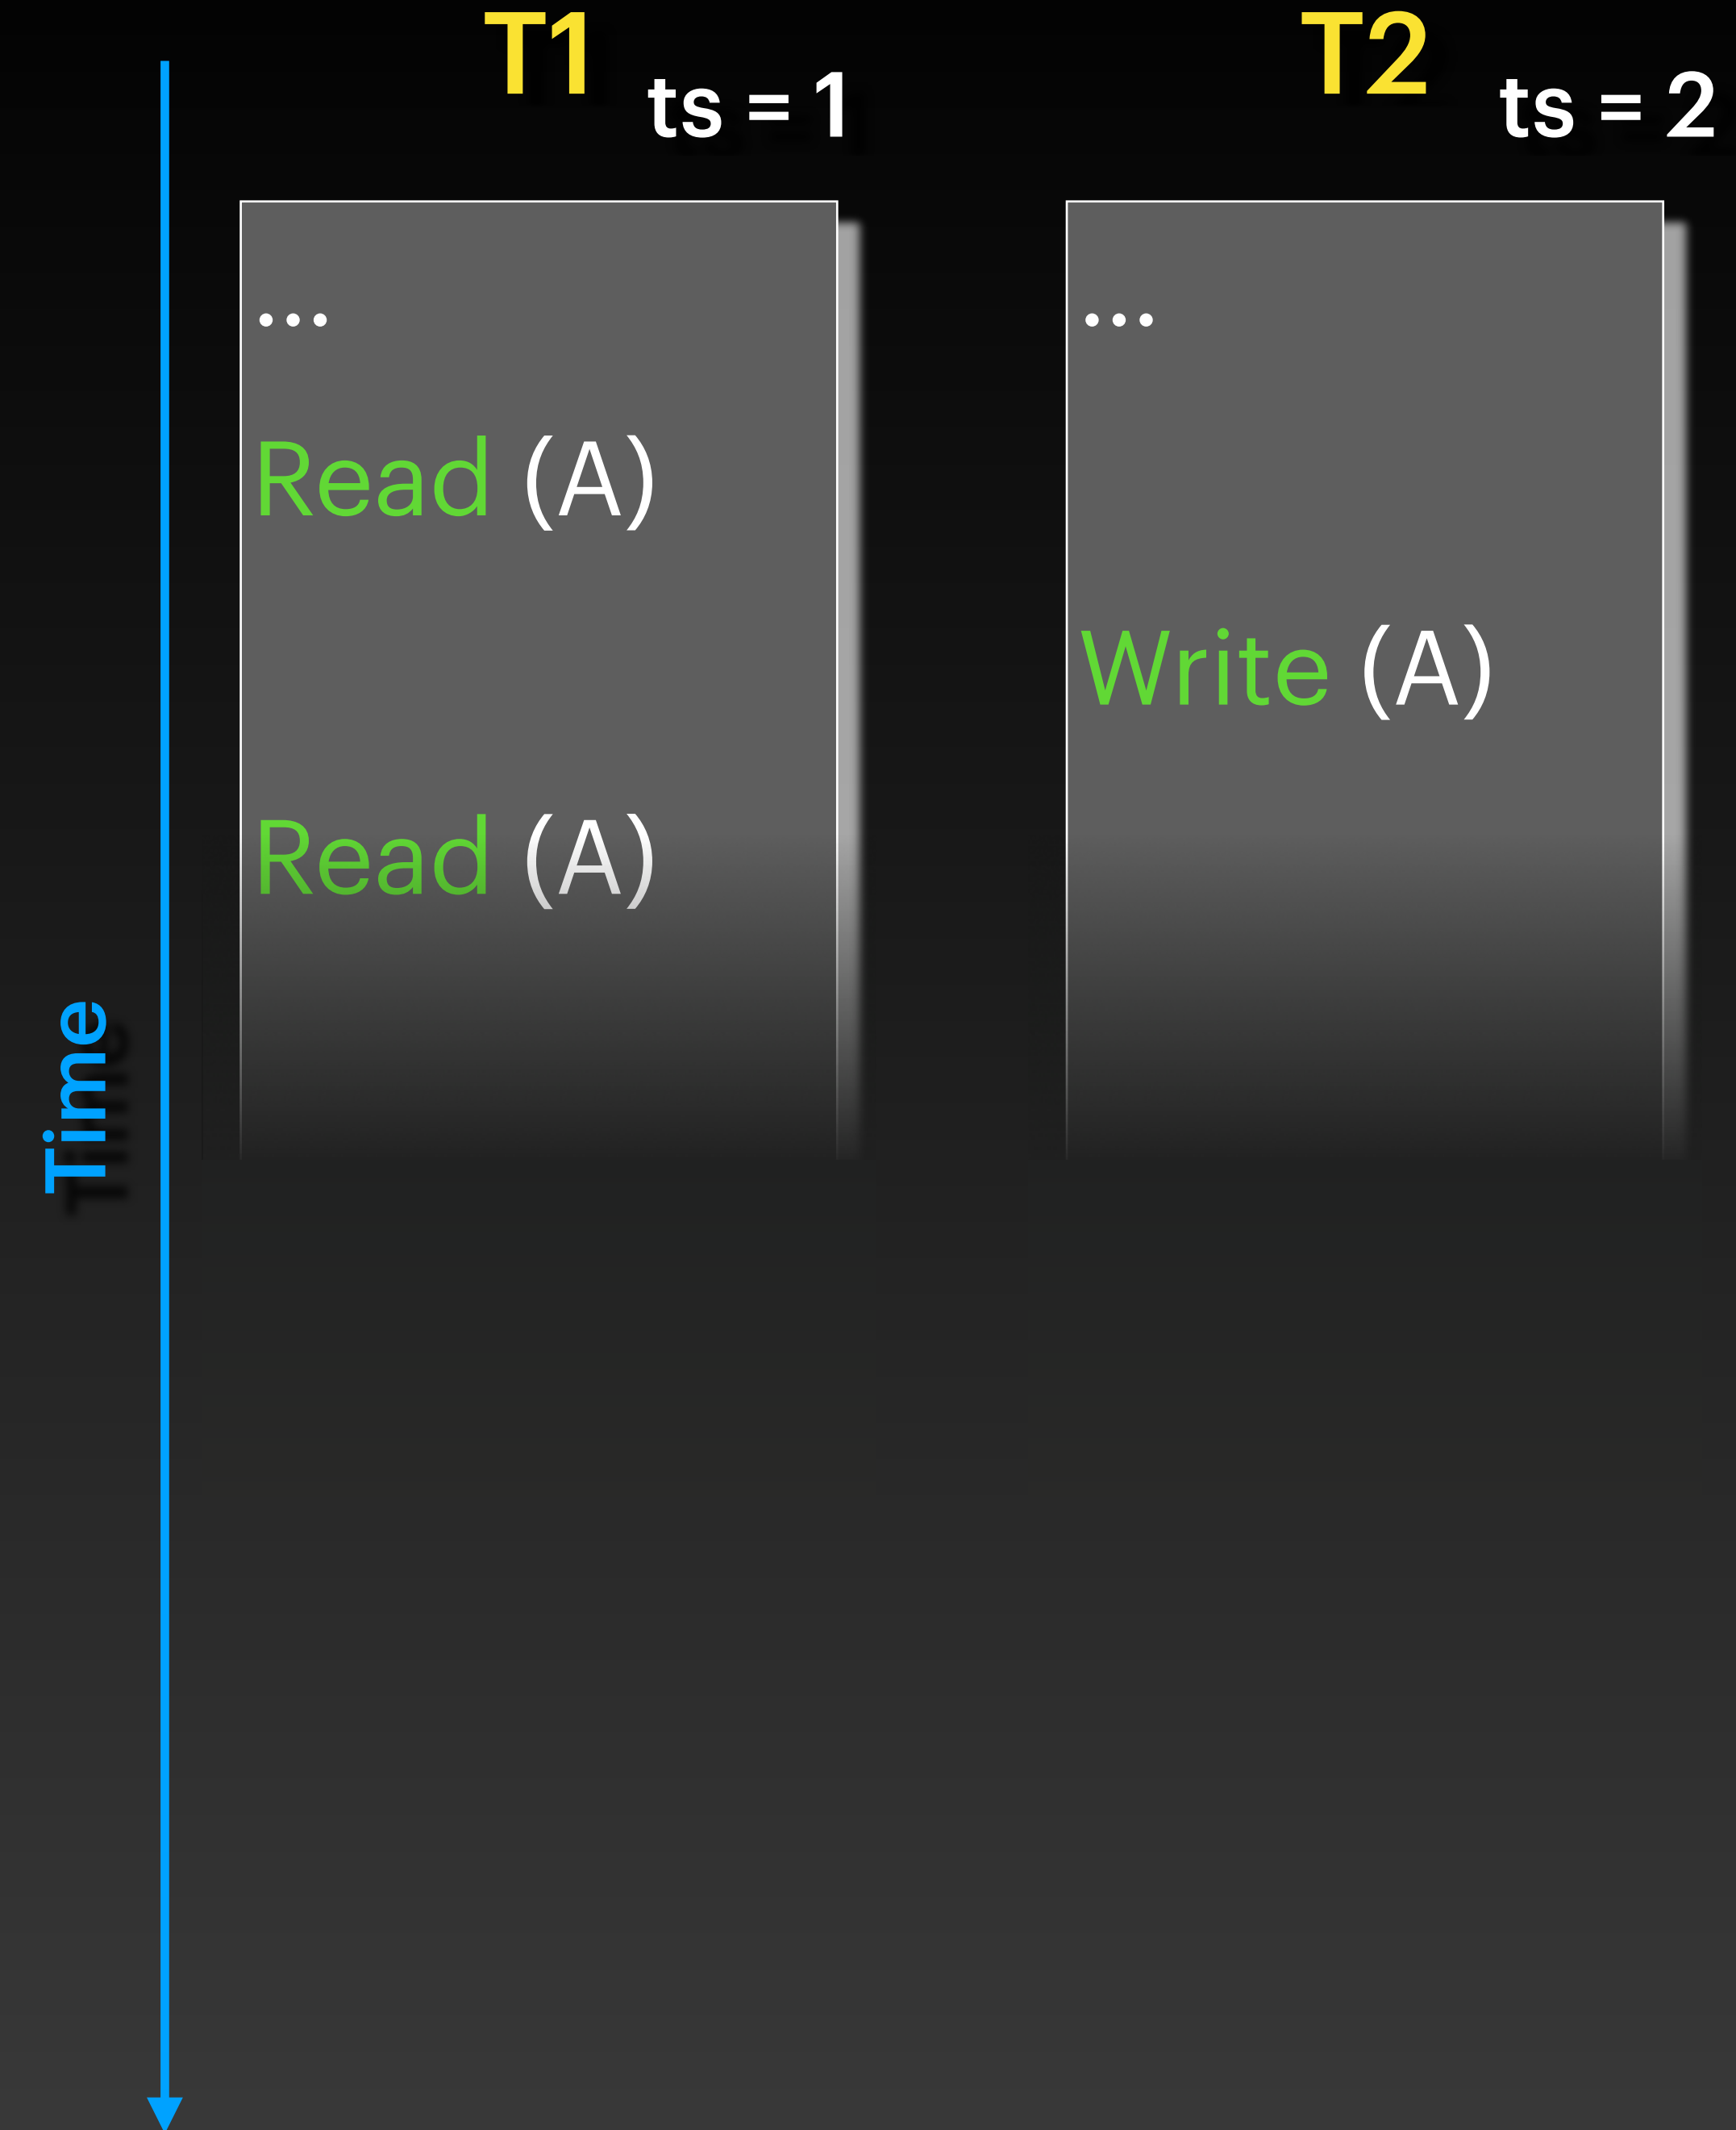
- For each data item X , we maintain several versions X_0, X_1, X_2, \dots
- For each version, we maintain:
 - *Read timestamp* **R-TS**
 - Largest timestamp of a transaction that read this version
 - *Write timestamp* **W-TS**
 - Timestamp of the transaction that wrote this version

MVCC: Reads

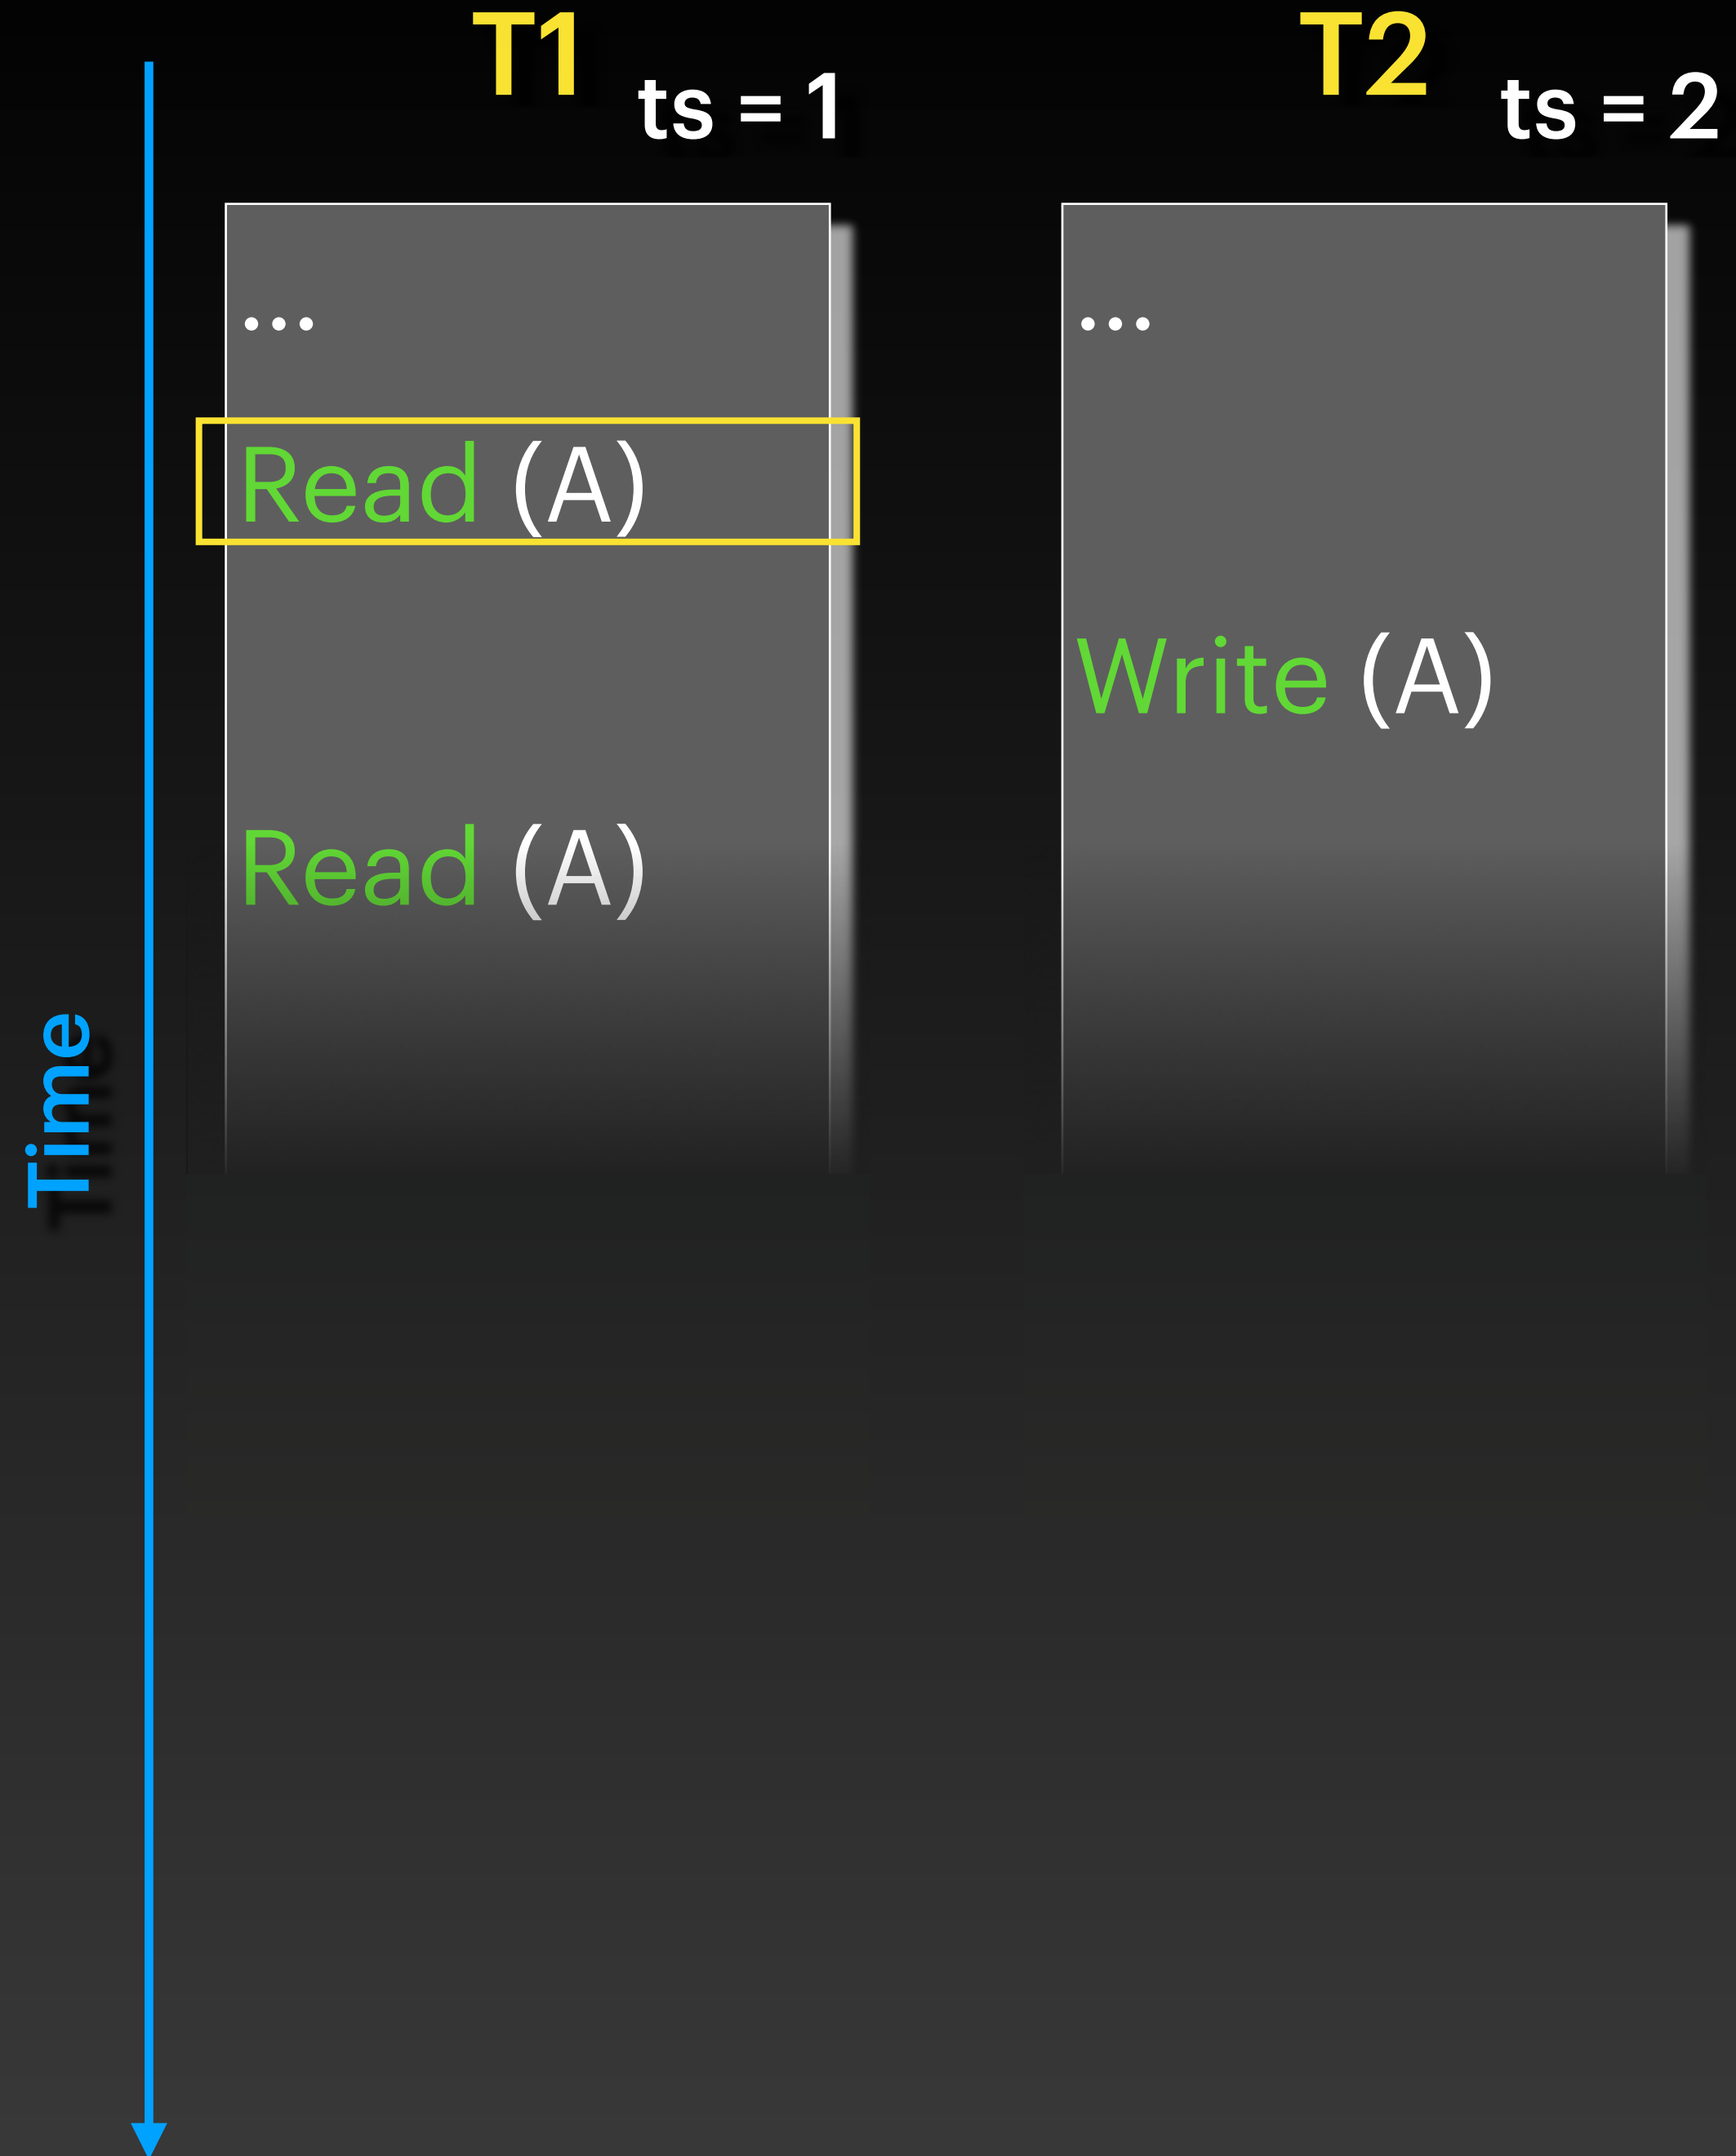
- Transaction T wants to read object X
- Read the latest version X_i that was written before
 T started (or during T)
 - i.e. $W-TS(X_i)$ has the highest value $\leq ts(T)$
- Set $R-TS(X_i)$ to $\max(R-TS(X_i), ts(T))$

MVCC: Writes

- Transaction T wants to write object X
- Let X_i be the version that has the highest $W-TS(X_i)$ that is also $\leq ts(T)$
- If $R-TS(X_i) > ts(T)$, then Abort T
- Otherwise, create a new version X_j
 - Set $R-TS(X_j) = W-TS(X_j) = ts(T)$



Version	Value	R-TS	W-TS
A ₀	100	0	0



Version	Value	R-TS	W-TS
A_0	100	1	0

Time

T1 $ts = 1$

...

Read (A)

Read (A)

T2 $ts = 2$

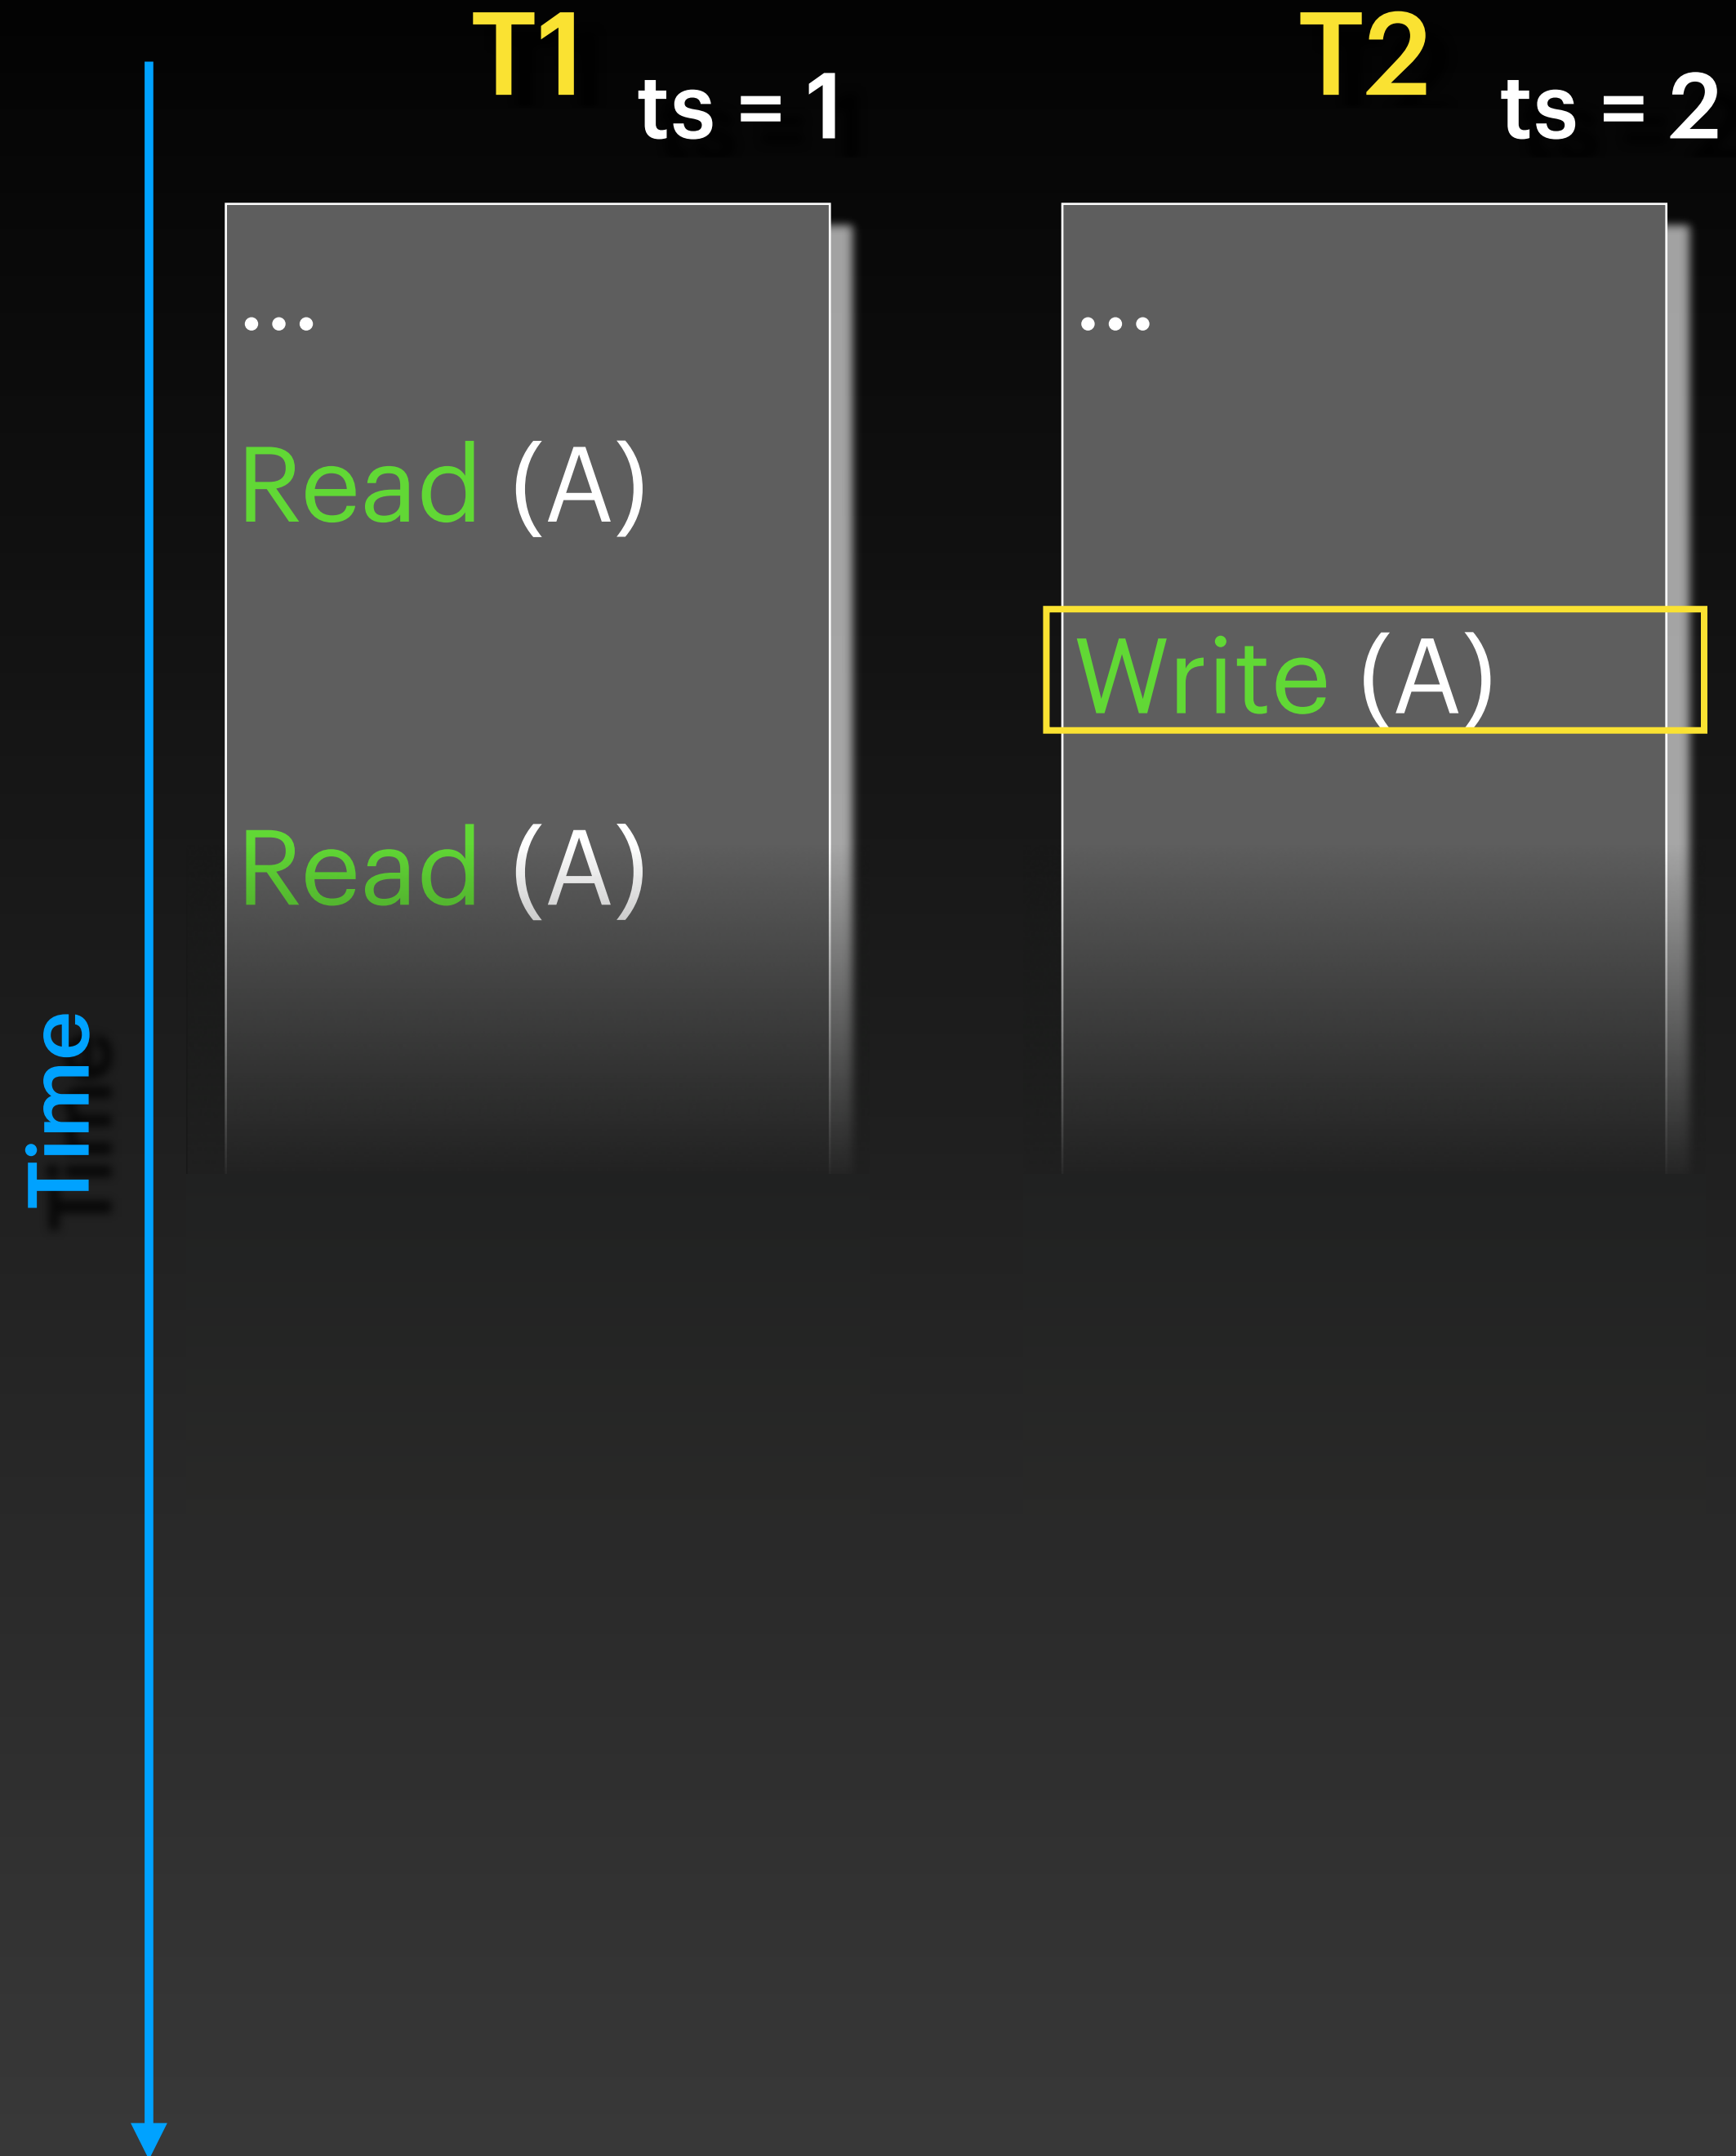
...

Write (A)

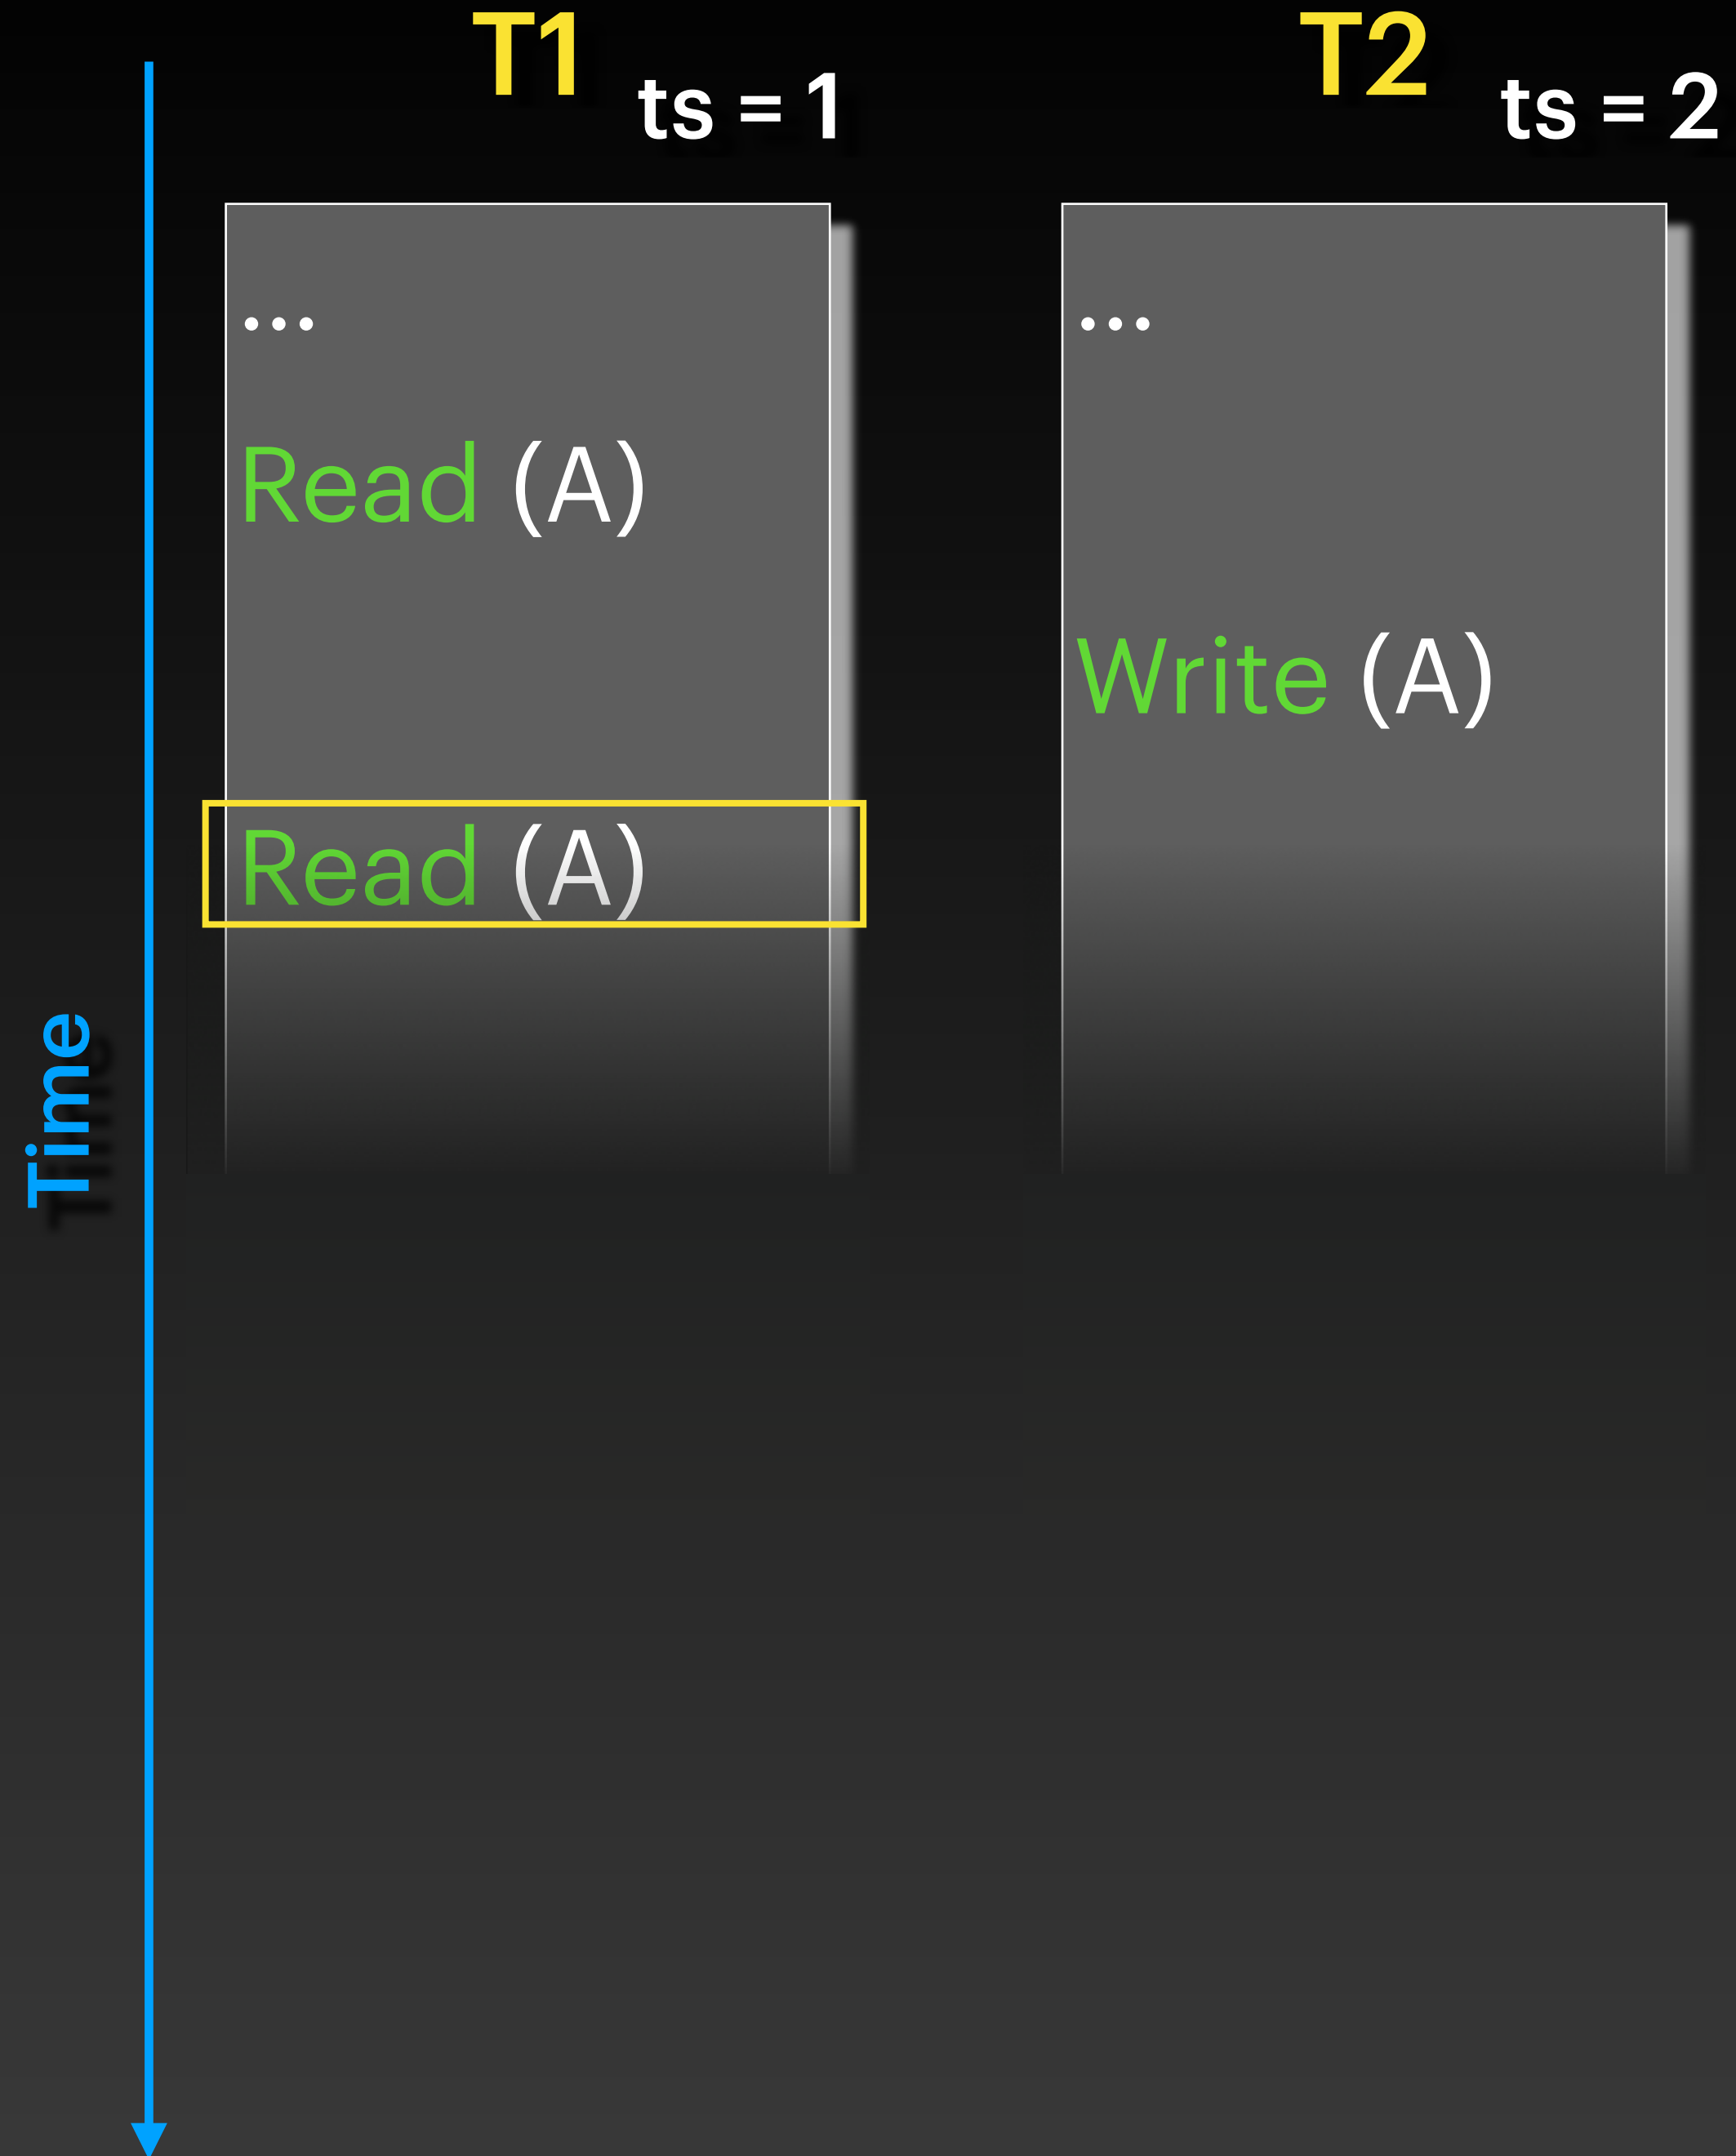
Version	Value	R-TS	W-TS
A₀	100	1	0

$R-TS < ts(T)$

Highest W-TS
less than or
equal $ts(T)$



Version	Value	R-TS	W-TS
A_0	100	1	0
A_1	500	2	2



Version	Value	R-TS	W-TS
A_0	100	1	0
A_1	500	2	2

Highest W-TS
 $\leq ts(T)$

Time

T1 ts = 1

Read (A)

Write (A)

Read (A)

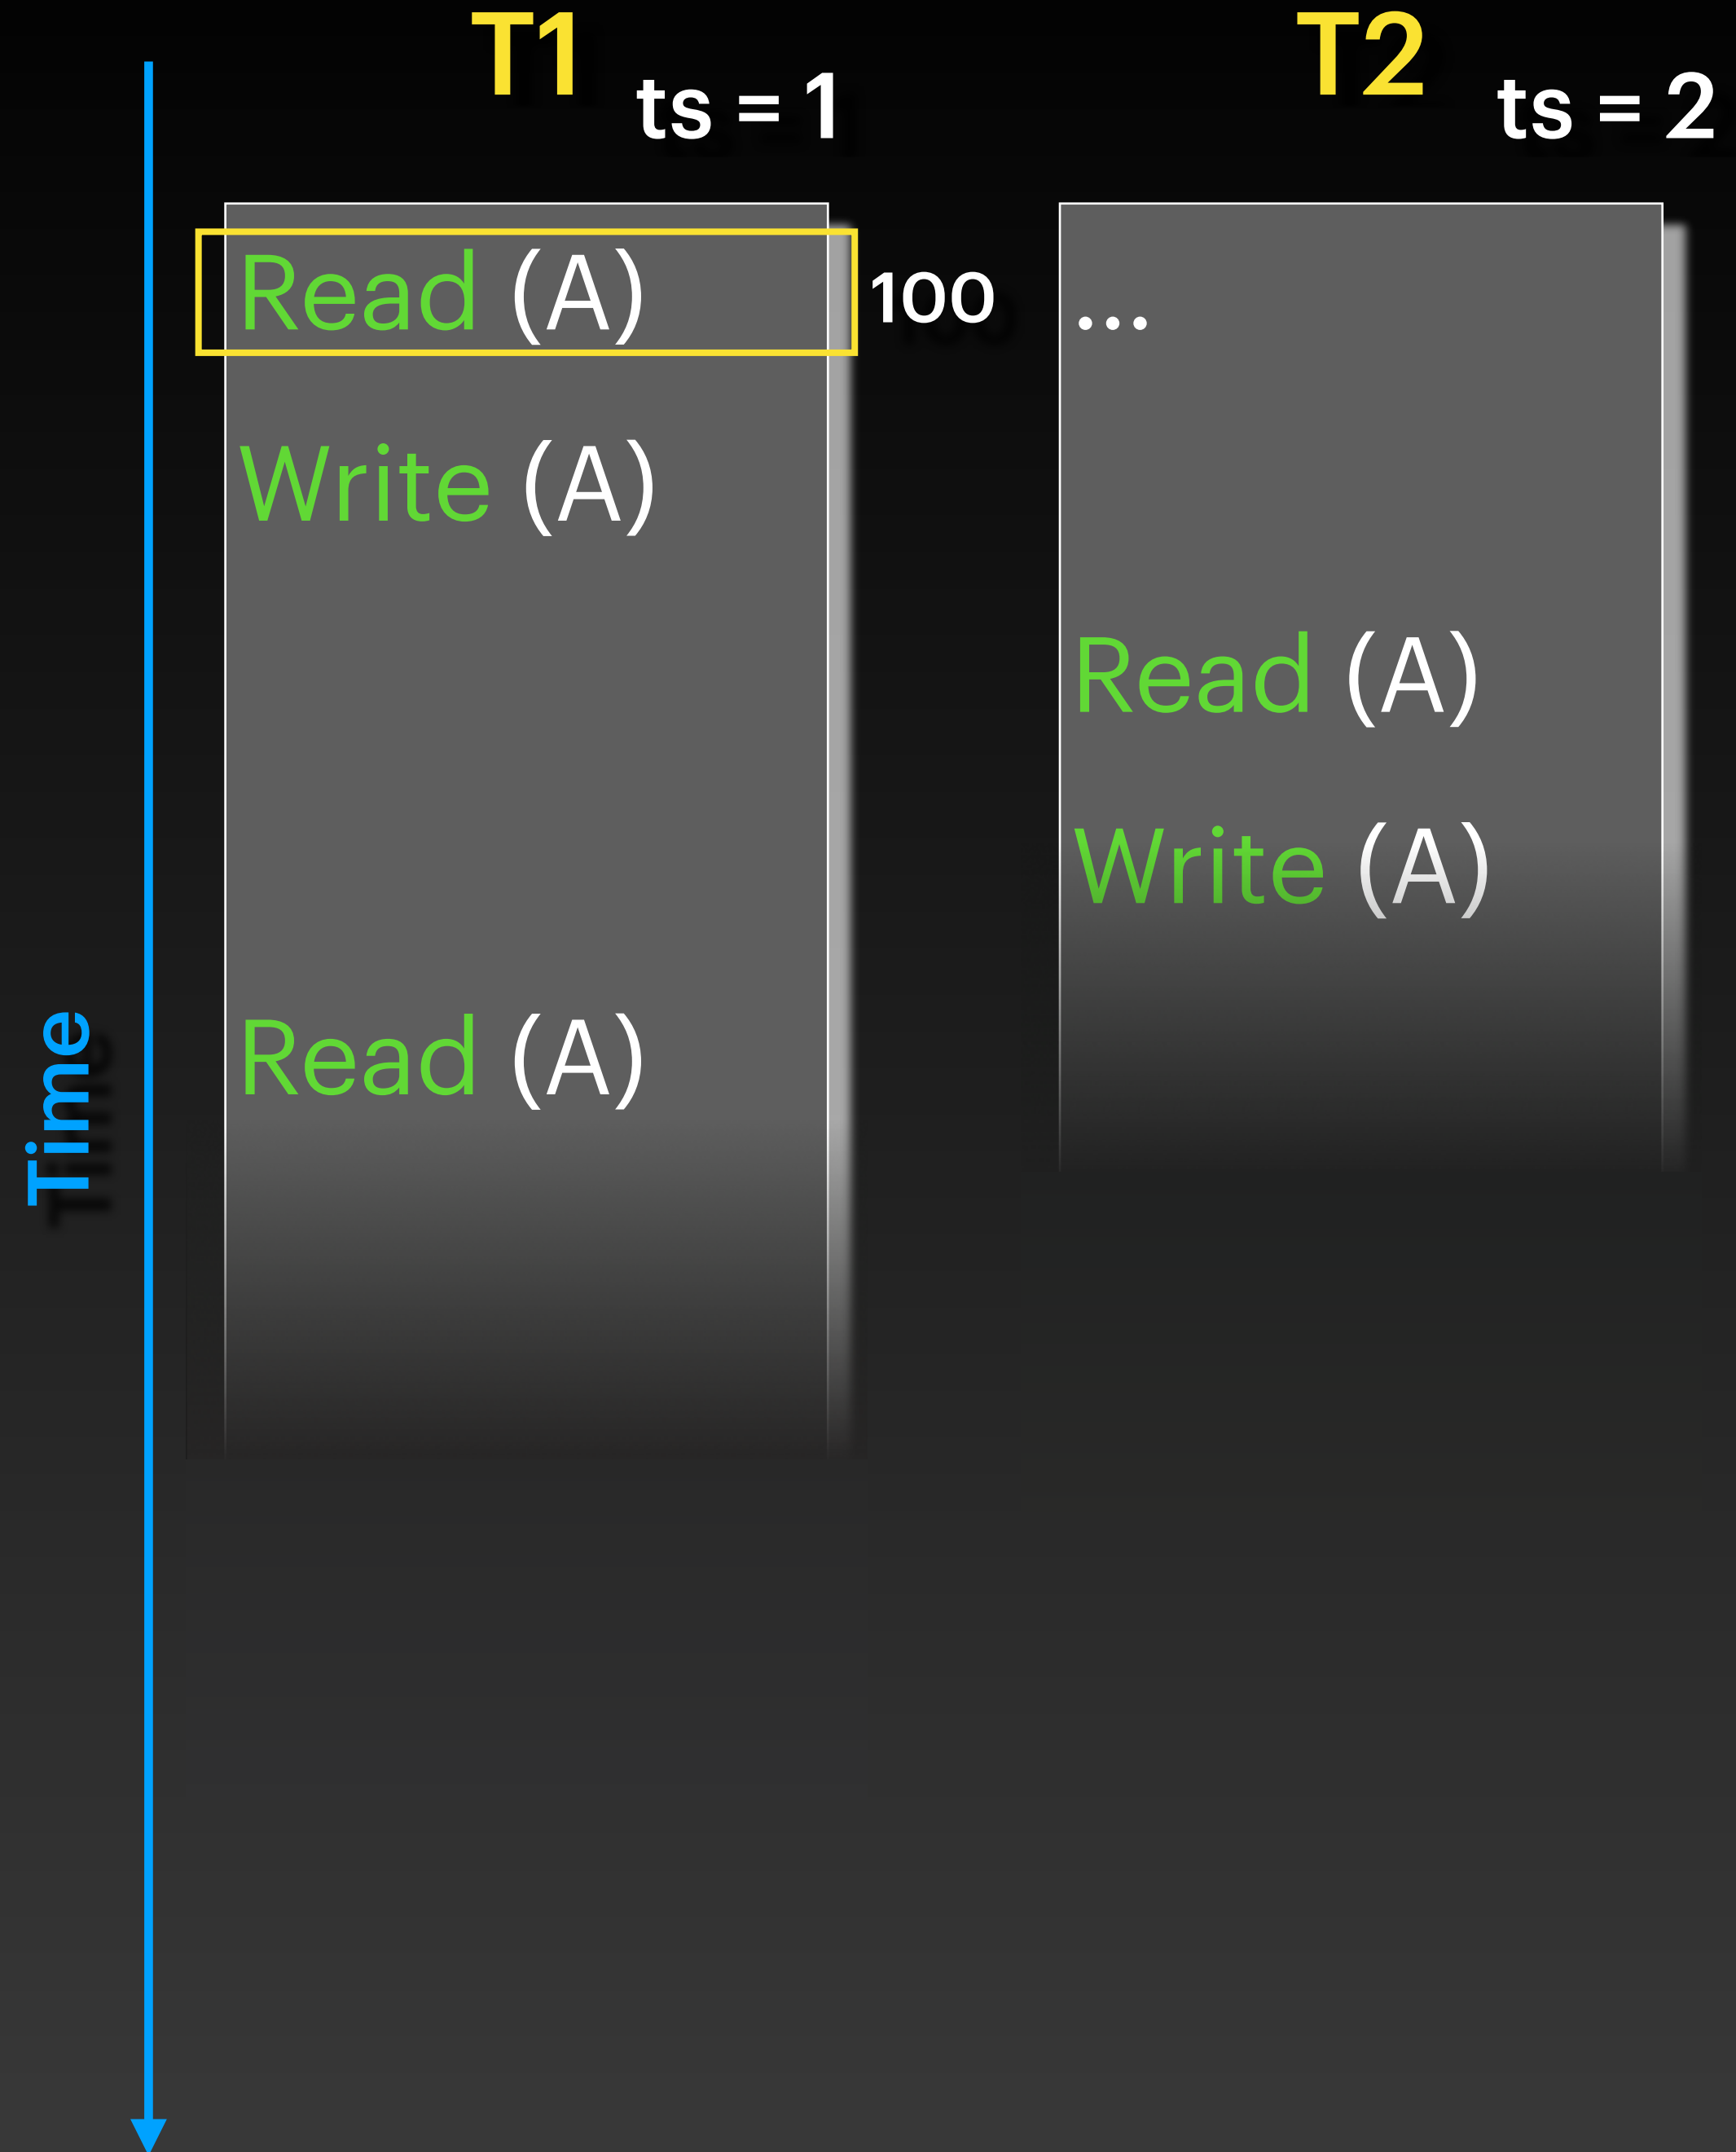
T2 ts = 2

...

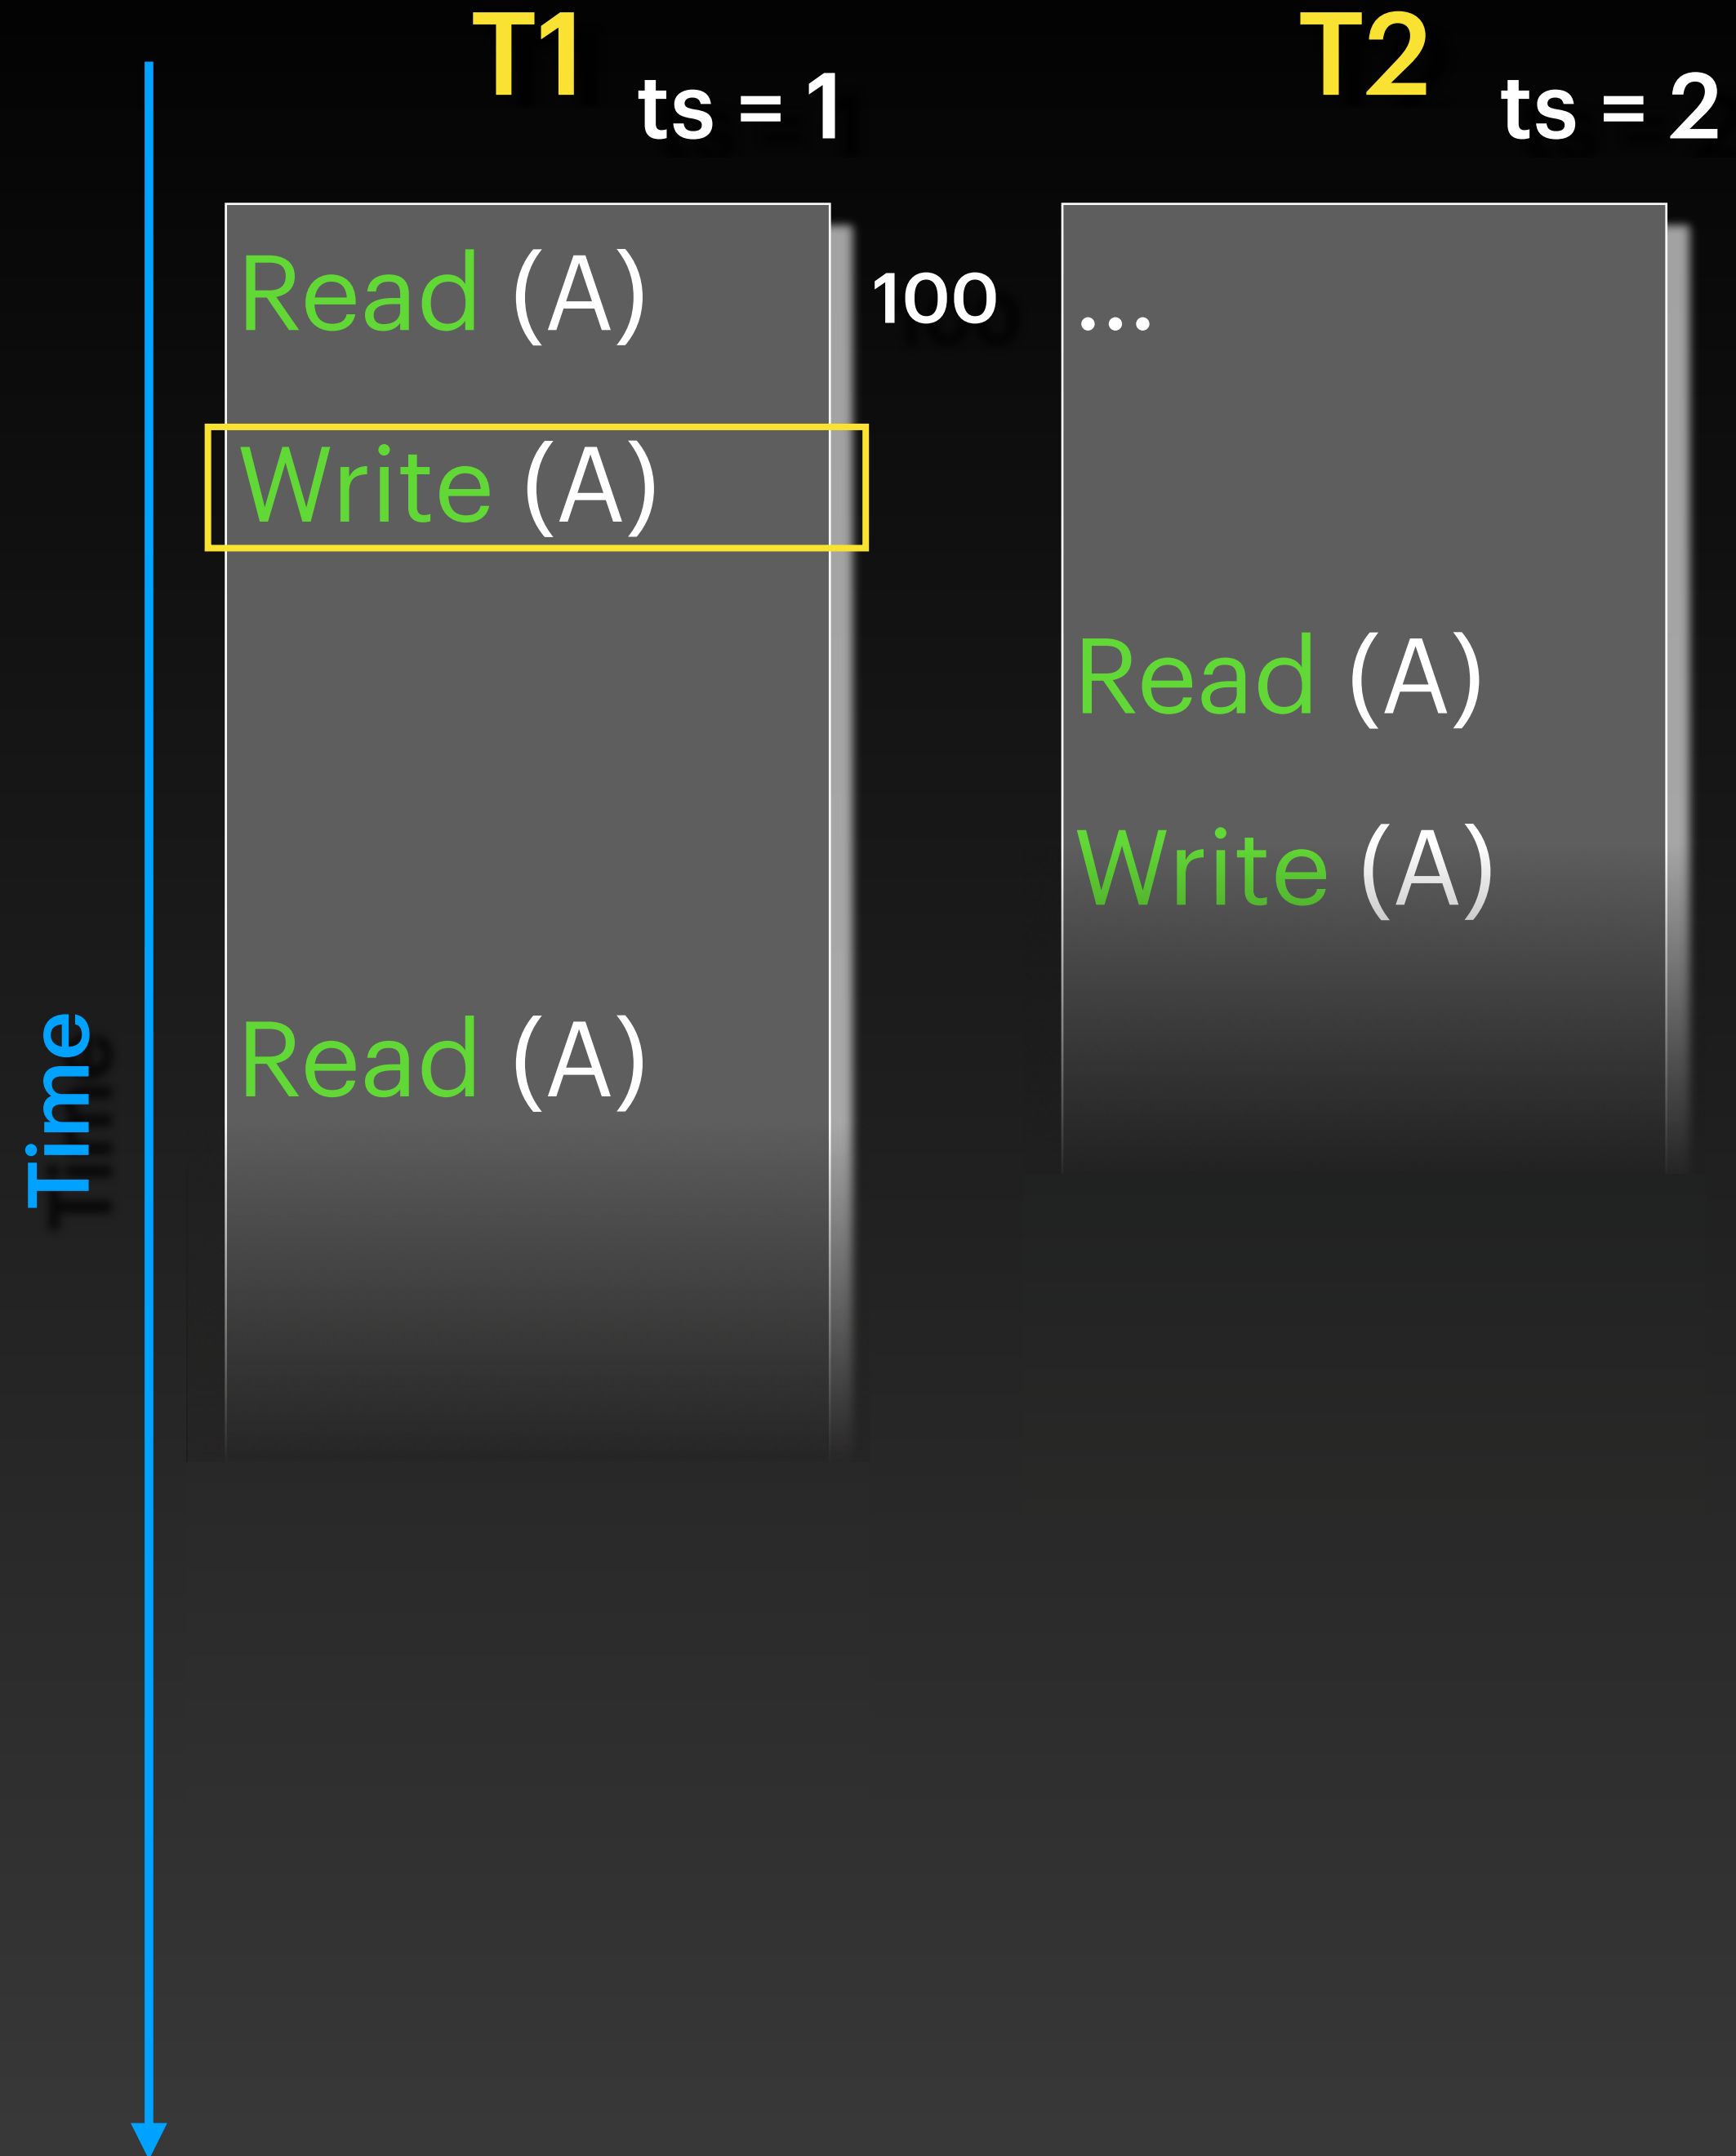
Read (A)

Write (A)

Version	Value	R-TS	W-TS
A₀	100	0	0

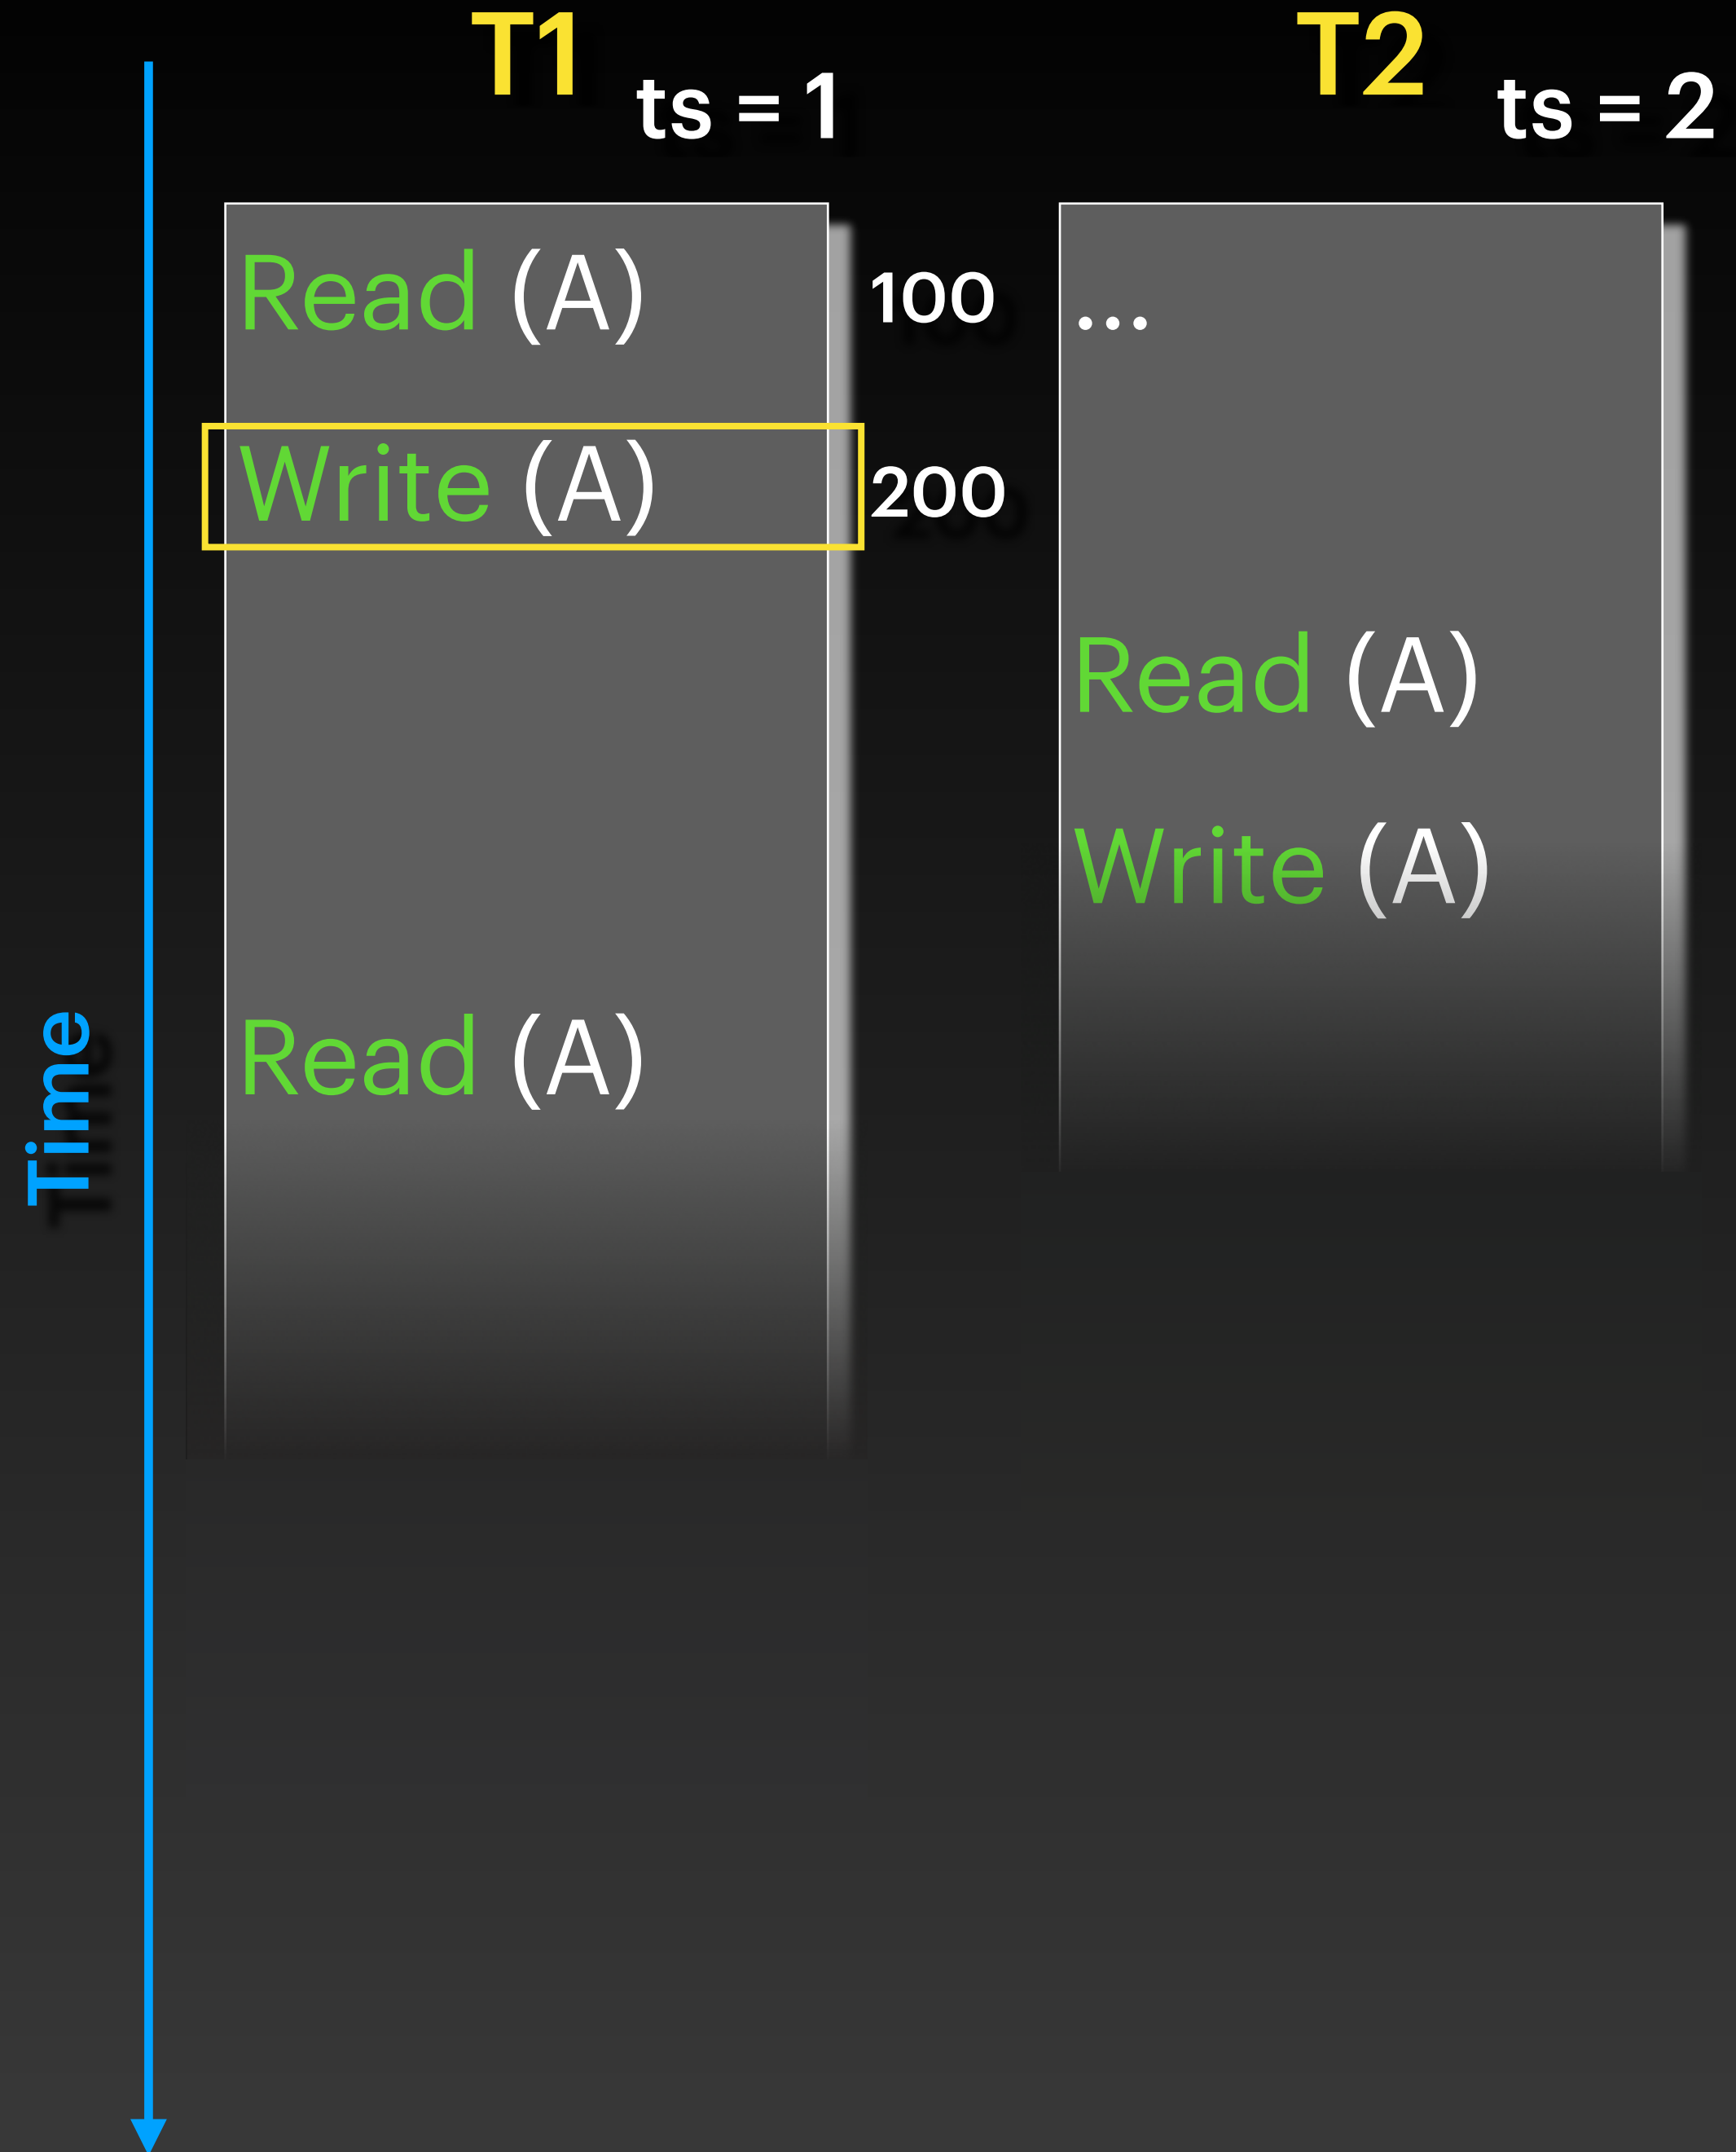


Version	Value	R-TS	W-TS
A_0	100	1	0

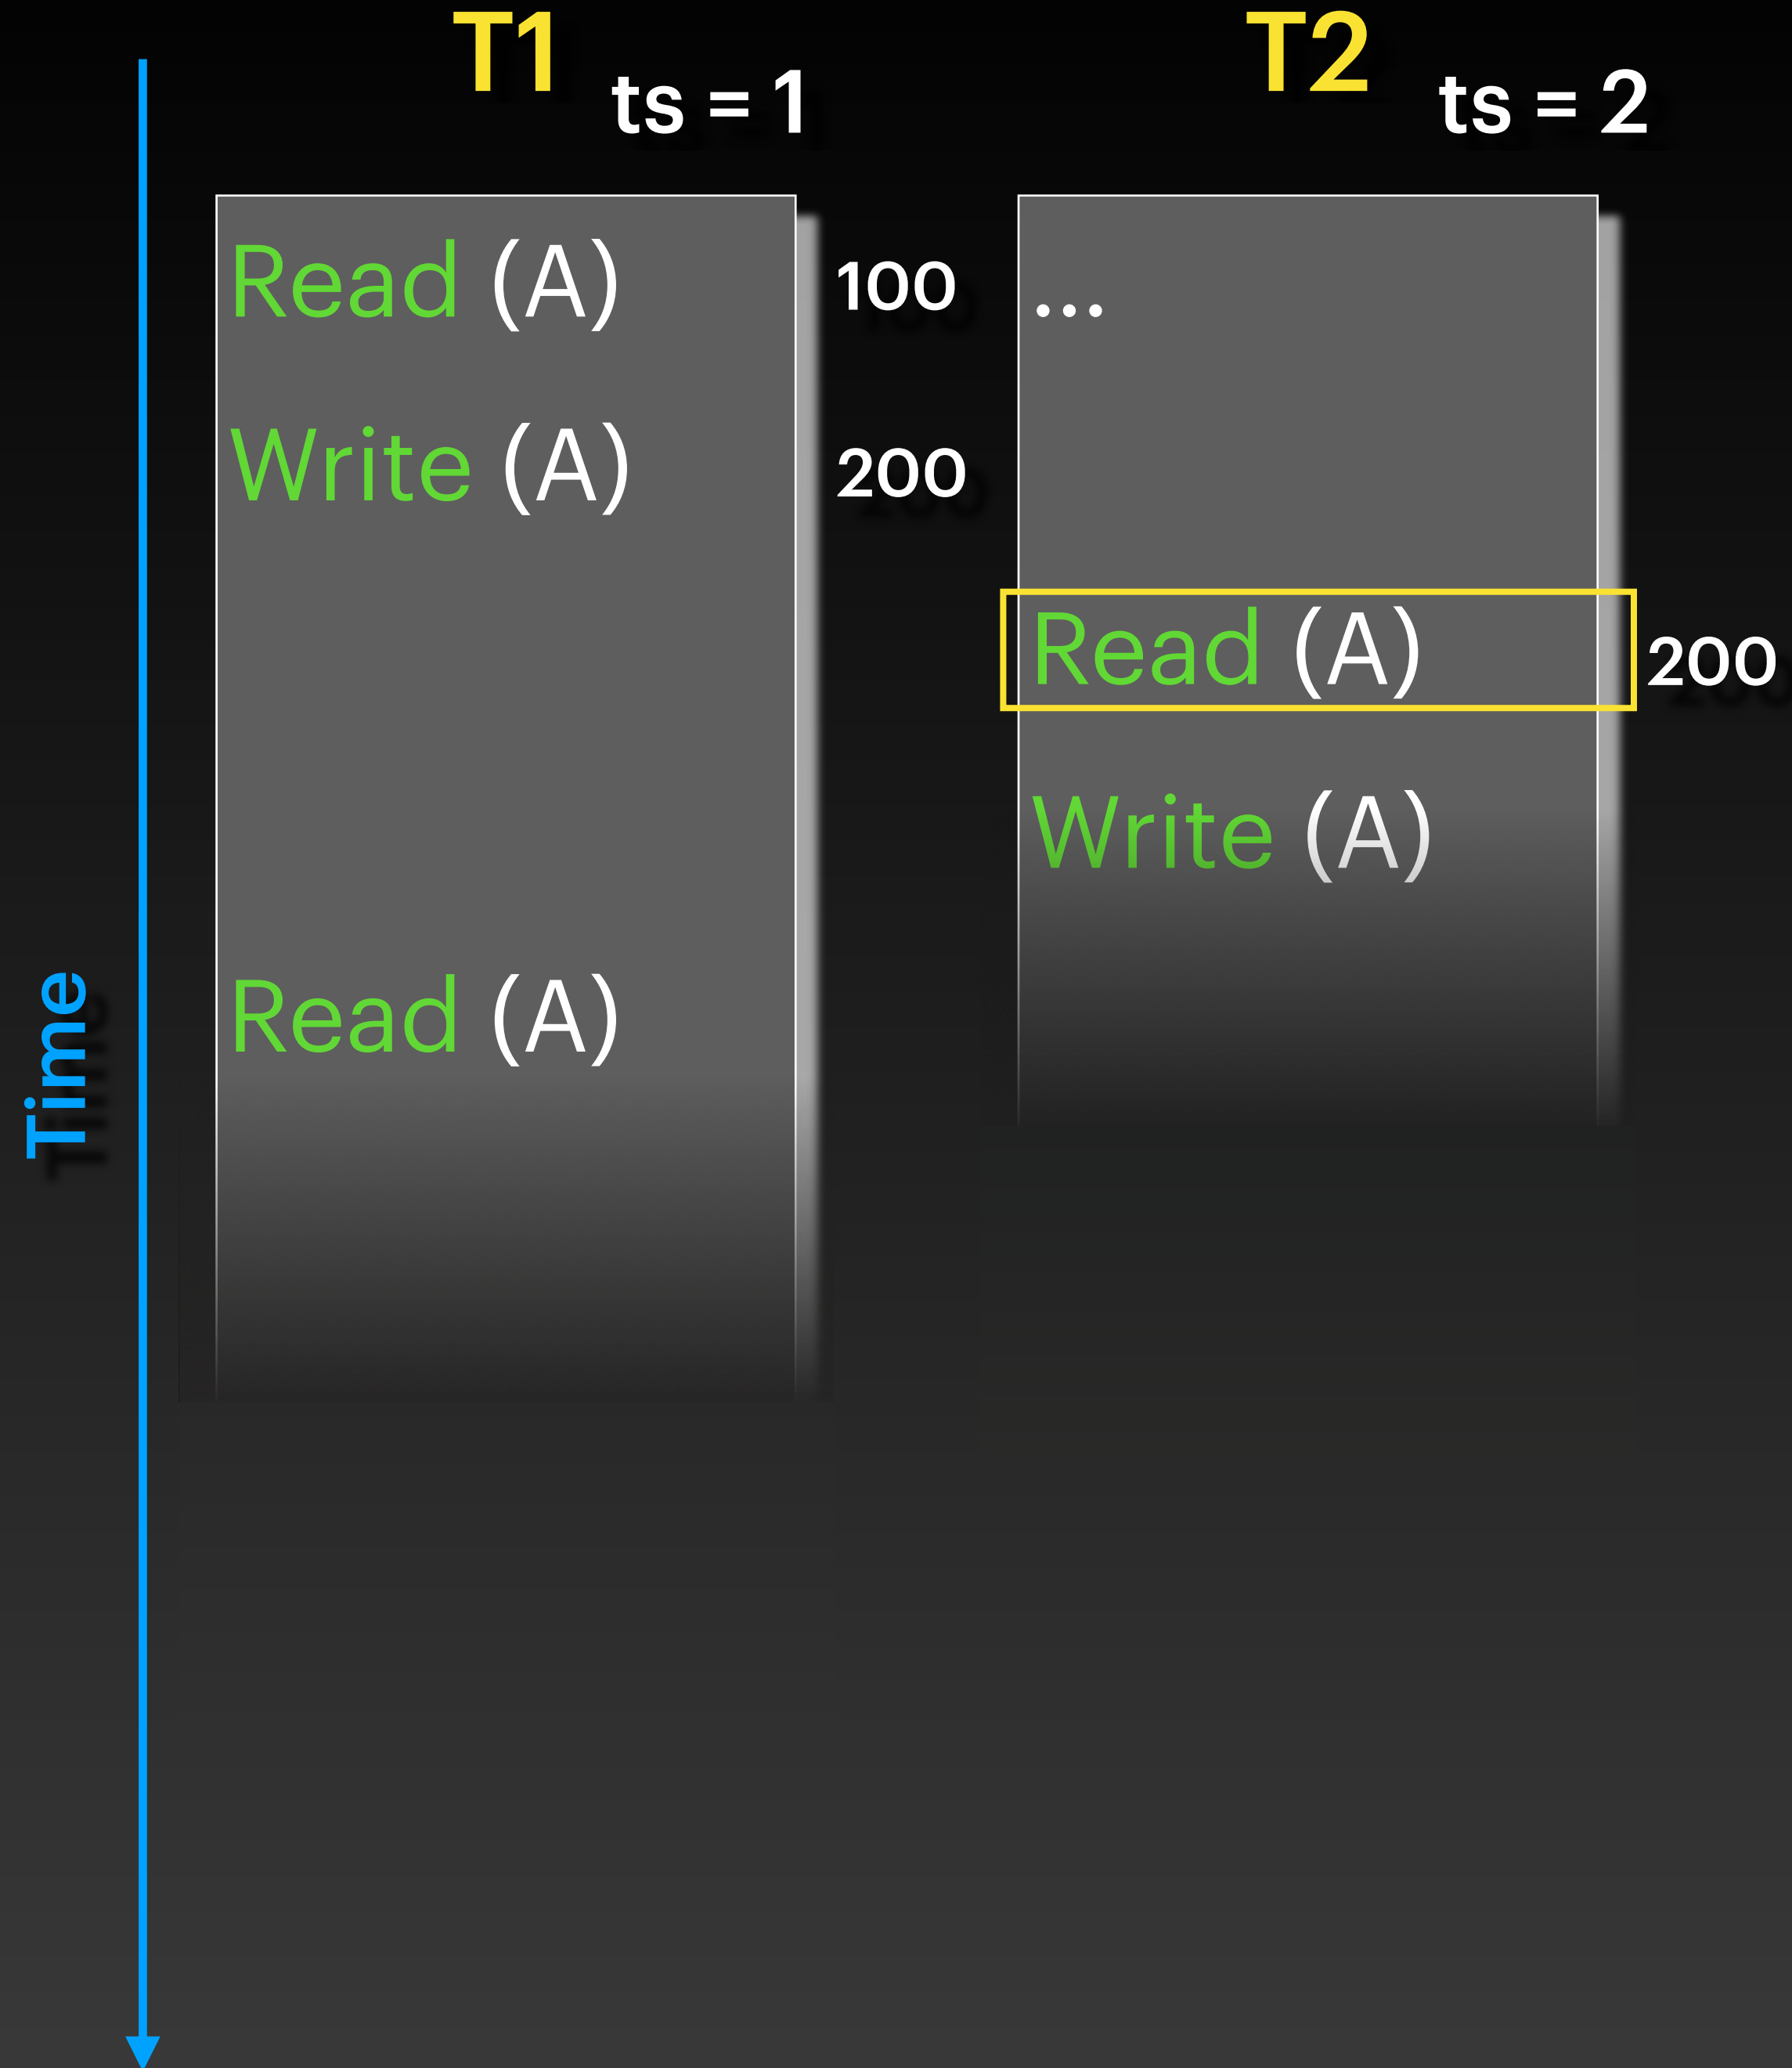


Version	Value	R-TS	W-TS
A_0	100	1	0

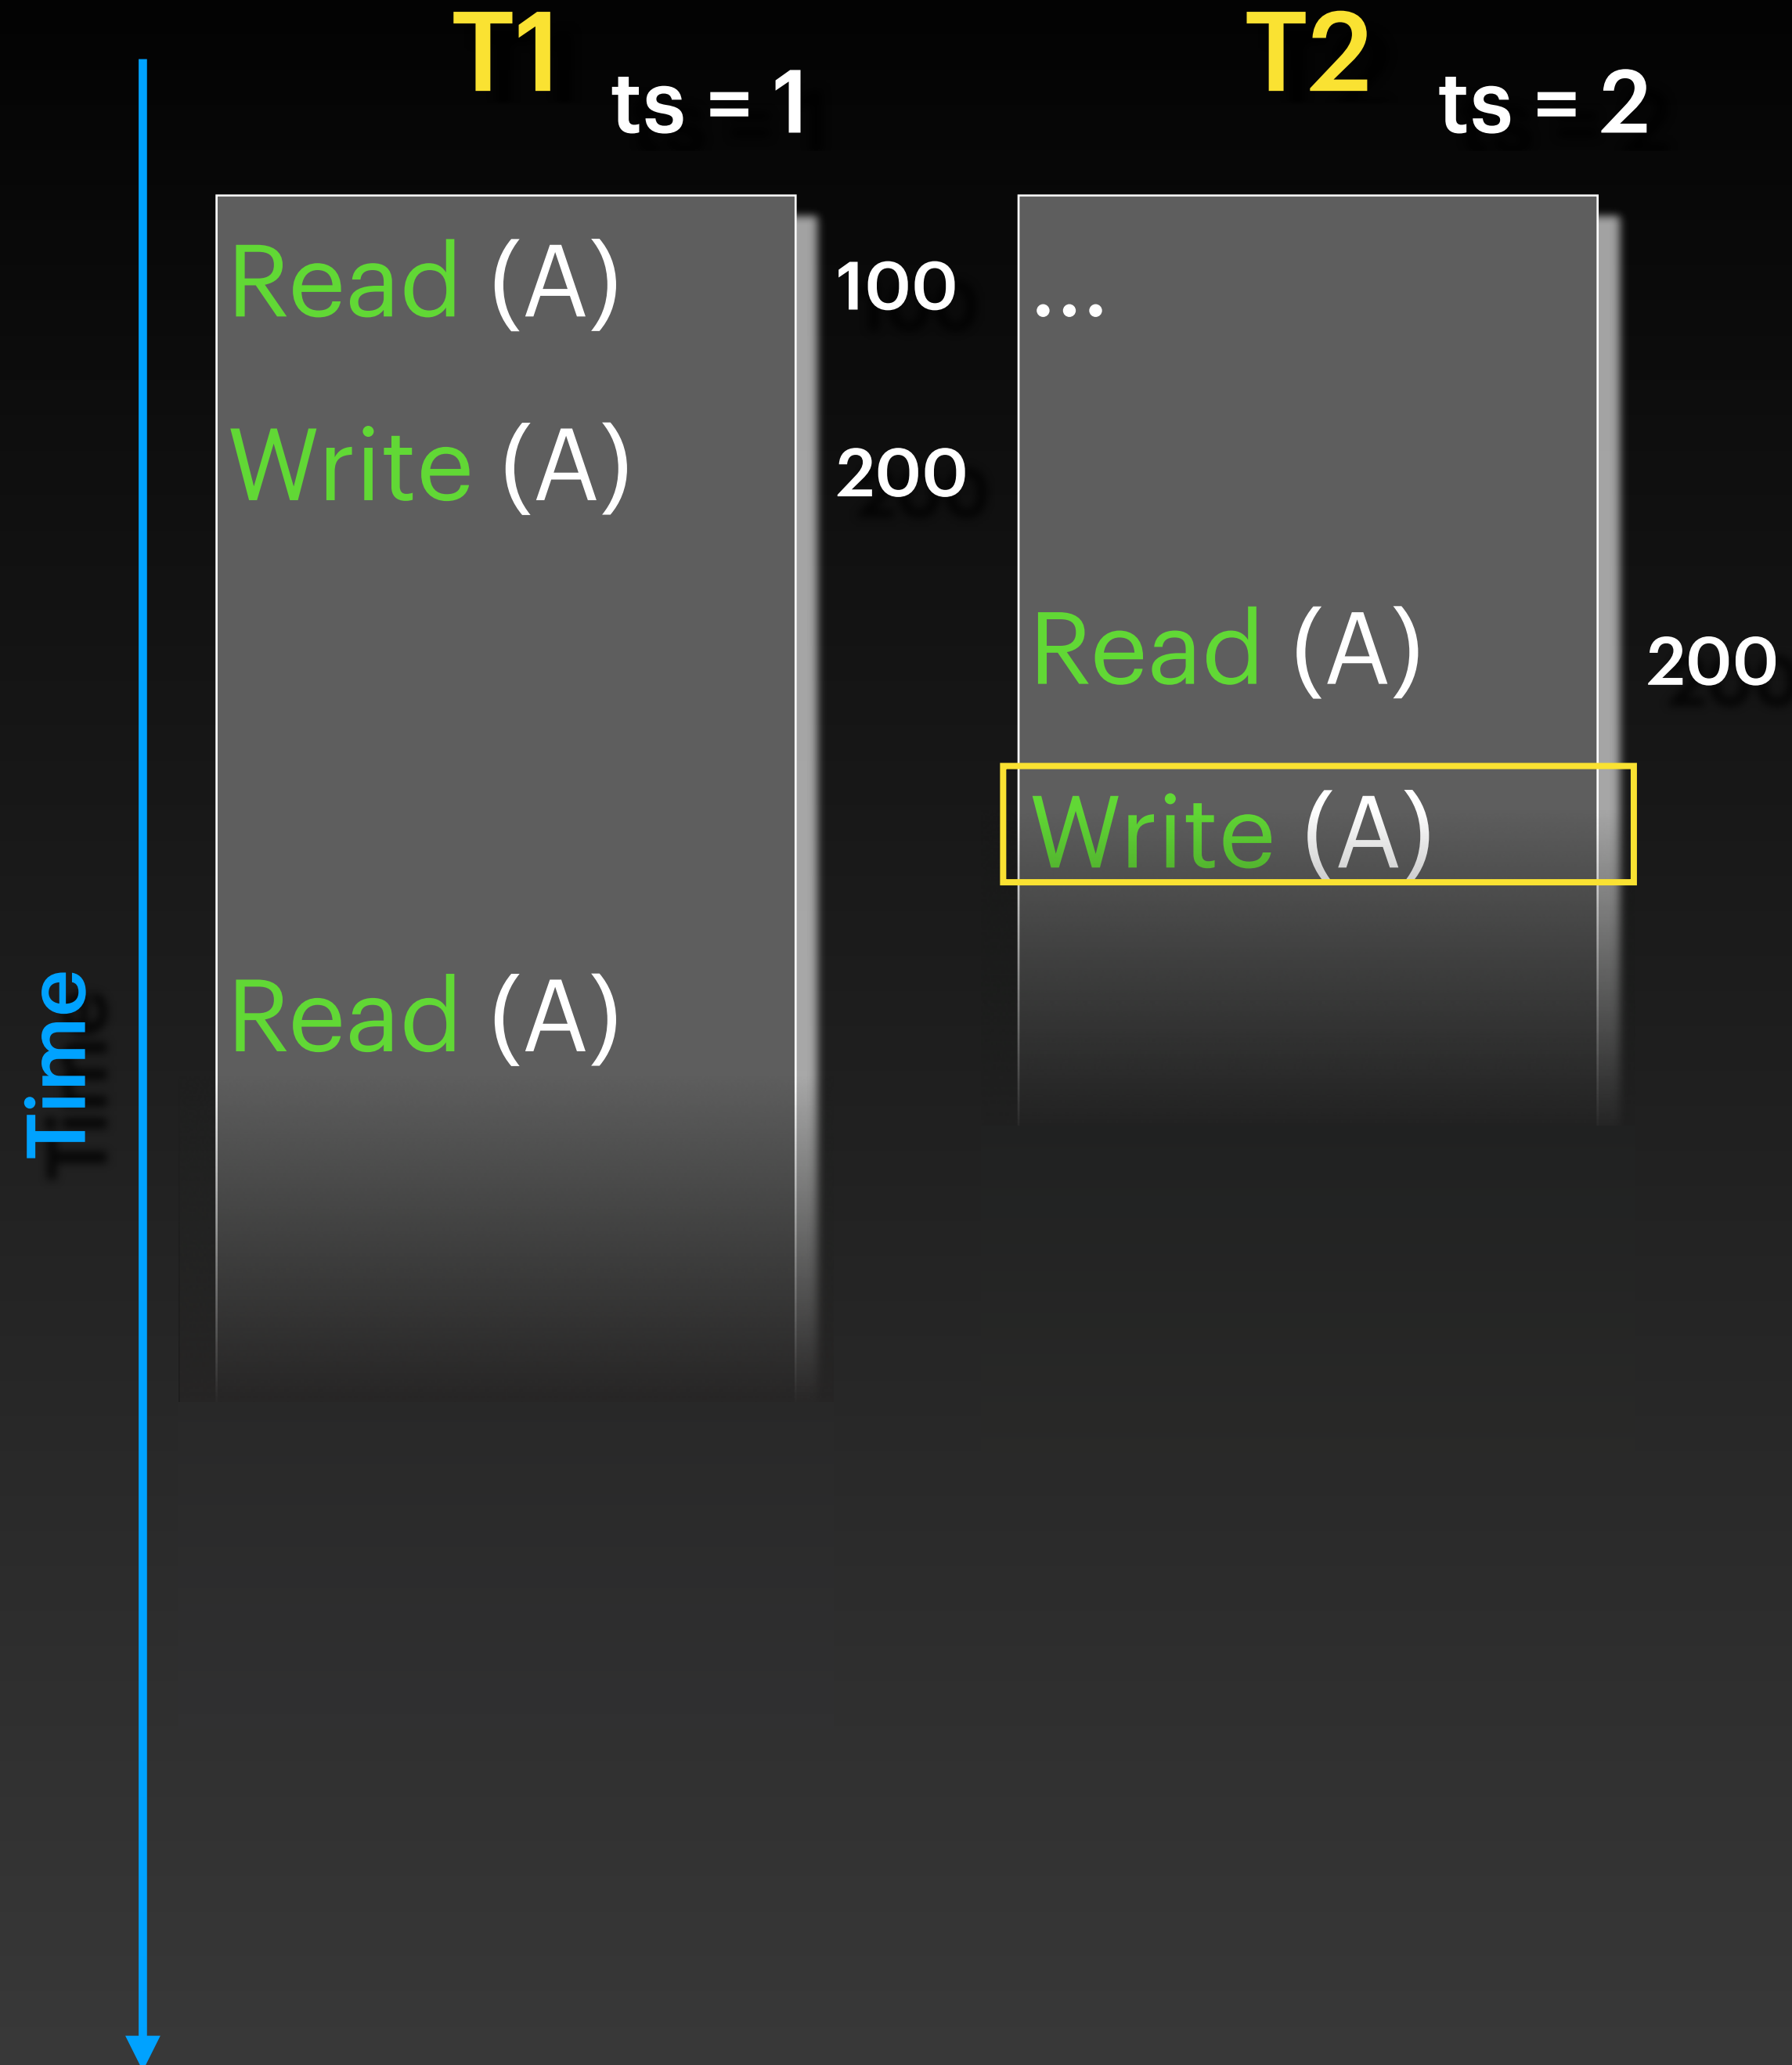
$R-TS \leq ts(T)$



Version	Value	R-TS	W-TS
A_0	100	1	0
A_1	200	1	1



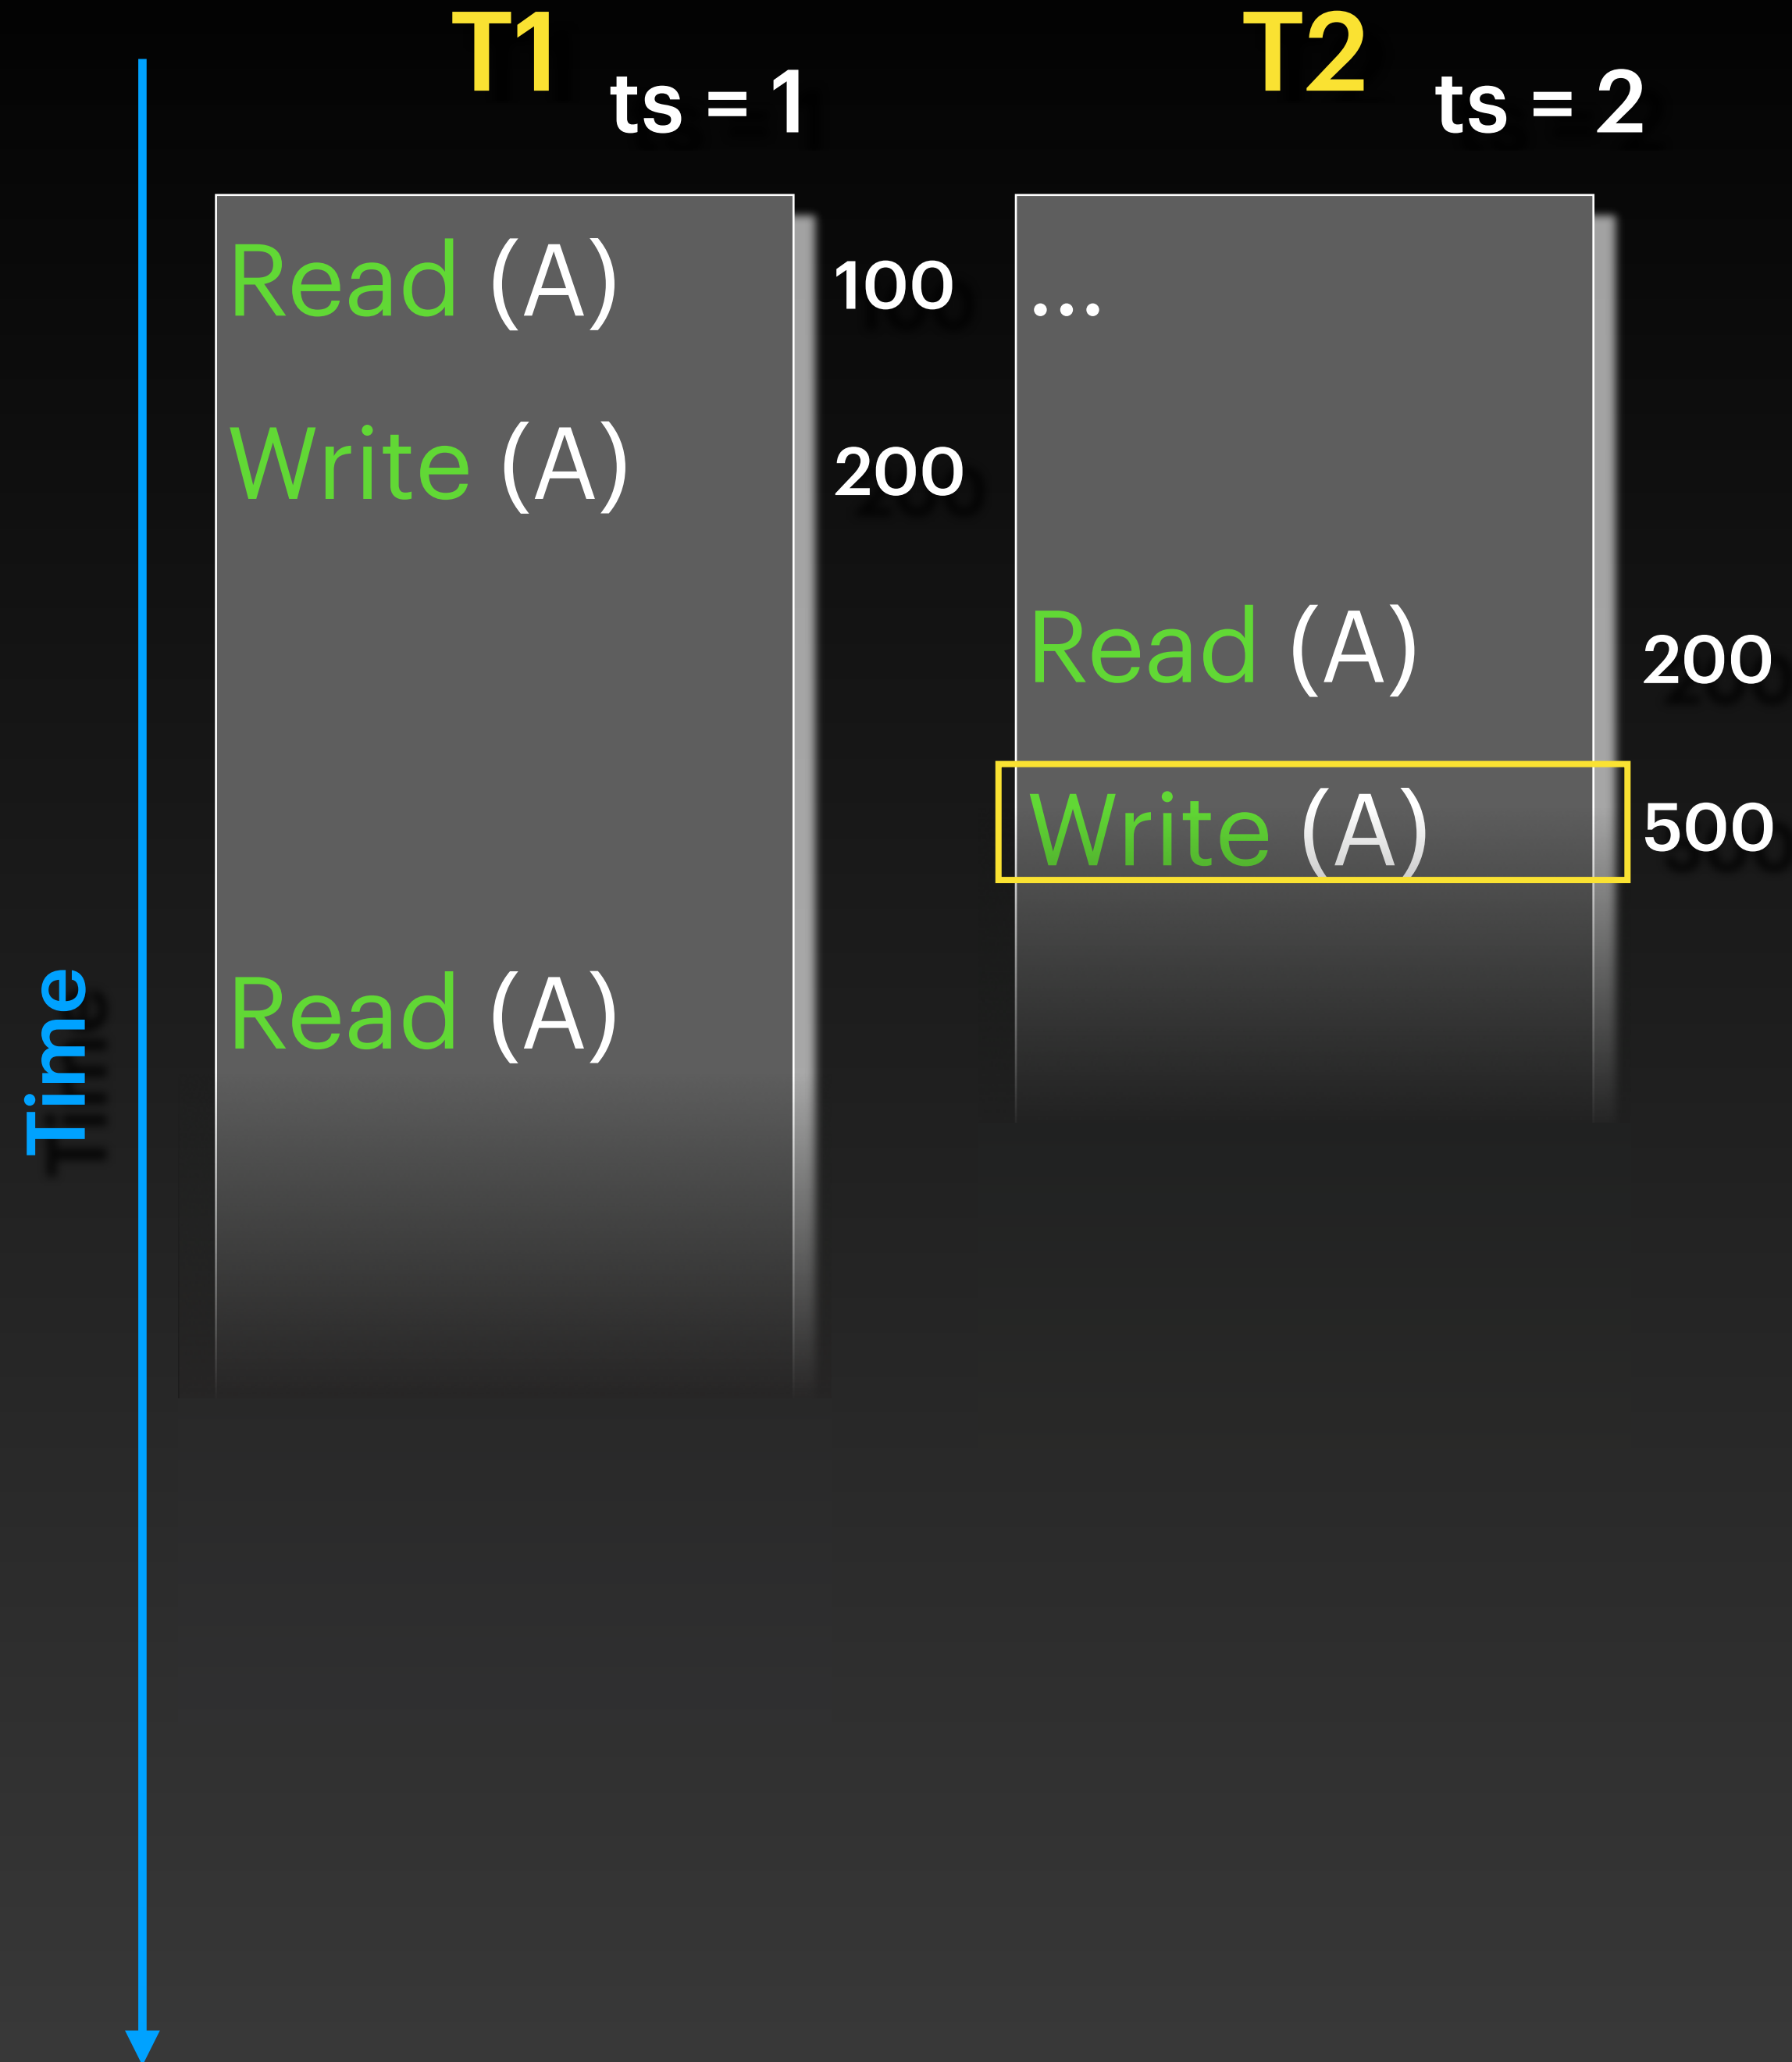
Version	Value	R-TS	W-TS
A_0	100	1	0
A_1	200	2	1



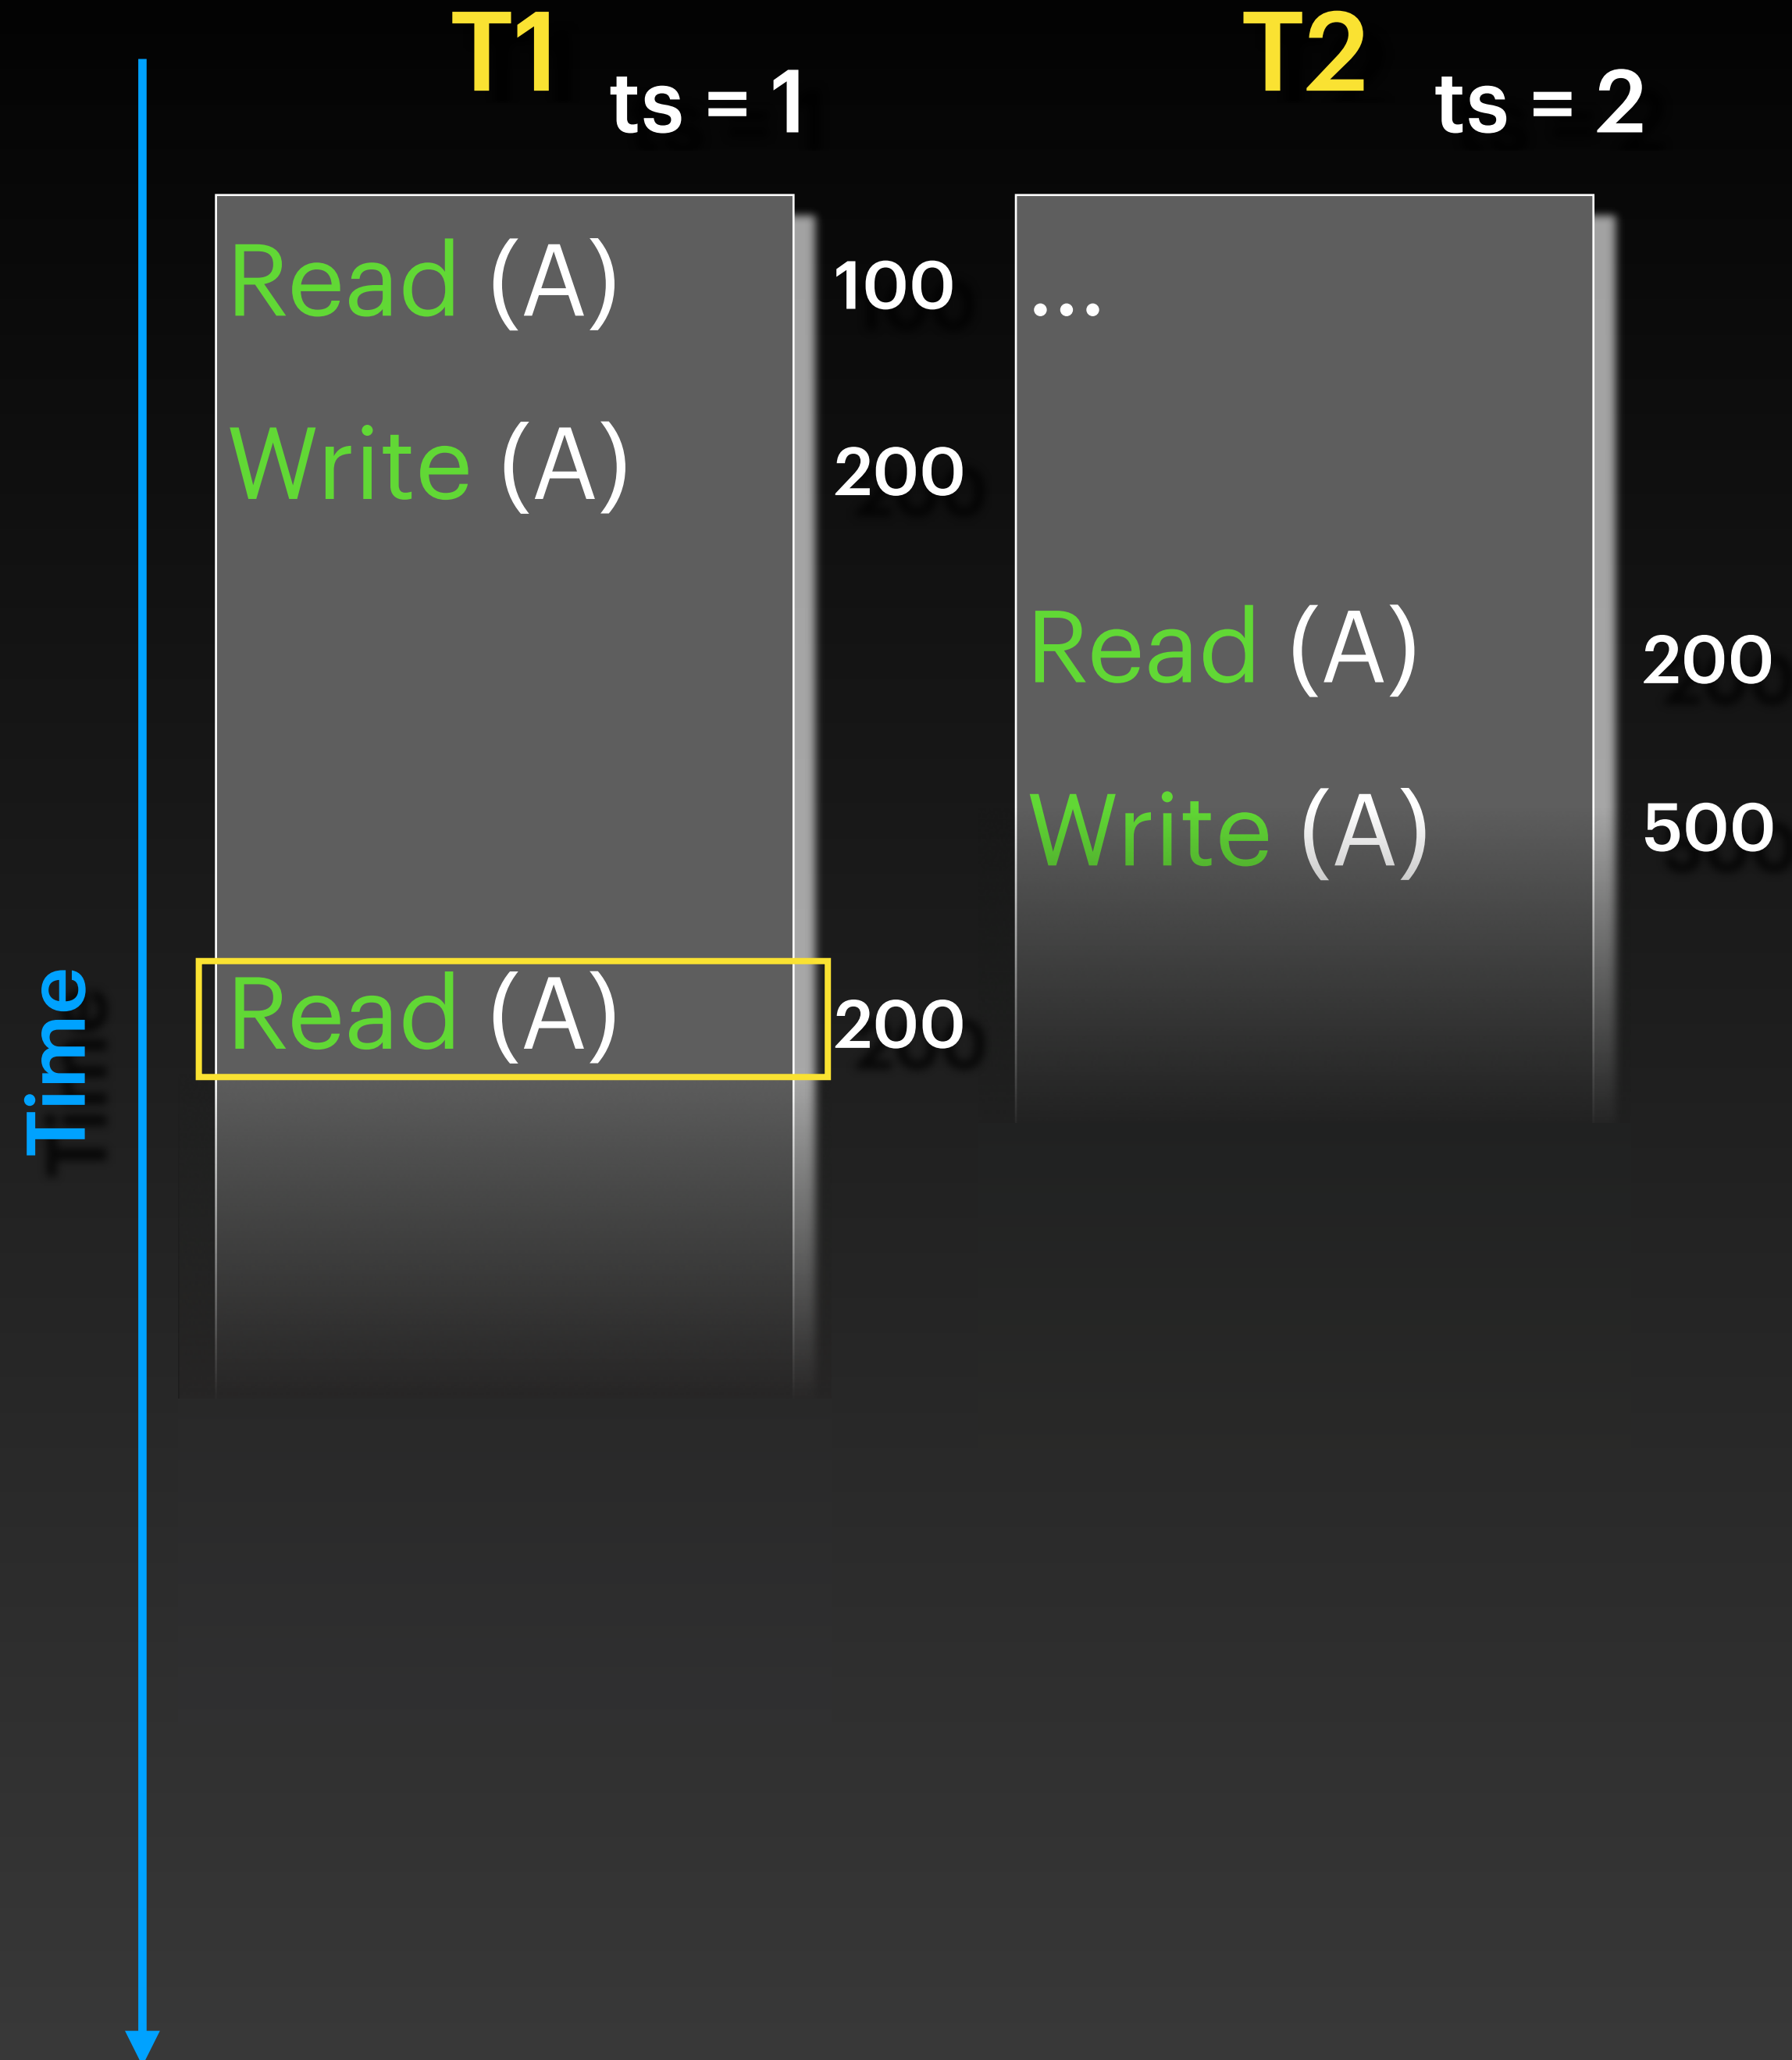
Version	Value	R-TS	W-TS
A_0	100	1	0
A_1	200	2	1

R-TS $\leq ts(T)$

Highest W-TS $\leq ts(T)$

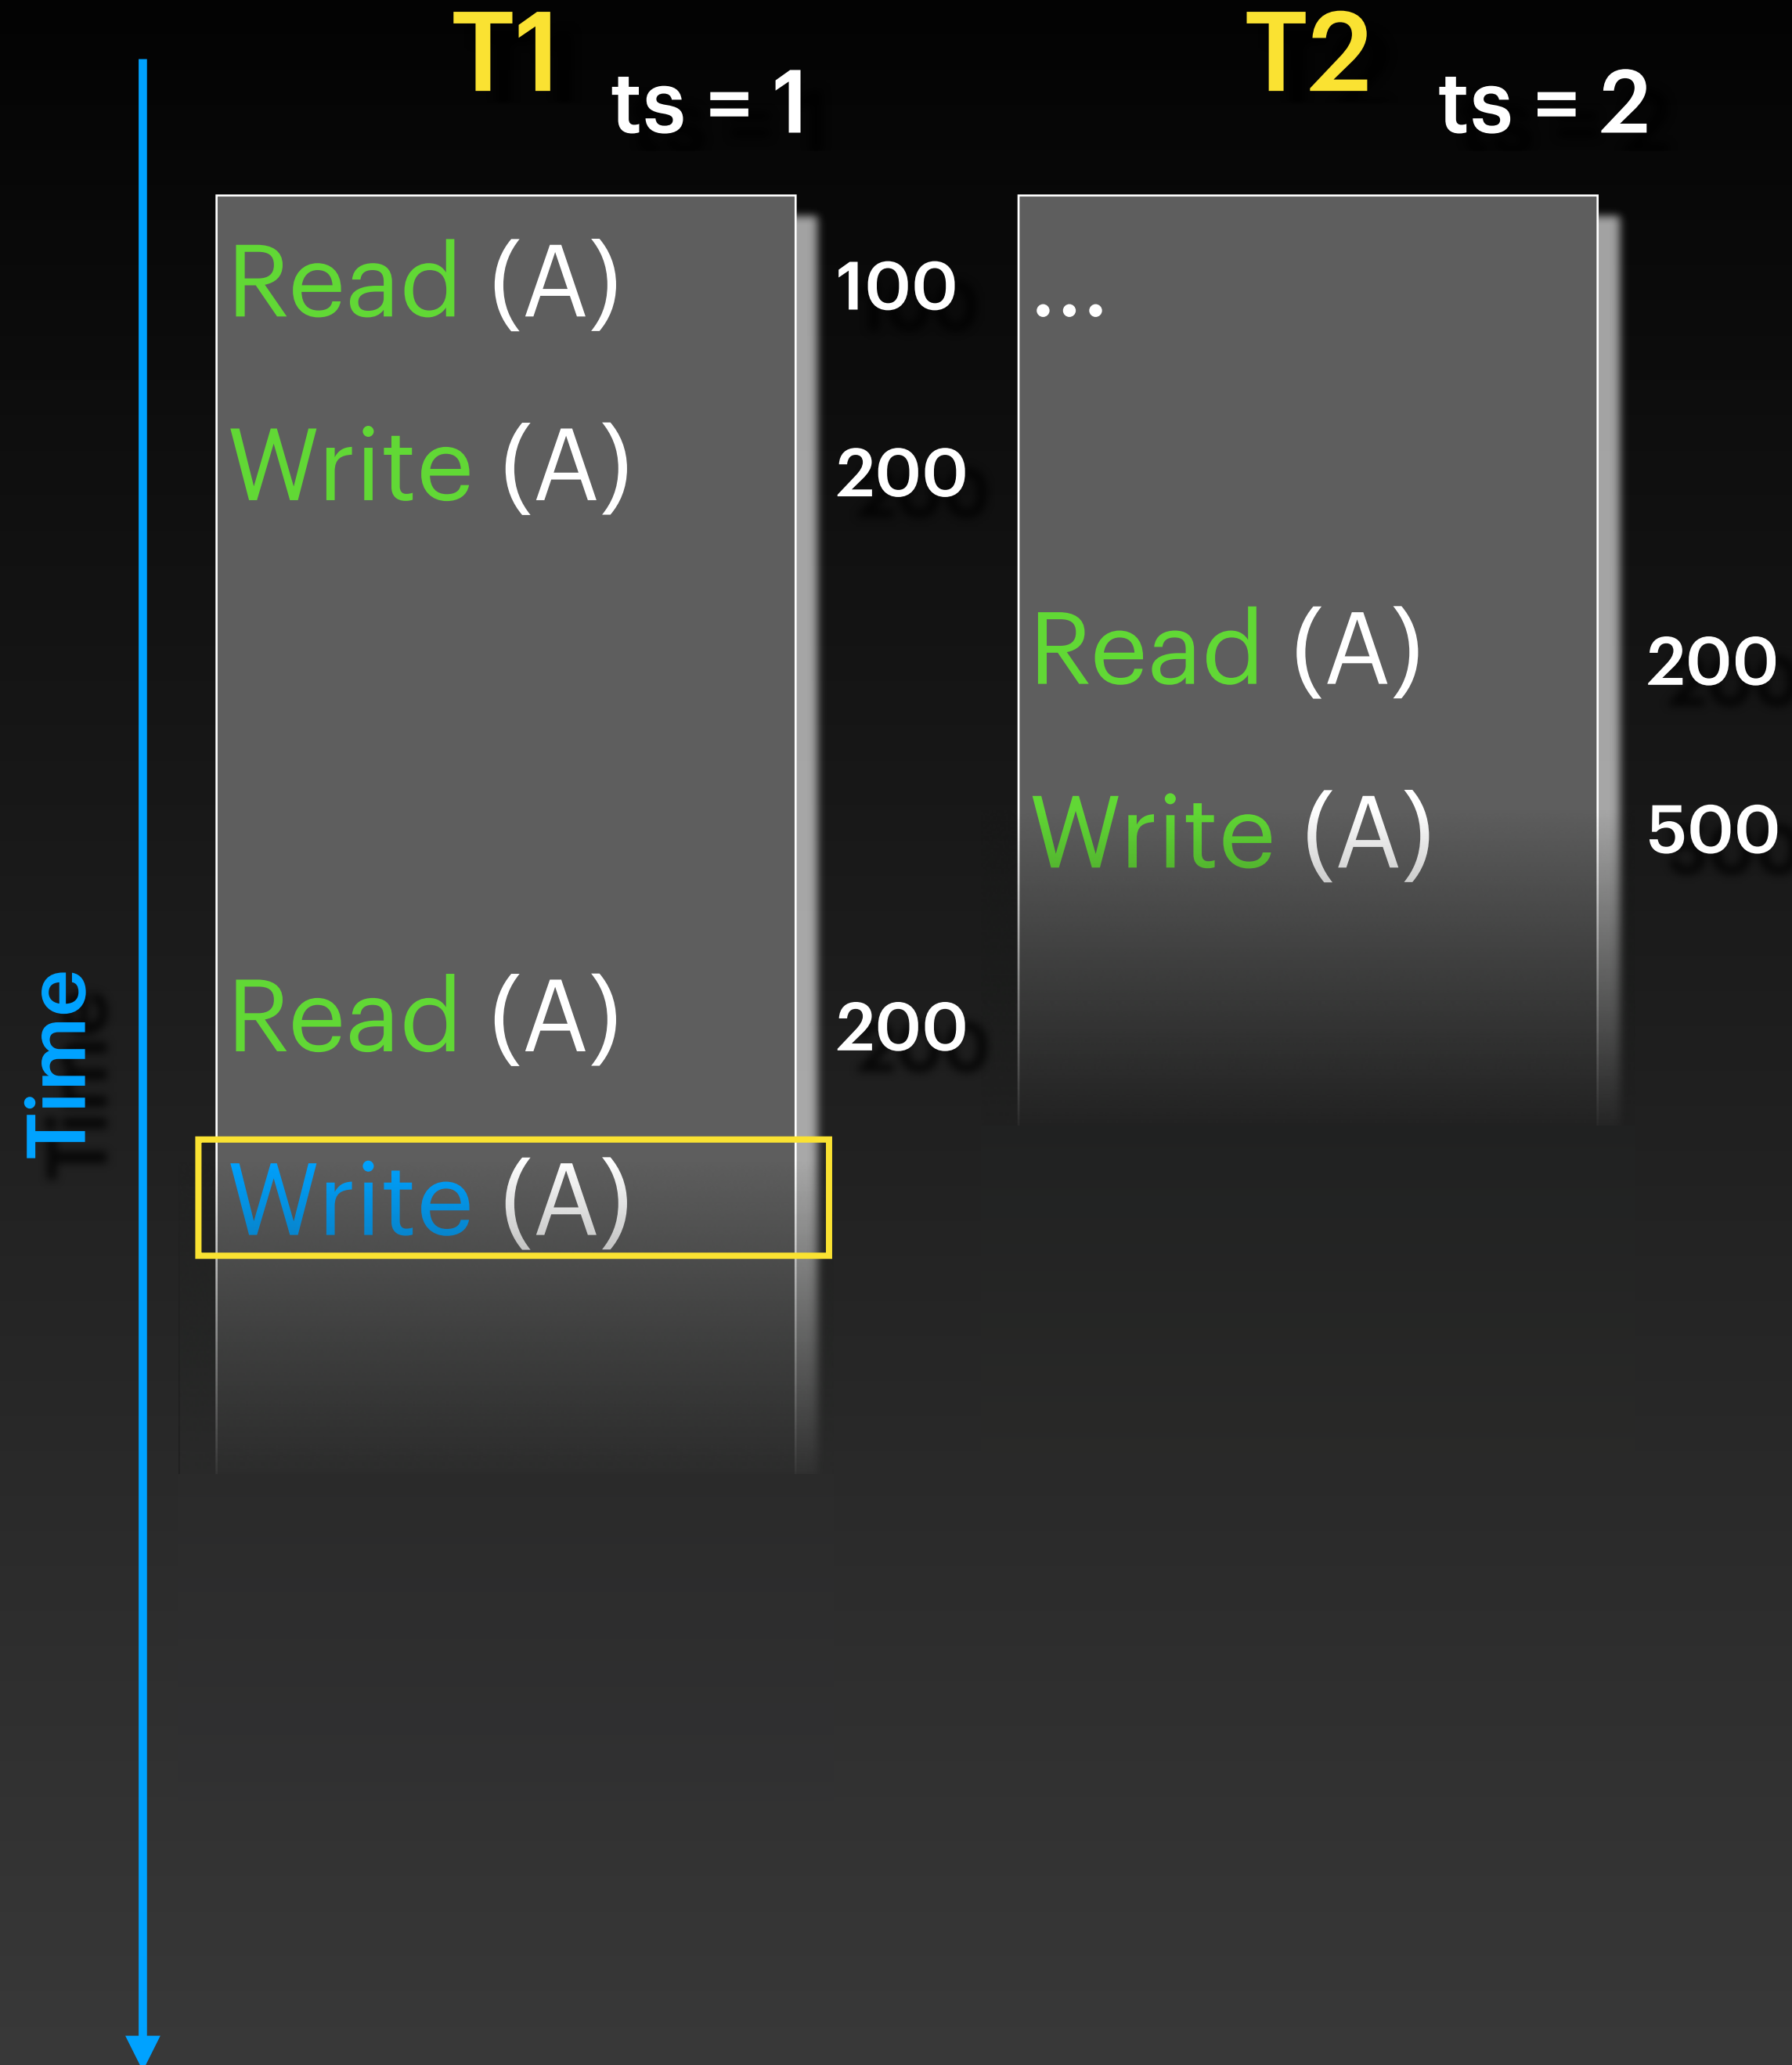


Version	Value	R-TS	W-TS
A_0	100	1	0
A_1	200	2	1
A_2	500	2	2



Version	Value	R-TS	W-TS
A_0	100	1	0
A_1	200	2	1
A_2	500	2	2


Highest W-TS
 $\leq ts(T)$



Version	Value	R-TS	W-TS
A_0	100	1	0
A_1	200	2	1
A_2	500	2	2

Highest W-TS
 $\leq ts(T)$

$R-TS > ts(T)$



Version	Value	R-TS	W-TS
A_0	100	1	0
A_1	200	2	1
A_2	500	2	2

Highest W-TS
 $\leq ts(T)$

$R-TS > ts(T)$