


# DATABASE STORAGE PART 2

*Amr Elhelw's*  
**TECH**  
**VAULT**




# Database Storage


Database Files



Pages

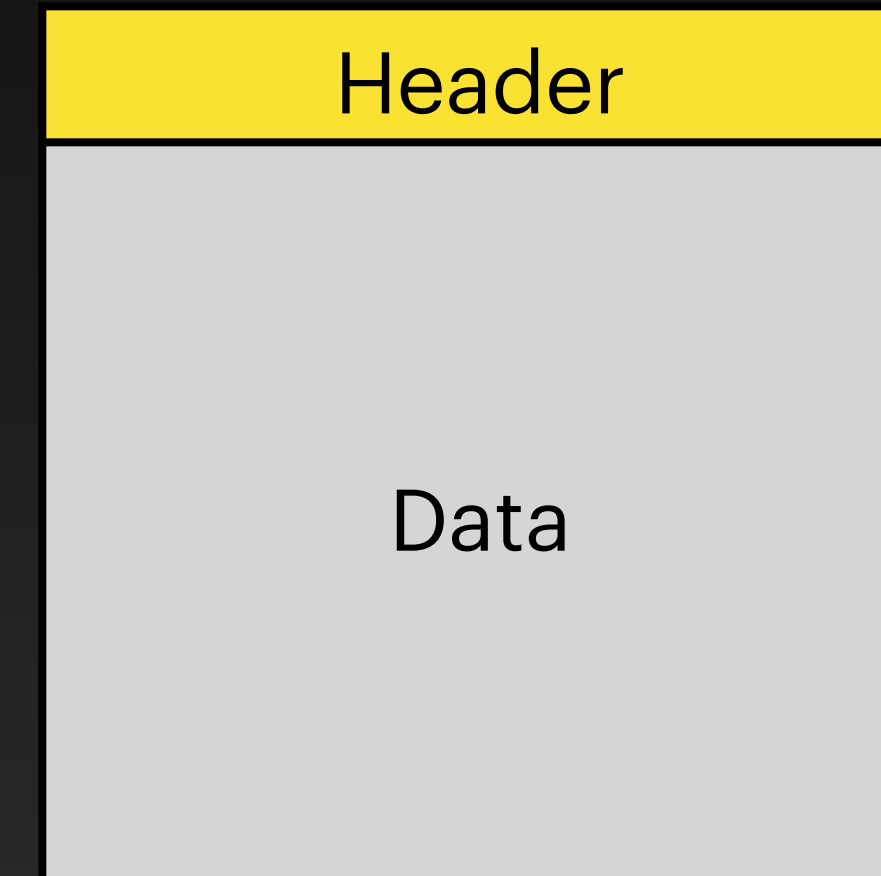


Tuples



# Page Header

- Exists in every page
- Metadata
  - Page size
  - Compression/encoding info
  - DBMS version
  - ...



# Naive Approach


- Maintain # of tuples in page
- Append new tuples at the end
  - Use the # of tuples to determine “the end” (assuming equal-size tuples)
- **Challenges**
  - Deletion
  - Variable-sized tuples

Start  
of data →

Num tuples = 2	
	Tuple #1
	Tuple #2

# Slotted Pages

- Header maintains:
  - # slots
  - Offset of last used slot
- Each slot points to corresponding tuple





# Slotted Pages

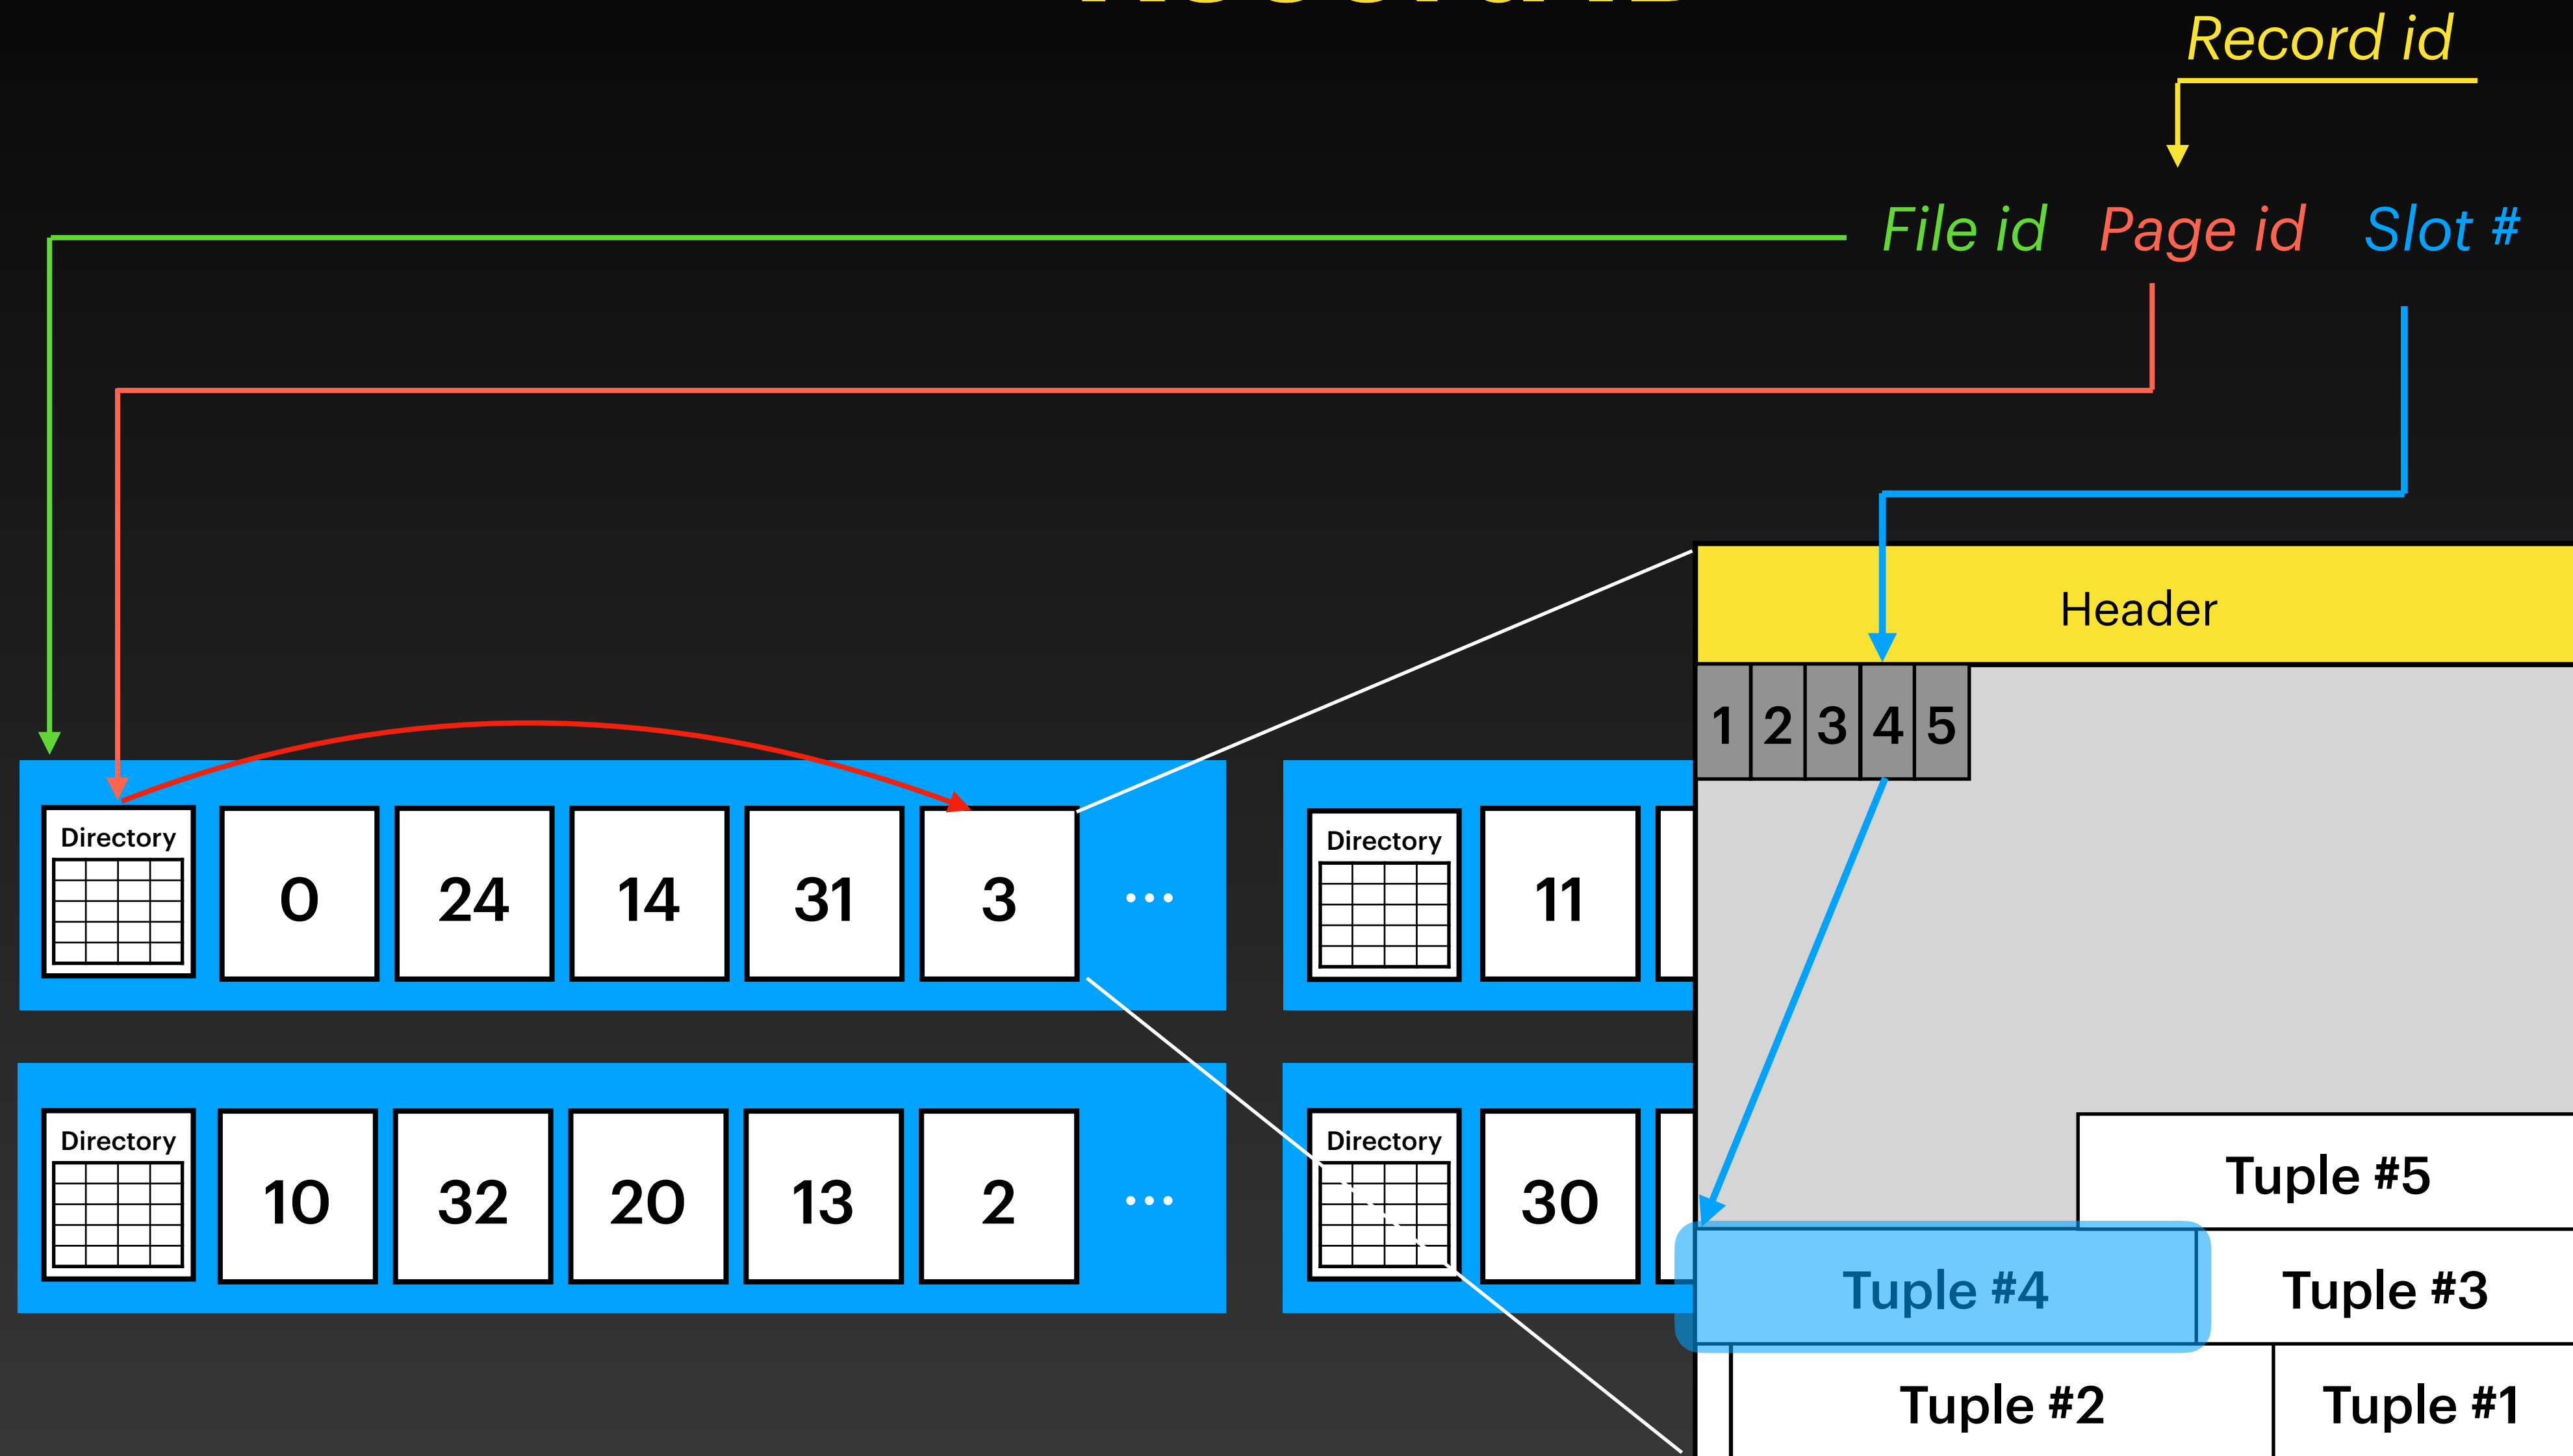
- Header maintains:
  - # slots
  - Offset of last used slot
- Each slot points to corresponding tuple
- Tuples are added starting from the end of the page

Header						
1	2	3	4	5	6	
			Tuple #6		Tuple #5	
Tuple #4				Tuple #3		
	Tuple #2				Tuple #1	

# Record ID

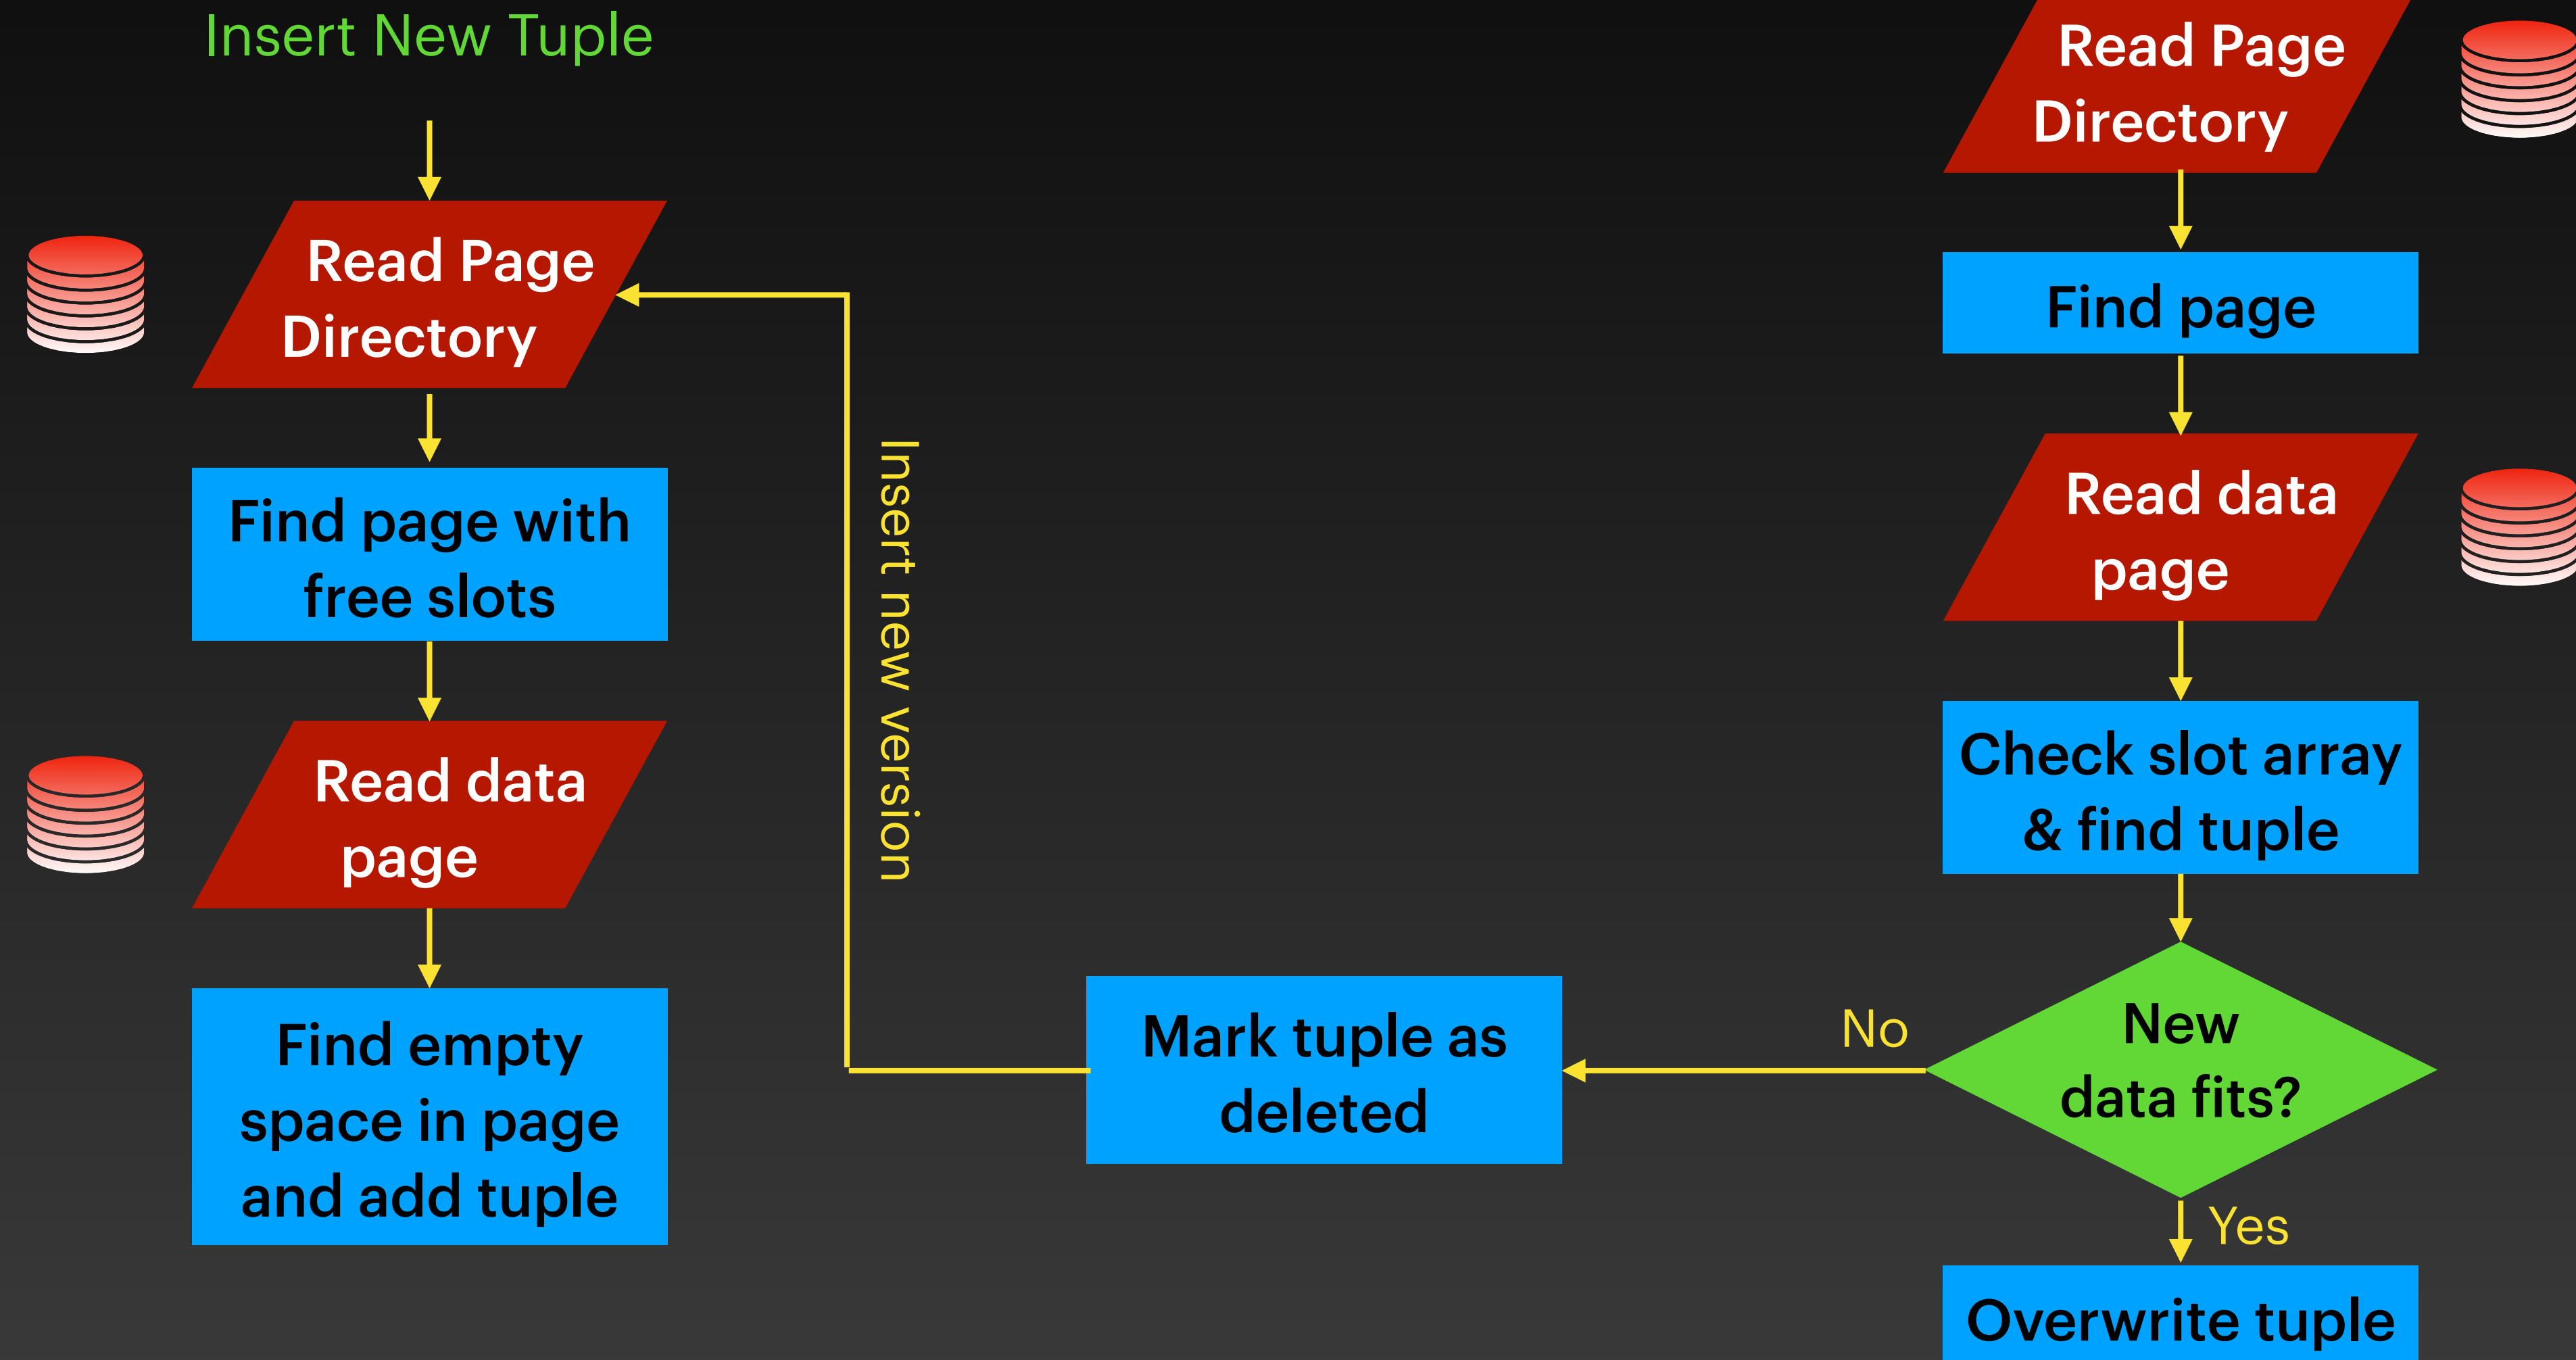
- Unique for every tuple
- Represents physical location of tuple
  - Some encoding of (file id, page id, slot #)
  - Internal “index”
- Usually not stored with the tuple

# Record ID





# Slotted Pages

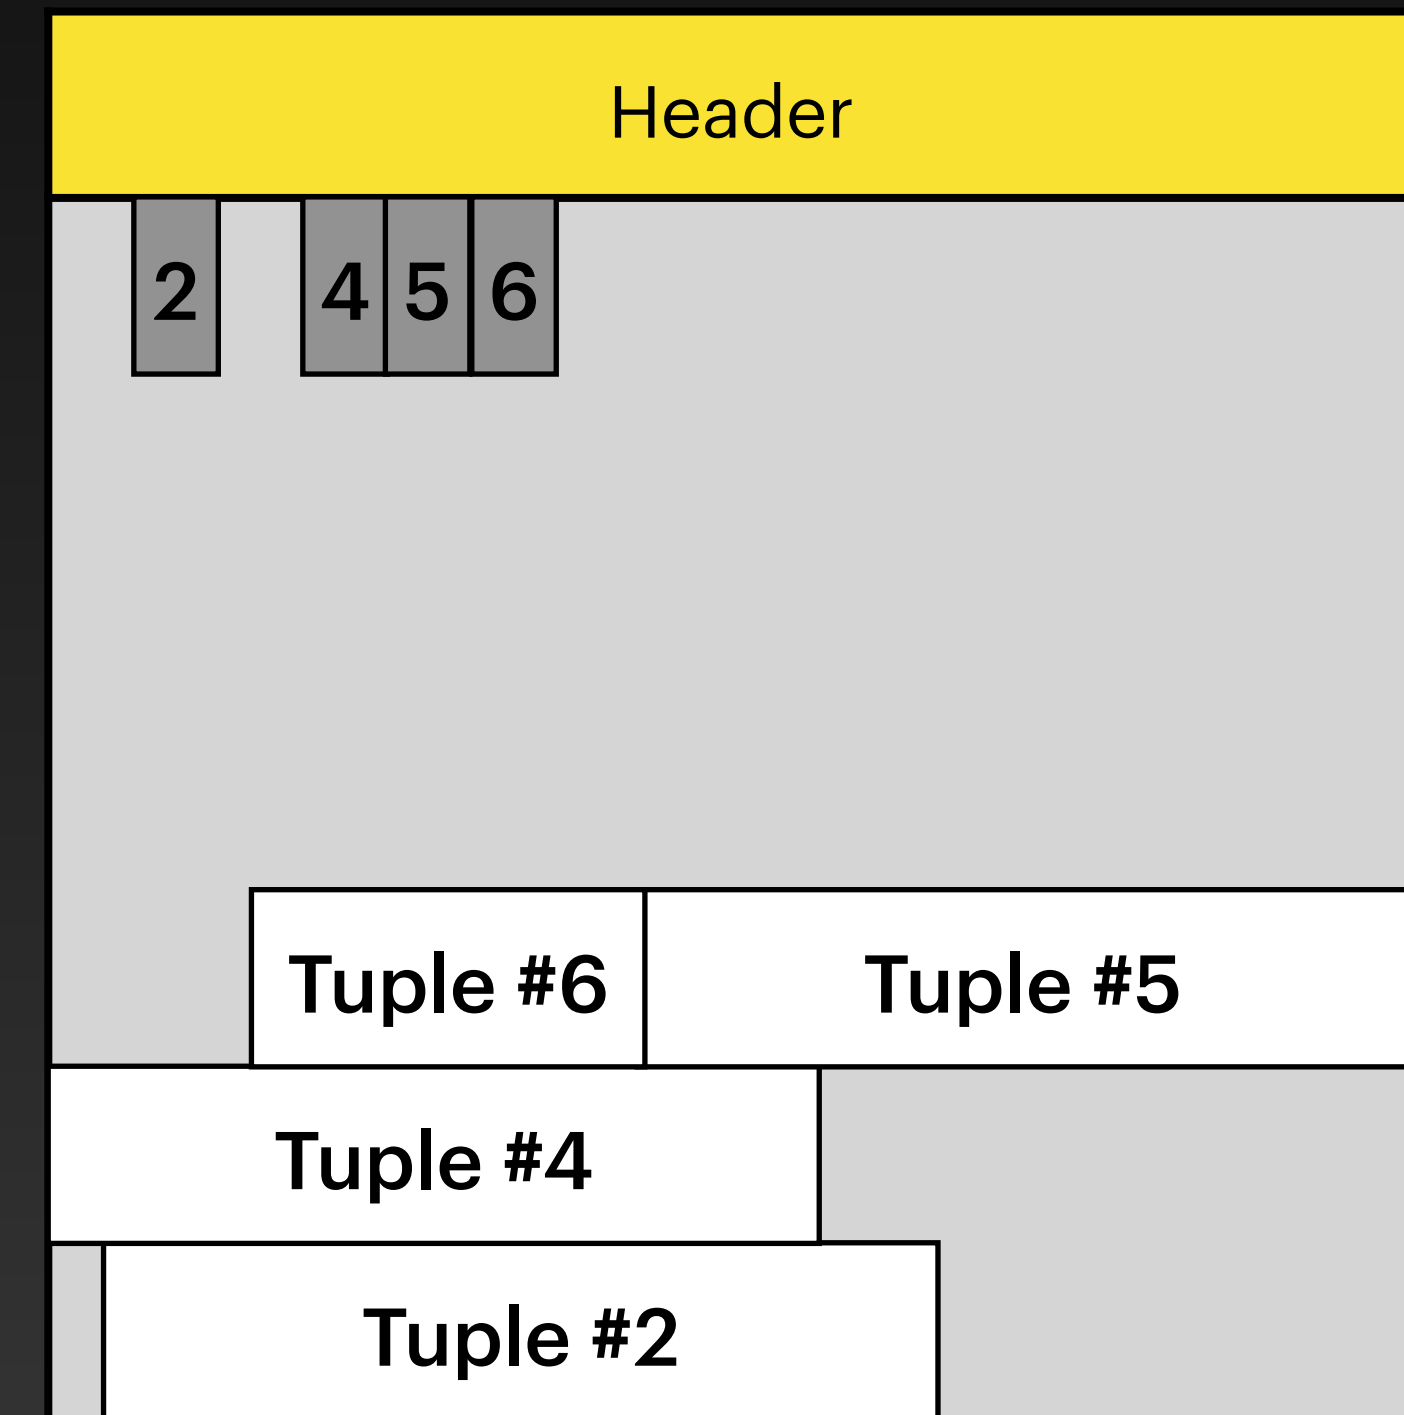


# Slotted Pages - Challenges

- **Challenge #1 - Too much Disk I/O**
  - 2 pages for an insert
  - 2-4 pages for an update
- **Challenge #2 - Random Disk I/O**
  - E.g. updating multiple tuples in separate pages

# Slotted Pages - Challenges

- **Challenge #3 - Fragmentation**
  - Unused tuple and slot spaces



# Log-Structured Storage

- Also known as Log-structure merge trees (LSM trees)
- Record “changes” rather than storing the tuples in pages
  - Each entry: **SET** or **DELETE** operation on a tuple
  - Must include the **tuple id**



Data Page (in memory)

**SET** #31 (val=10)

**SET** #32 (val=20)

**DEL** #30

**SET** #31 (val=15)

**SET** #33 (val=42)


**SET** #31 (val=25)

# Log-Structured Storage



Memory

SET	...
DEL	...
SET	...
SET	...
DEL	...
SET	...




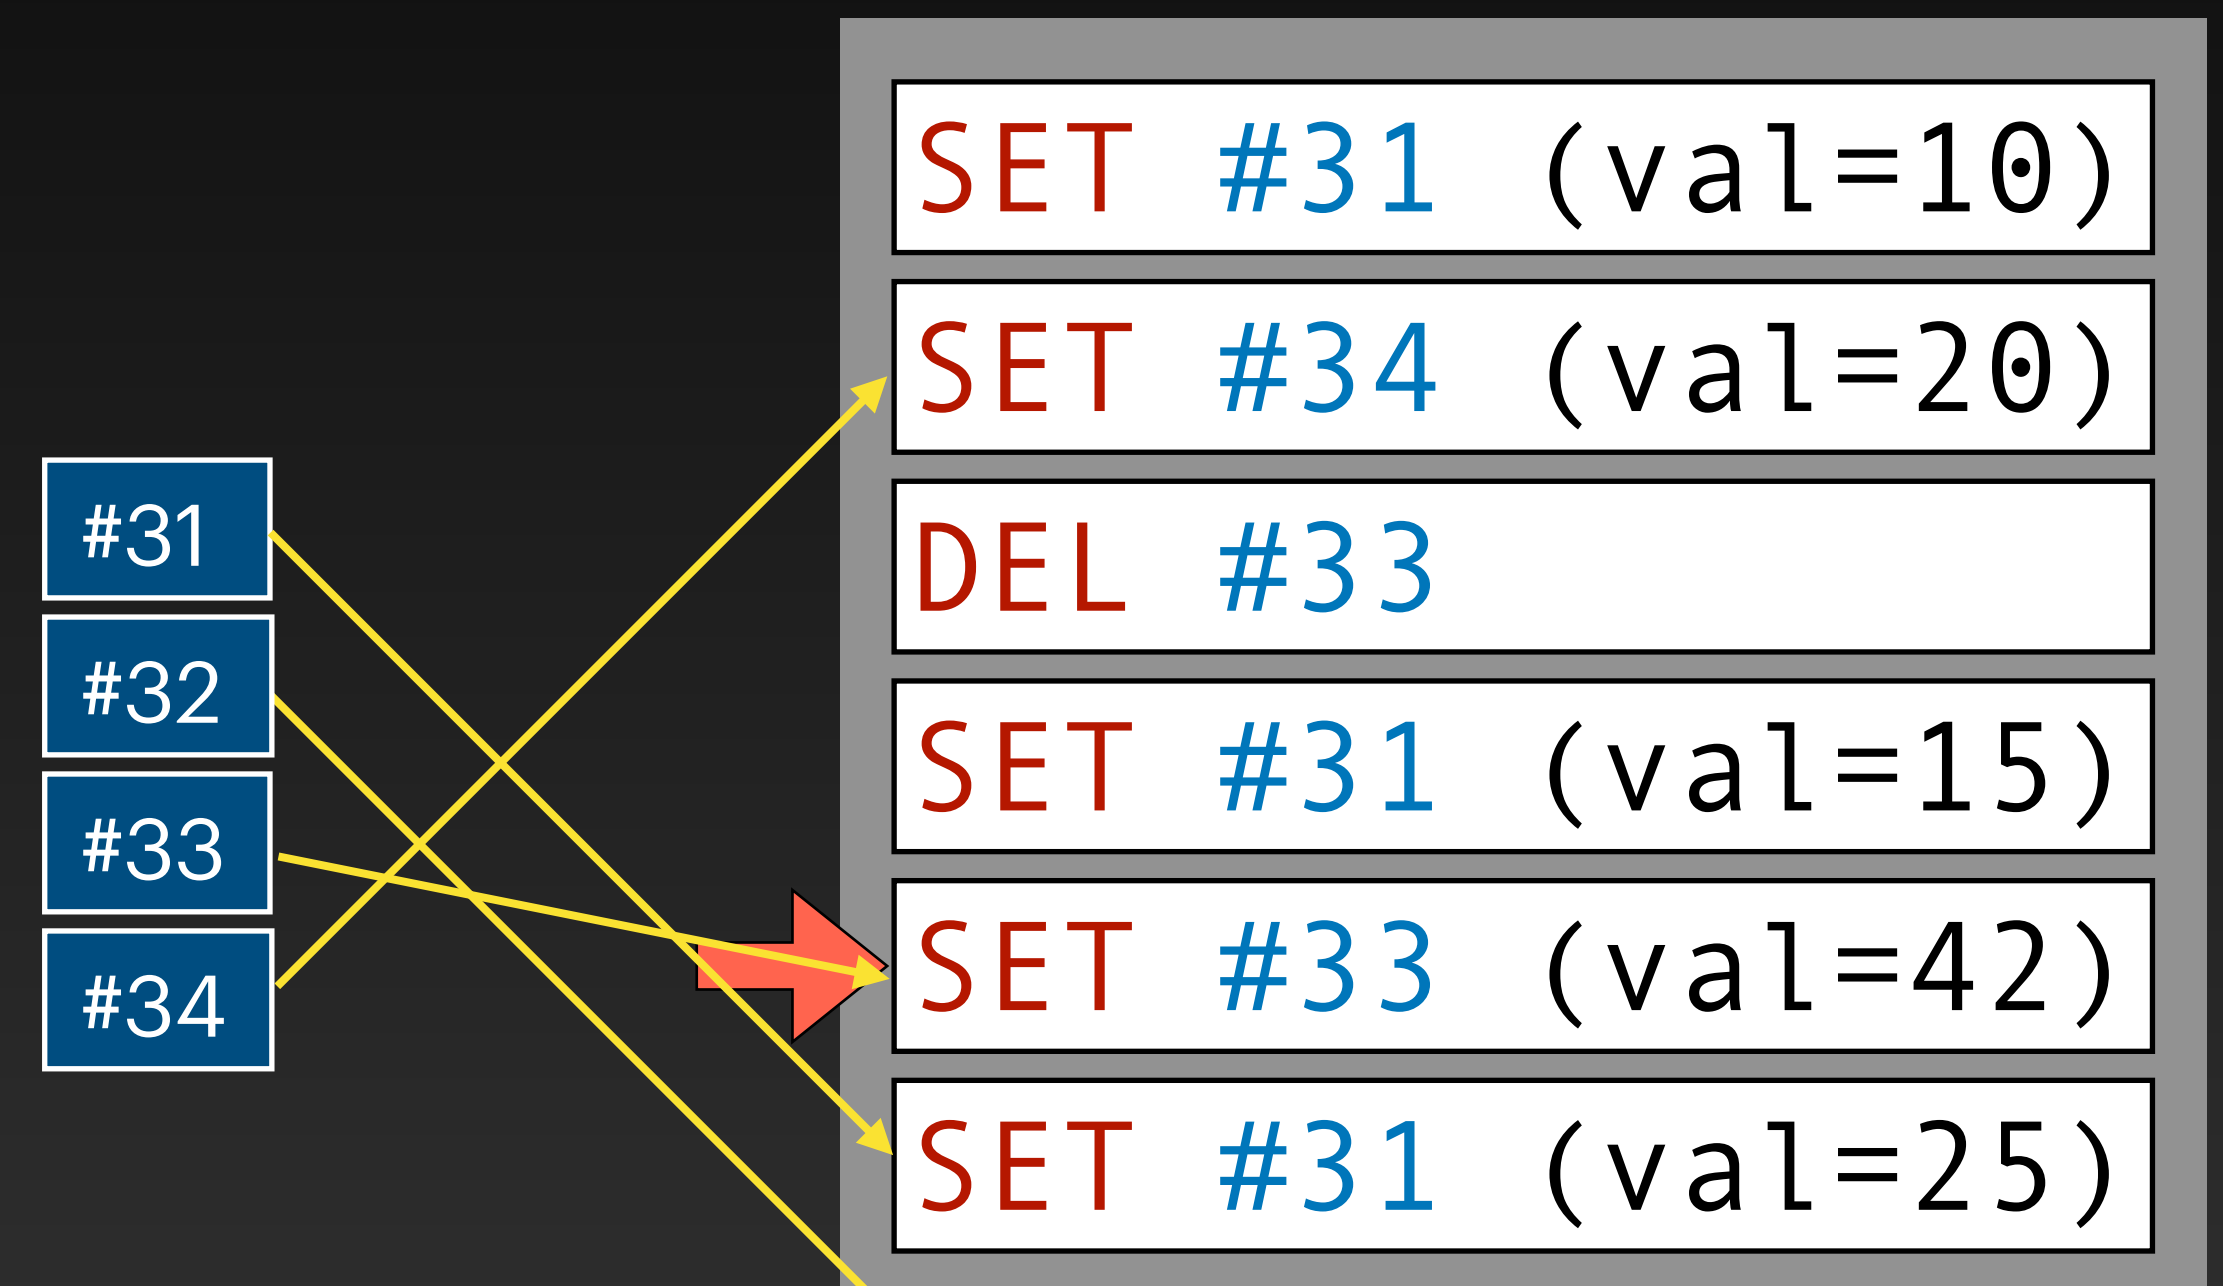
Disk



# Log-Structured Storage

Read #33

- Read tuple with given id
  - Find newest entry with that id
  - Scan log from end to beginning
- Scanning is inefficient
  - Use index





# Compaction

SET #31 (val=10)

SET #34 (val=20)

DEL #33

SET #31 (val=15)

SET #33 (val=42)

SET #31 (val=25)

SET #33 (val=11)


DEL #34

SET #35 (val=30)

SET #32 (val=32)

DEL #35

SET #32 (val=16)



SET #31 (val=25)

SET #33 (val=11)

DEL #34

DEL #35

SET #32 (val=16)


# Index-Organized Tables (IOT)

- Store the actual data inside an index (primary key)
- Faster reads
- Slower writes
- Reduced storage




# Database Storage

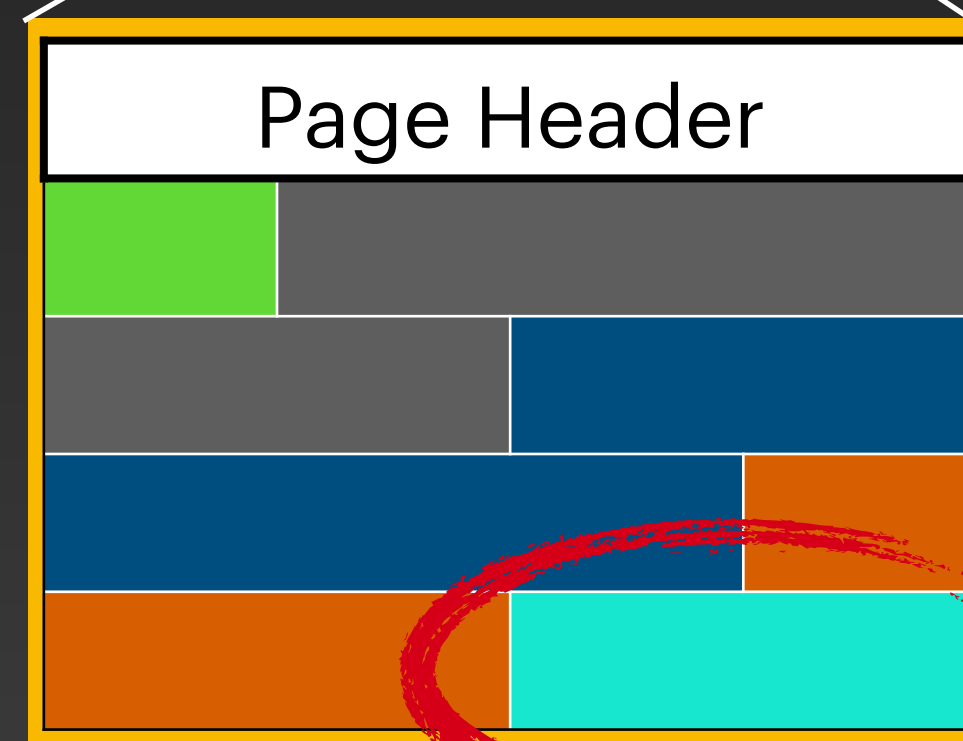
Database Files



Pages



Tuples




# Tuple Storage

- A sequence of bytes
- Header + Data



- Interpreted by the DBMS using schema information


# Tuple Storage



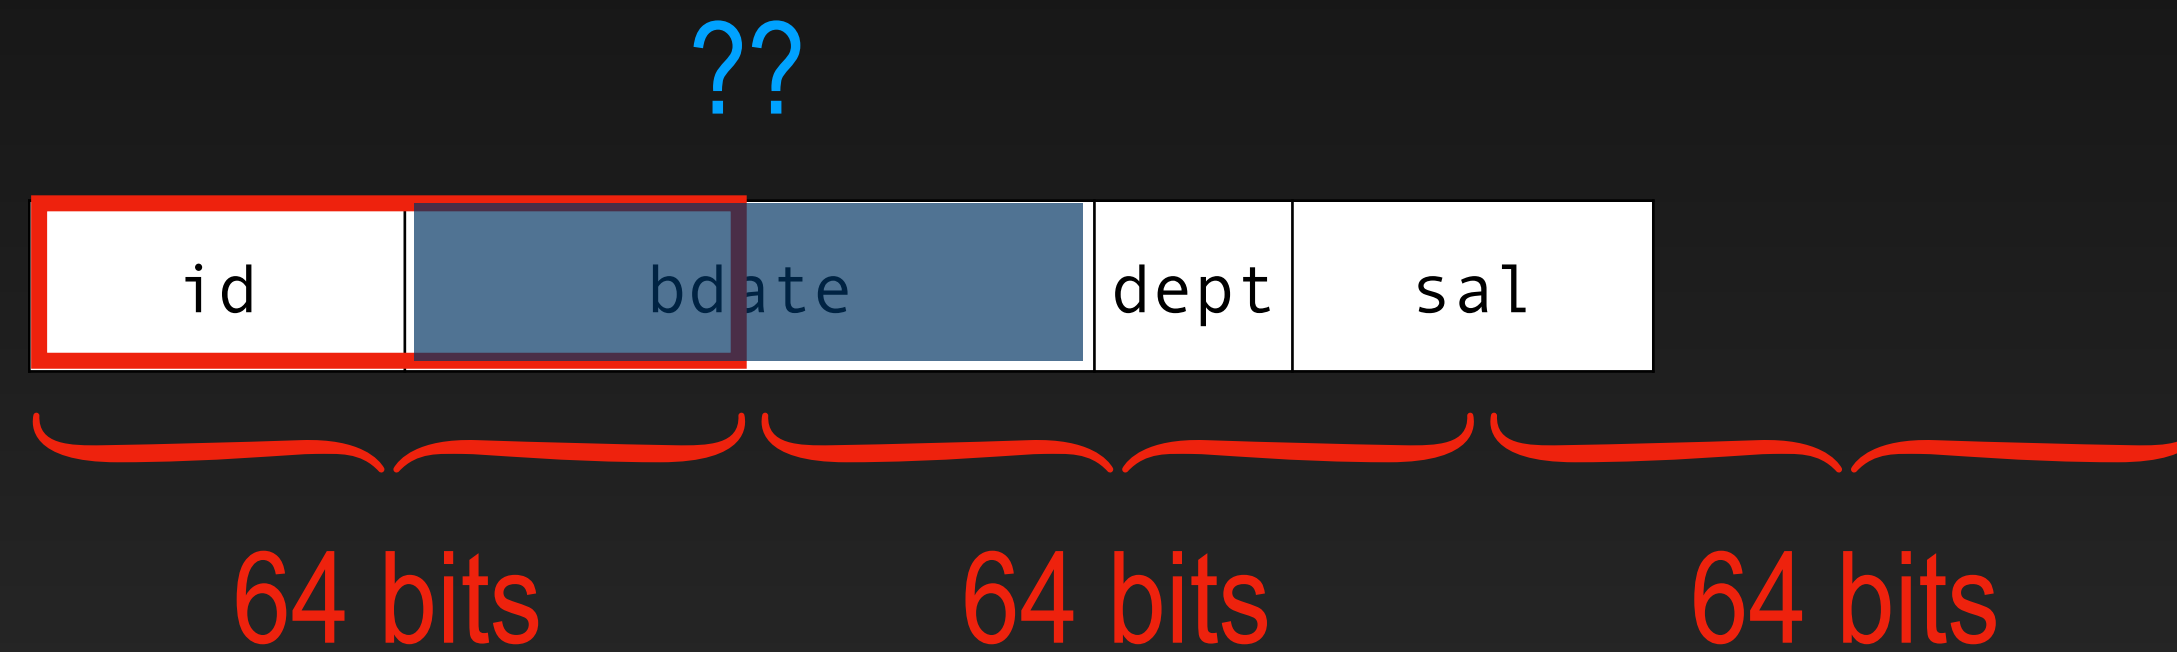
```
id: int  
bdate: timestamp  
dept: char(2)  
sal: int
```

id	bdate	dept	sal
----	-------	------	-----

# Tuple Storage



id: int	32 bits
bdate: timestamp	64 bits
dept: char(2)	16 bits
sal: int	32 bits





# Word-alignment: Padding

<code>id: int</code>	32 bits
<code>bdate: timestamp</code>	64 bits
<code>dept: char(2)</code>	16 bits
<code>sal: int</code>	32 bits

