



# QUERY EXECUTION

Parsing & Analysis

SQL

Analyzer

Initial  
Plan

Execution  
Engine

Result

Metadata

atalog

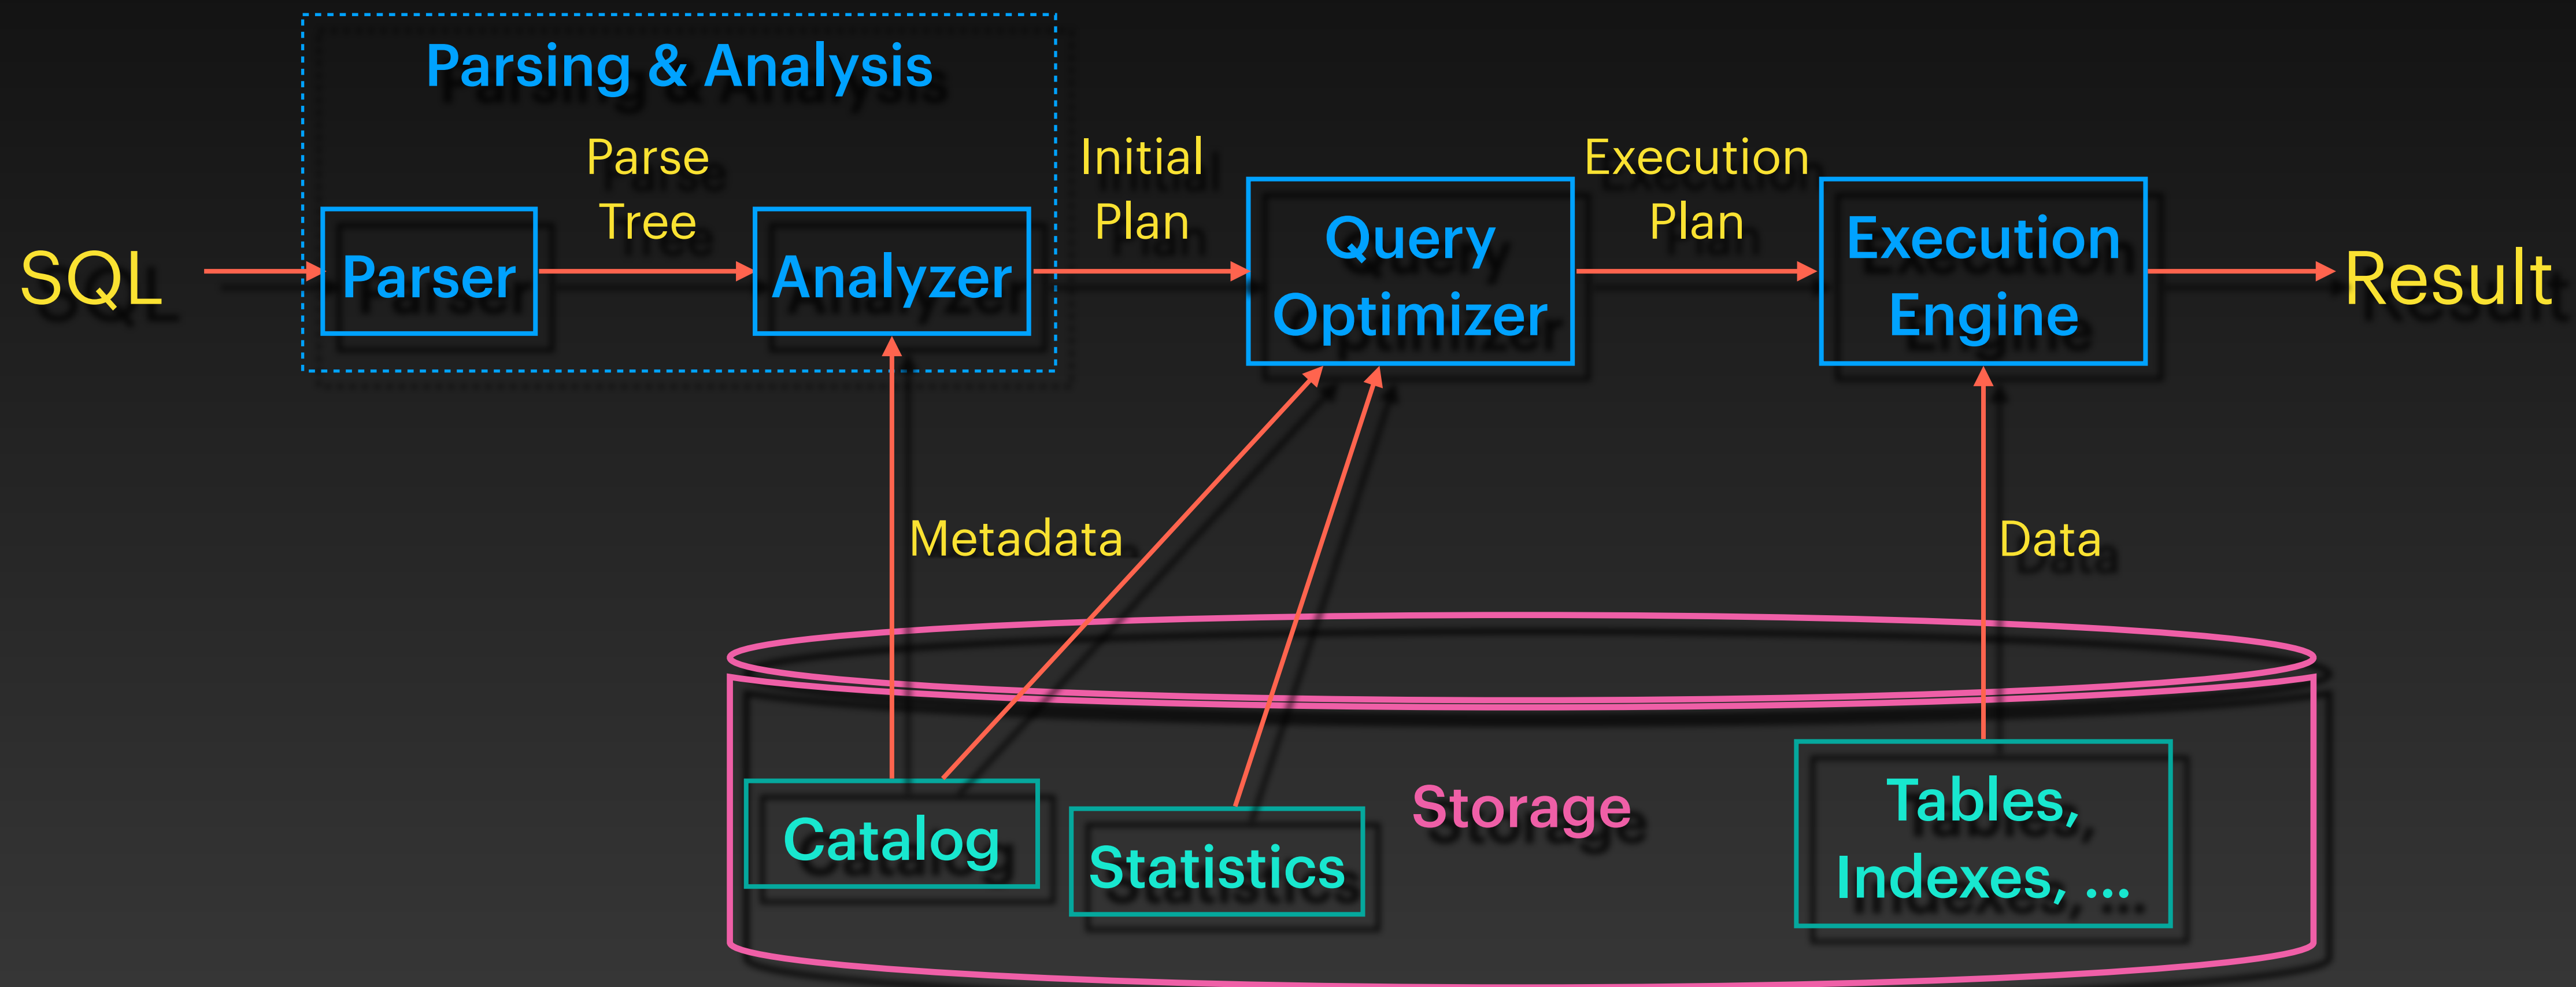
Statistics

Storage

Tables,  
Indexes, ...

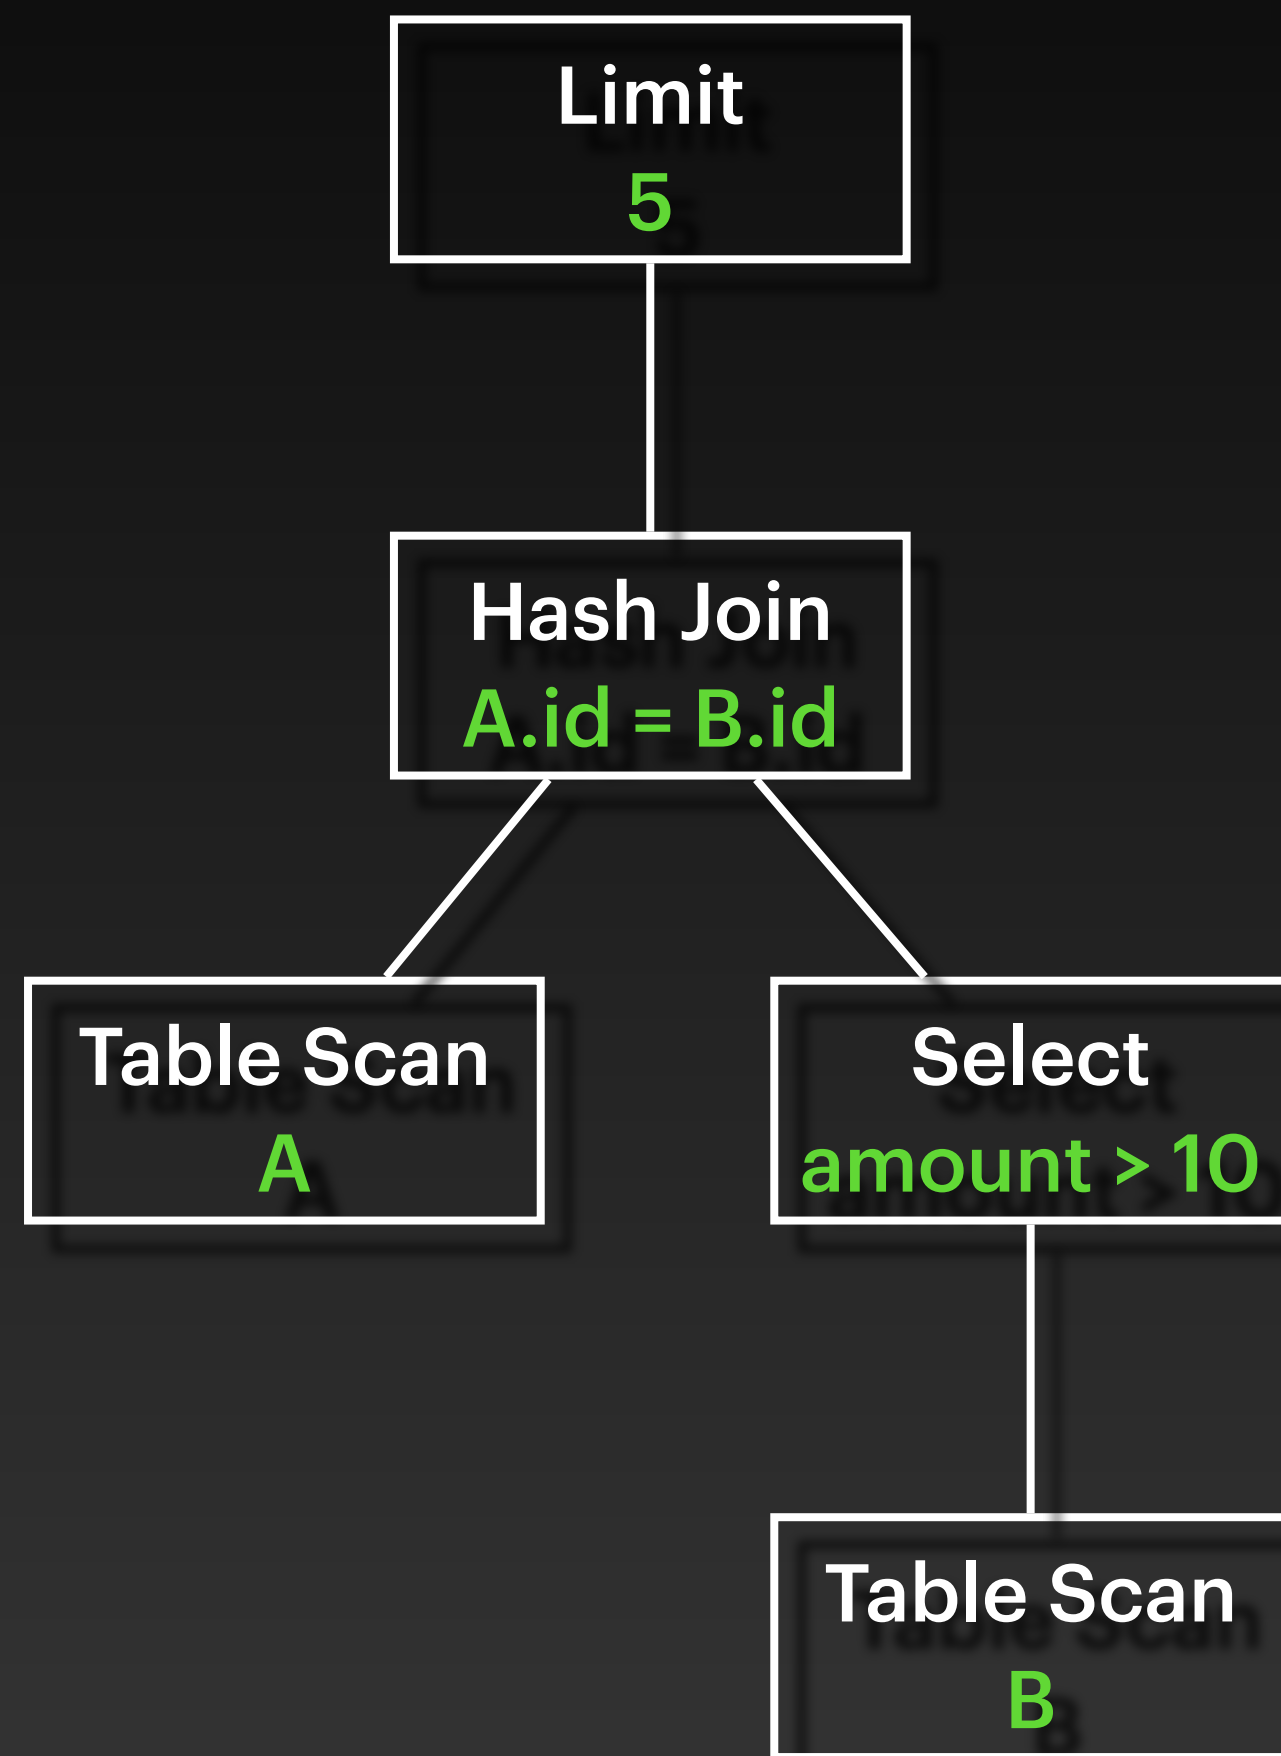
AMR ELHELW

# Query Engine





```
SELECT *  
FROM A JOIN B  
ON A.id = B.id  
WHERE B.amount > 100  
LIMIT 5
```



```
class TableScan : Node {  
    ...  
};  
  
class Select : Node {  
    ...  
};  
  
class HashJoin : Node {  
    ...  
};  
  
class Limit : Node {  
    ...  
};  
  
...
```

# Materialization Model

- Each operation processes all its input(s) at once, and produces its entire output
- Each Node has a function **getResults()** which returns all the output of this operation
- A node can call the functions of its input (child) nodes

```

class TableScan : Node {
    Table table;

    ResultSet getResults() {
        ResultSet out;
        for row in table;
            out.add(row);
        return out;
    }
};

```

```

class Select : Node {
    Node child;
    Condition cond;

    ResultSet getResults() {
        ResultSet out;
        for row in child.getResults():
            if check(cond, row):
                out.add(row);
        return out;
    }
};

```

```

class HashJoin : Node {
    Node left, right;
    JoinKey leftKey, rightKey;

    ResultSet getResults() {
        HashTable ht;
        for leftRow in left.getResults():
            ht.add(leftKey, leftRow);

        ResultSet out;
        for rightRow in right.getResults():
            for leftRow in ht.lookup(rightRow):
                out.add(join(leftRow, rightRow));
        return out;
    }
};

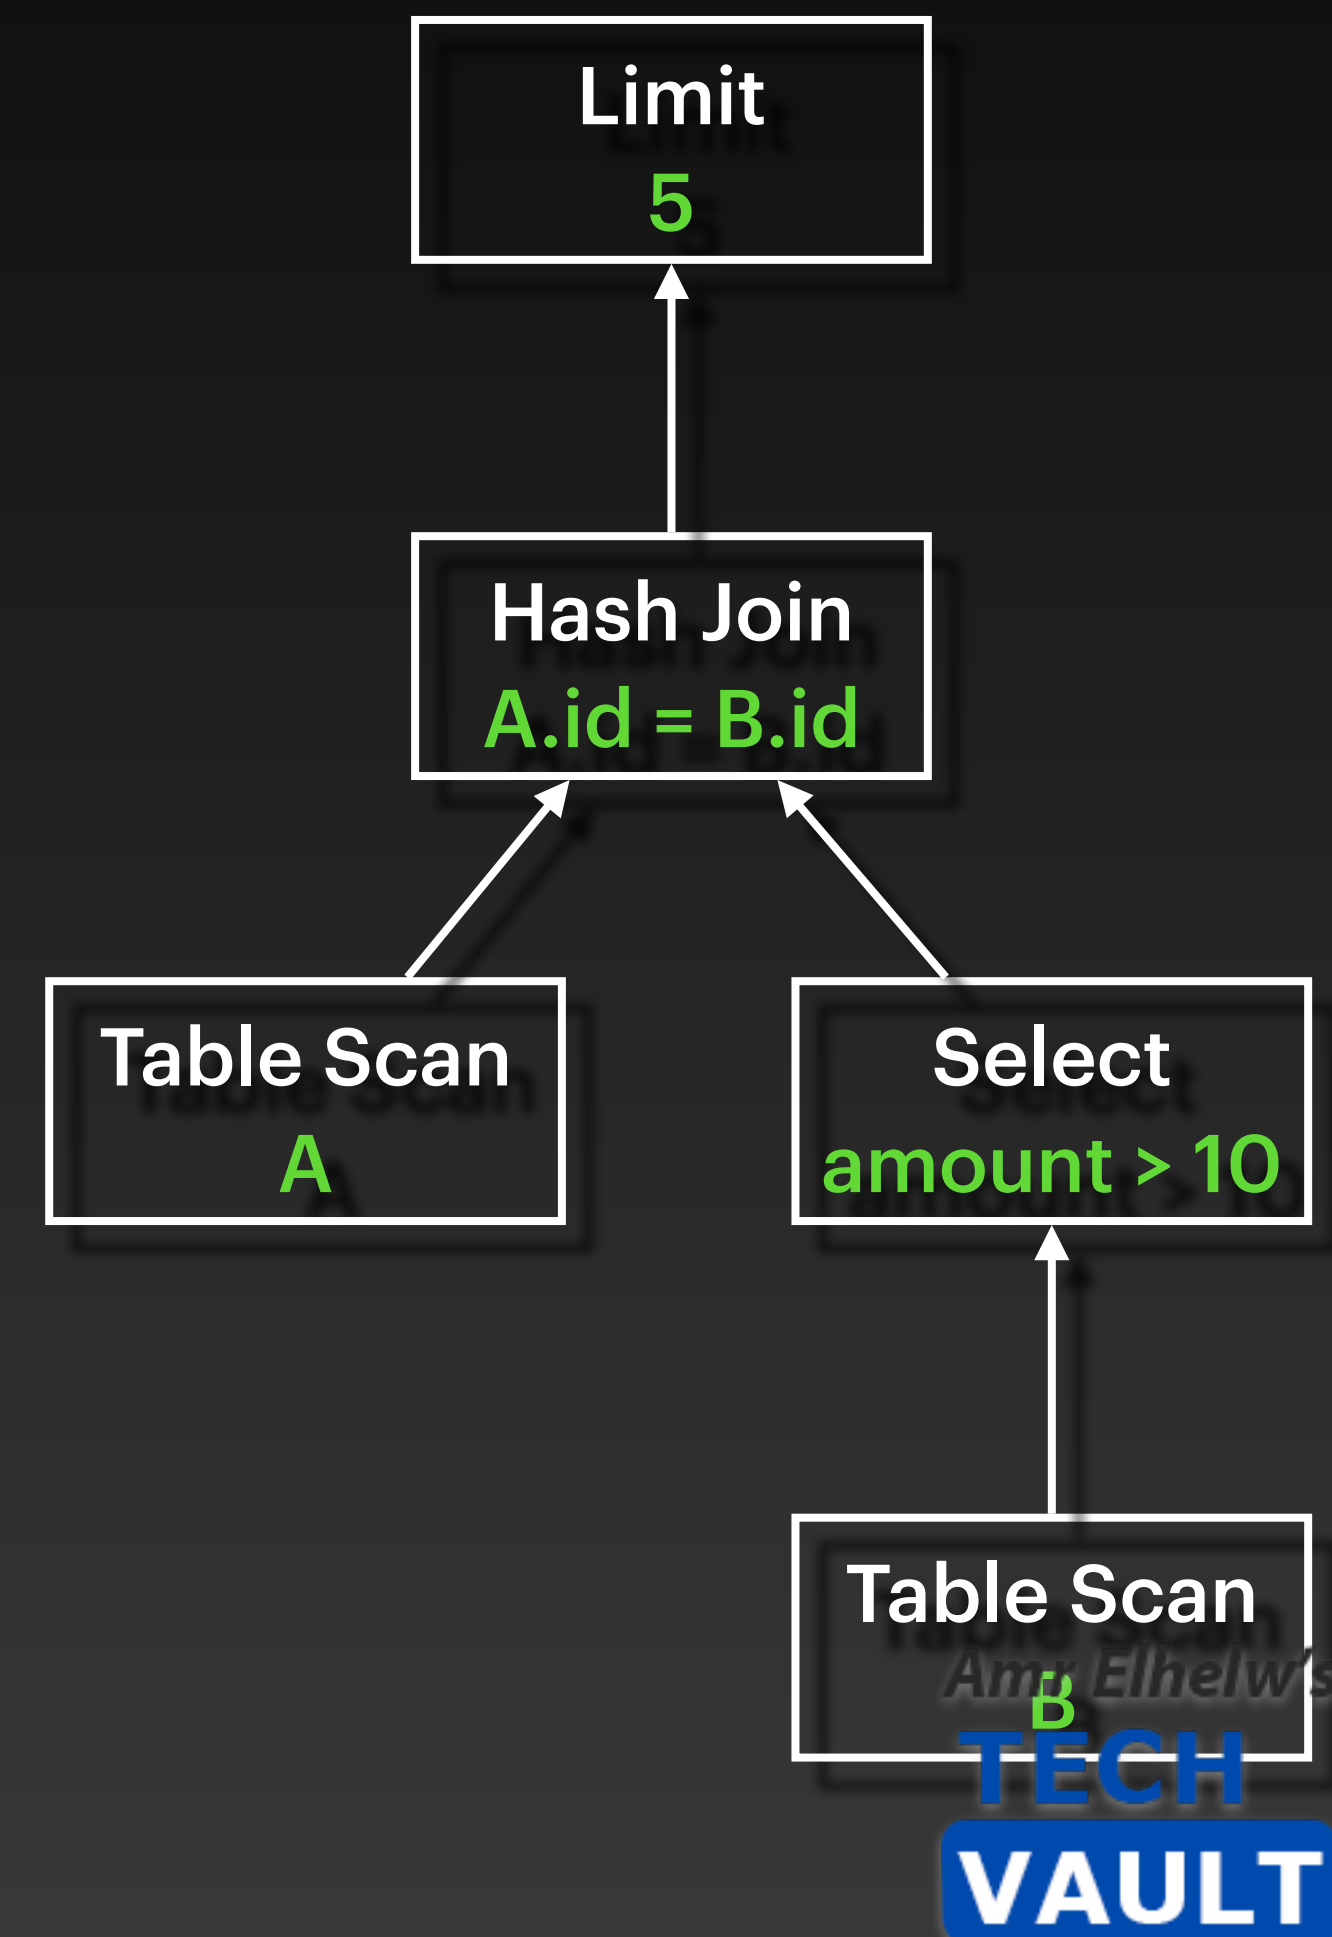
```

```

class Limit : Node {
    Node child;
    int n;

    ResultSet getResults() {
        ResultSet out;
        out.addFirst(child.getResults(), n);
        return out;
    }
};

```



```
root.getResults();
```

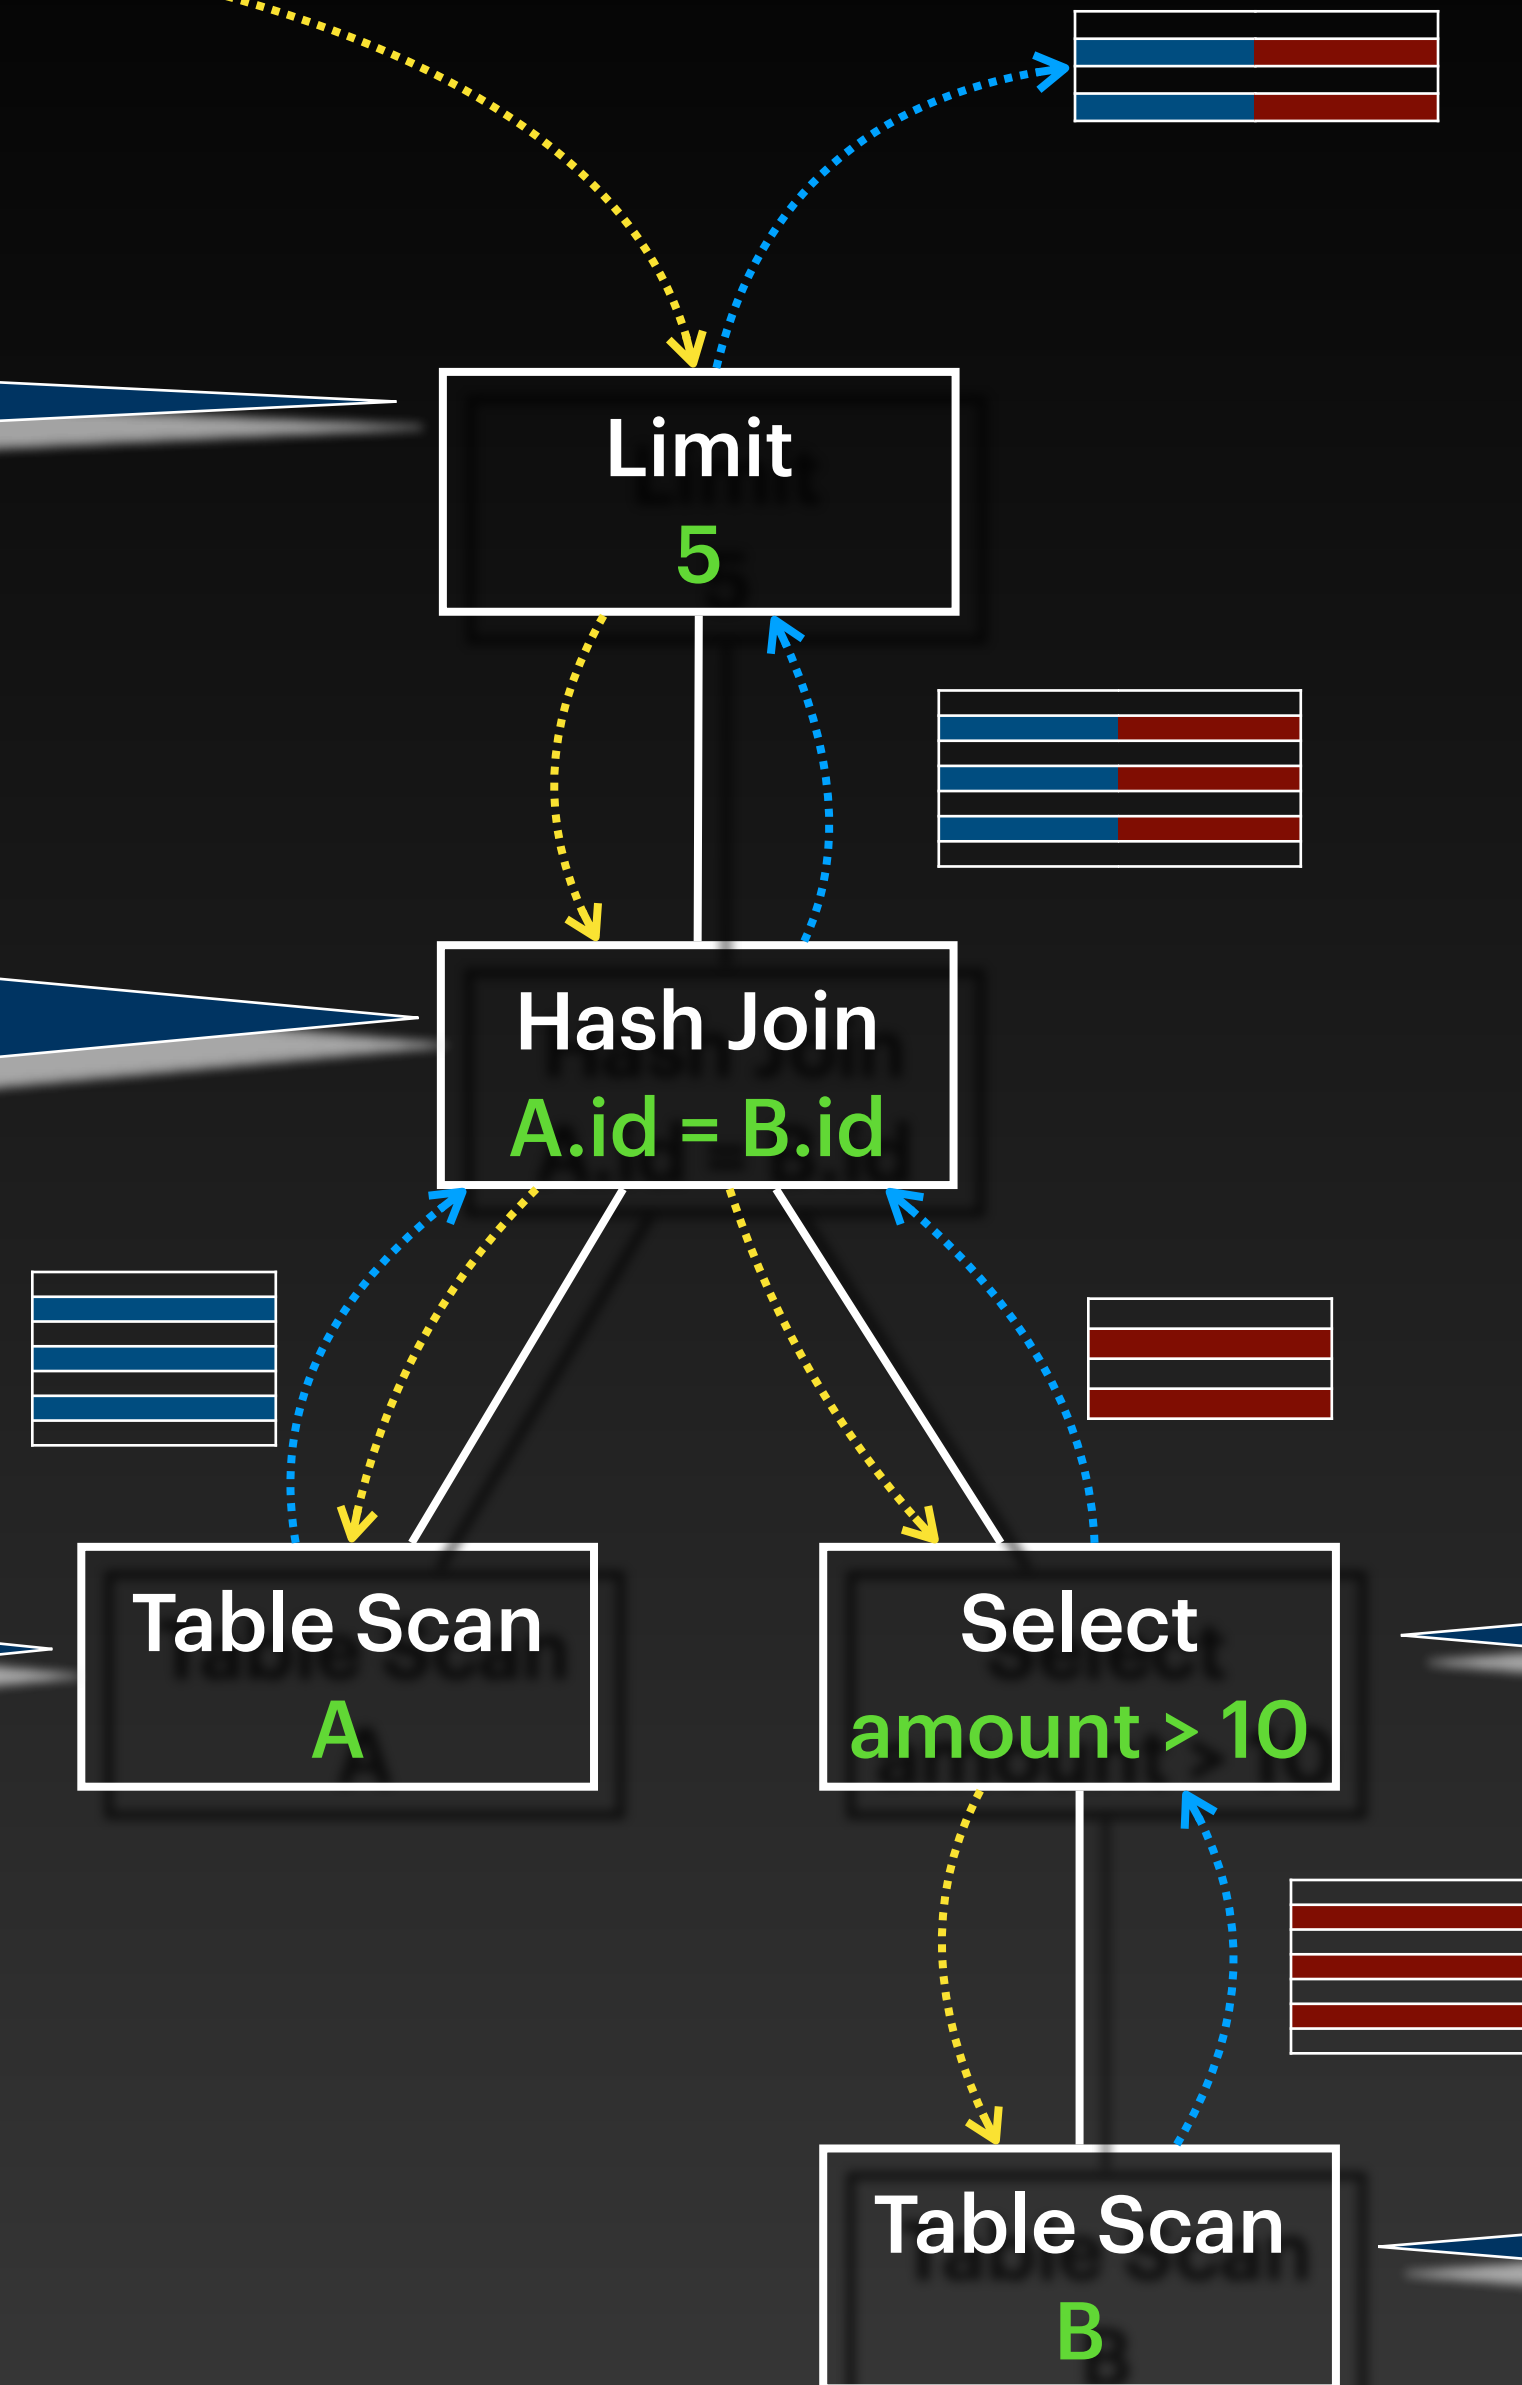
```
Limit::getResults() {  
  ResultSet out;  
  out.addFirst(child.getResults(), n);  
  return out;  
}
```

```
HashJoin::getResults() {  
  HashTable ht;  
  for leftRow in left.getResults():  
    ht.add(leftKey, leftRow);  
  
  ResultSet out;  
  for rightRow in right.getResults():  
    for leftRow in ht.lookup(rightRow):  
      out.add(join(leftRow, rightRow));  
  return out;  
}
```

```
TableScan::getResults() {  
  ResultSet out;  
  for row in table;  
    out.add(row);  
  return out;  
}
```

```
Select::getResults() {  
  ResultSet out;  
  for row in child.getResults():  
    if check(cond, row):  
      out.add(row);  
  return out;  
}
```

```
TableScan::getResults() {  
  ResultSet out;  
  for row in table;  
    out.add(row);  
  return out;  
}
```



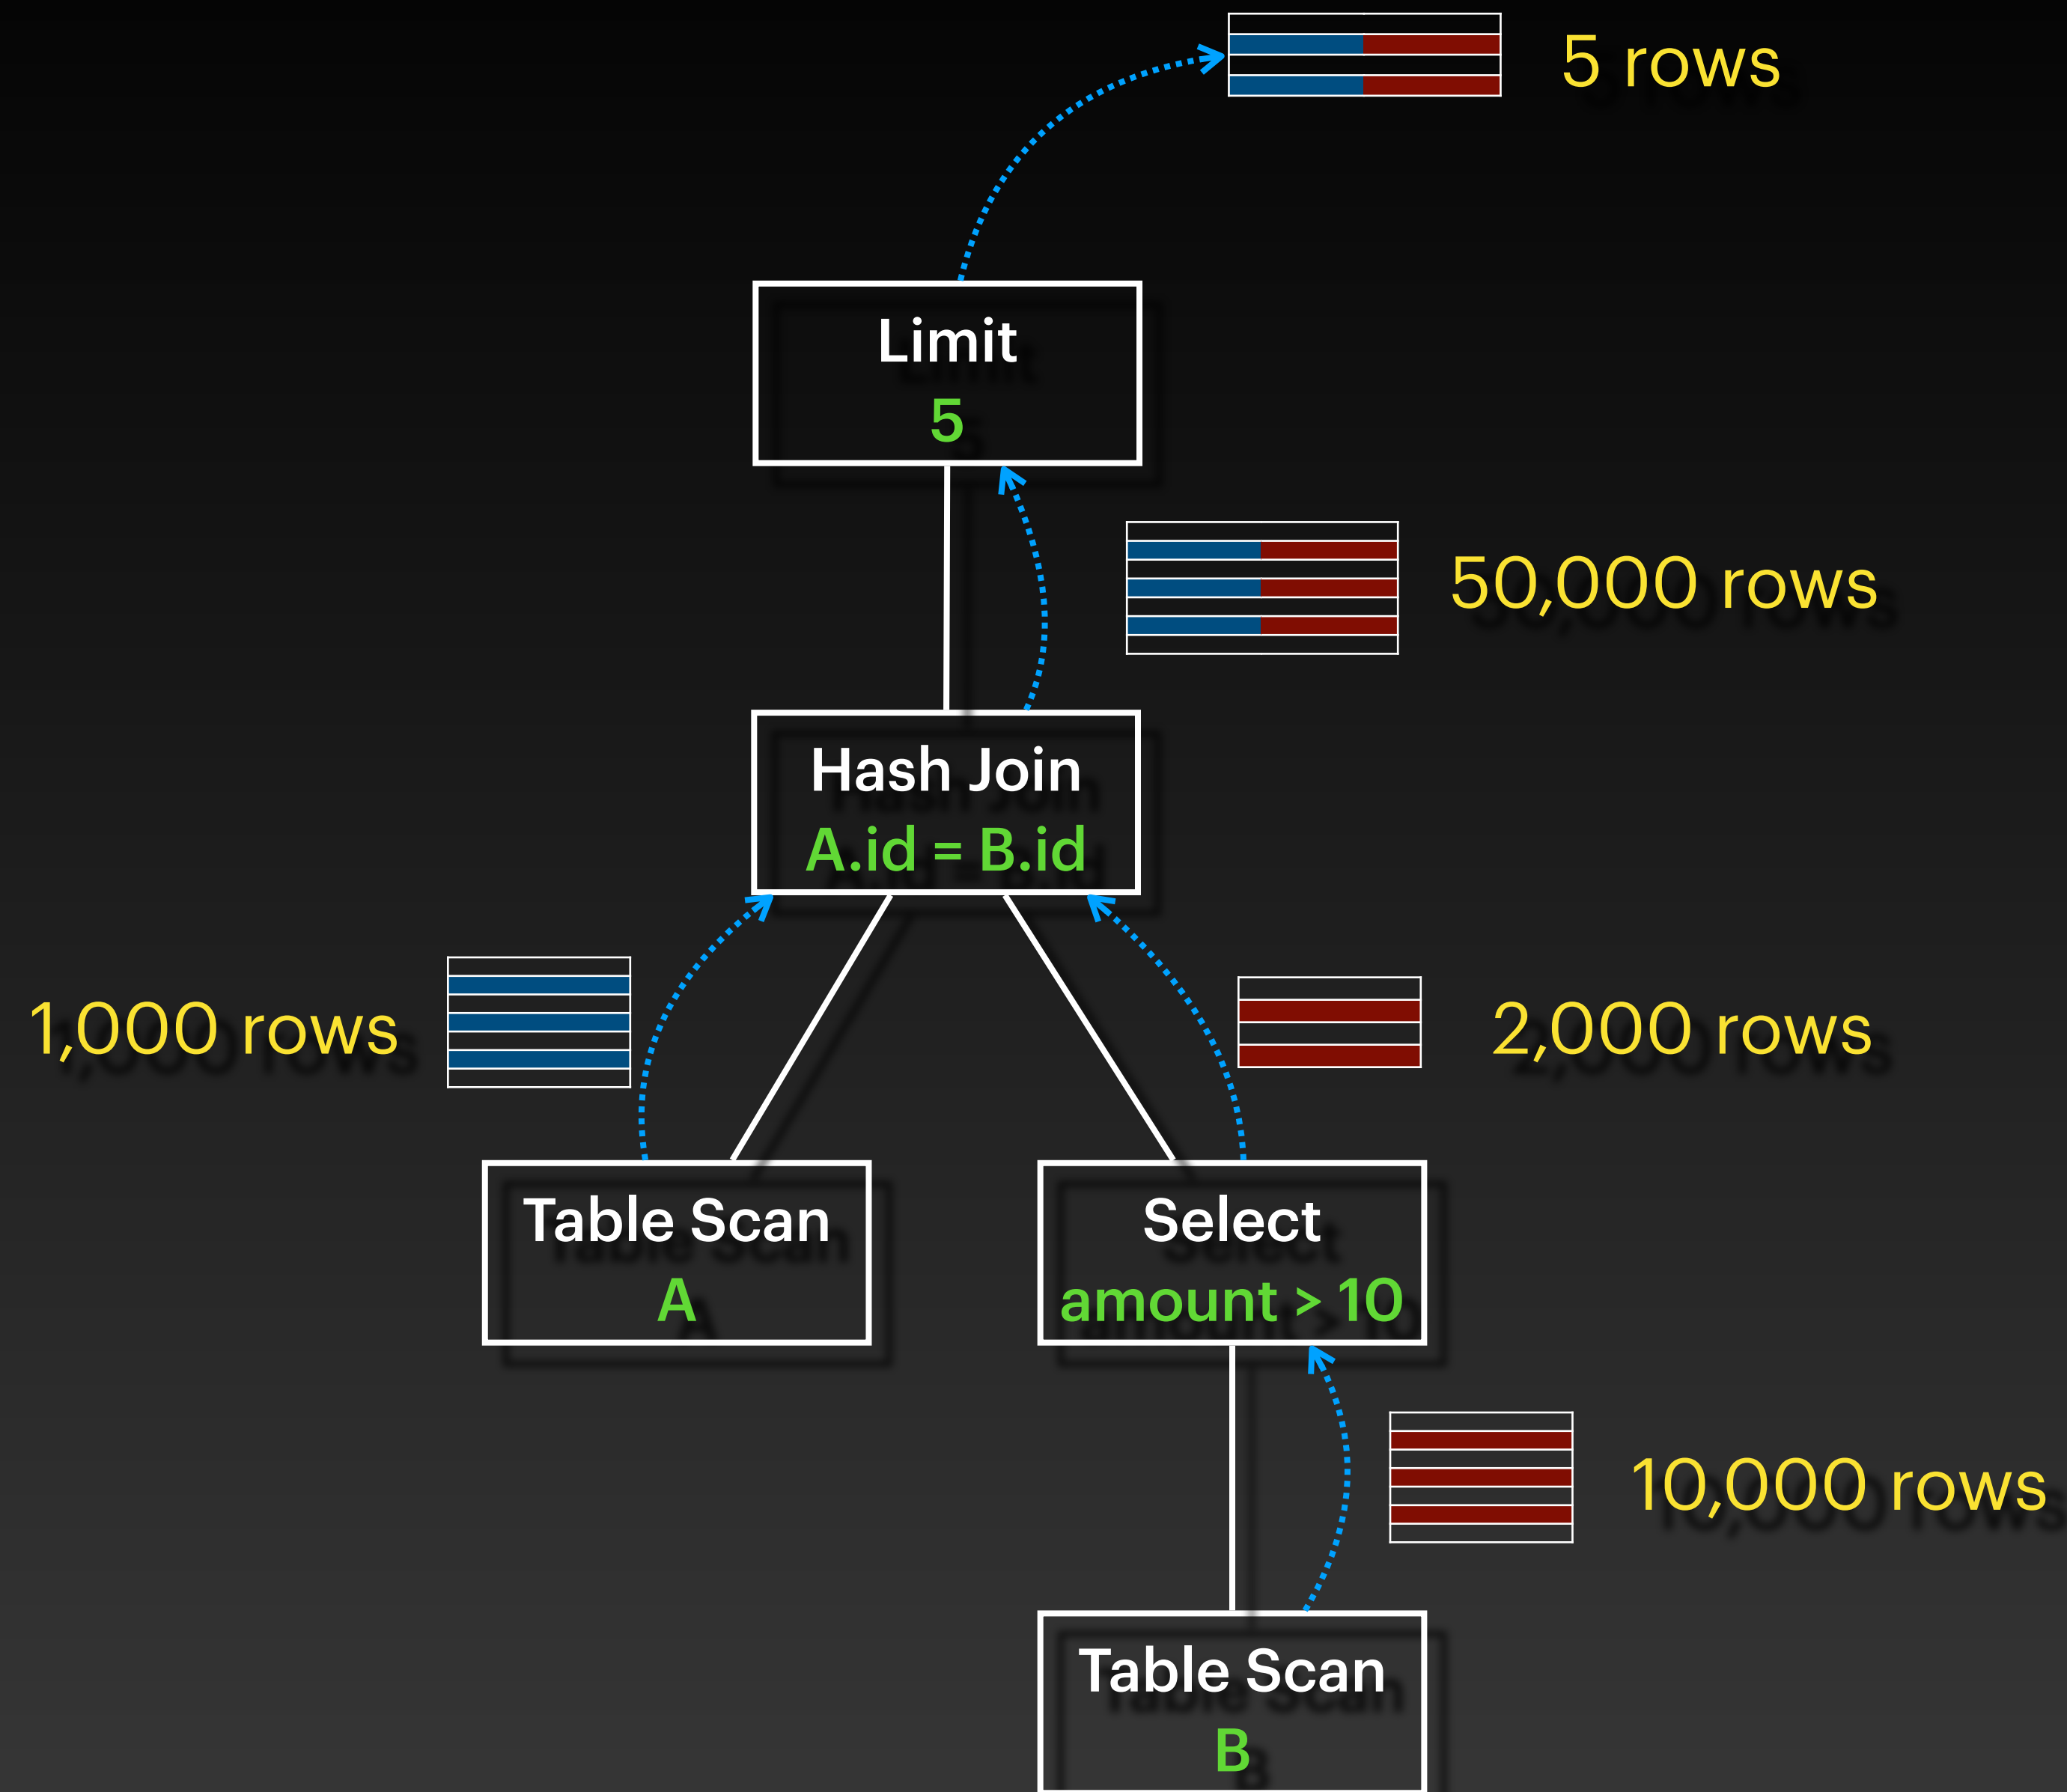


# Advantages

- Simple to implement
- No coordination required between operations
- Each operation is called only once
- Good for OLTP queries
  - Few operations
  - Small intermediate results

# Disadvantages

- Not good for OLAP/analytical queries
  - Complex queries, many operations
  - Large intermediate results
- May do some wasted work
  - Example: **LIMIT**



# Iterator Model



- Each operator implements **getNext()** function
  - Each call returns **next row** in result, or **EOF**
- Operators are “active” for longer
  - Need to maintain “state”
- Other important functions:
  - **open()** - initialization, any work needed before it can start returning results
  - **close()** - cleanup, release resources, etc.

```

class TableScan : Node {
    open() {
        initializeCursor();
    }

    Row getNext() {
        row = readRowAtCursor();
        if (row == EndOfTable):
            return EOF;
        advanceCursor();
        return row;
    };
};

```

```

class Limit : Node {
    Node child;
    int n;

    Row getNext() {
        if n == 0:
            return EOF;
        row = child.getNext();
        n--;
        return row;
    }
};

```

```

class Select : Node {
    Node child;
    Condition cond;

    Row getNext() {
        while (row = child.getNext()) != EOF {
            if check(cond, row):
                return row;
        }
        return EOF;
    }
};

```

```

class HashJoin : Node {
    Node left, right;
    JoinKey leftKey, rightKey;
    HashTable ht;
    Bool hashTableReady = false;

    Row getNext() {
        if (!hashTableReady) {
            while (leftRow = left.getNext()) != EOF:
                ht.add(leftKey, leftRow);
            hashTableReady = true;
        }
        while (rightRow = right.getNext()) != EOF {
            leftRow = ht.lookup(rightRow):
            return join(leftRow, rightRow);
        }
        return EOF;
    }
};

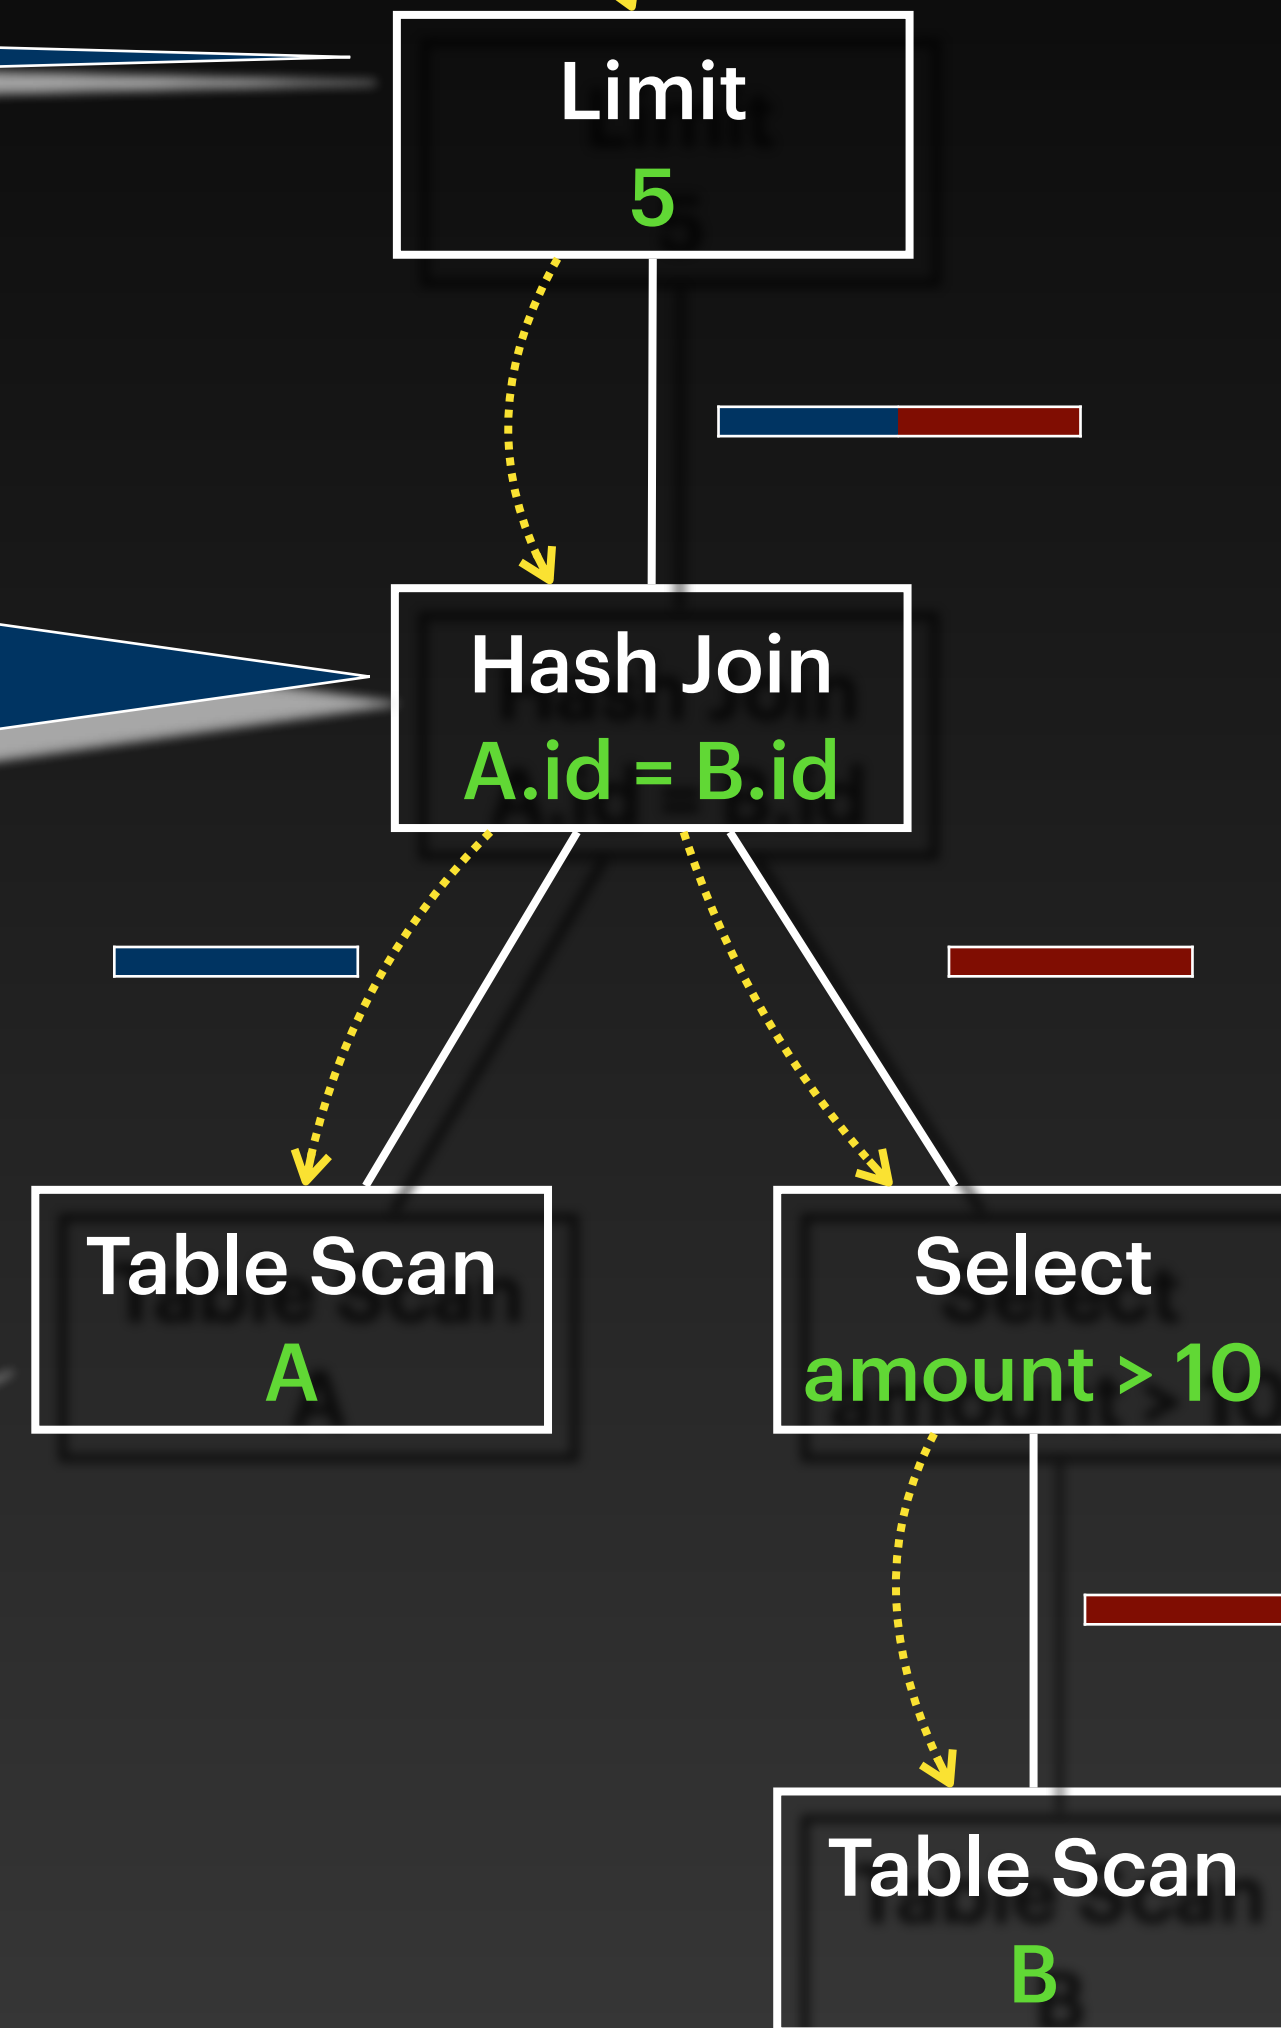
```

`root.getNext();`

```
if n == 0:
    return EOF;
row = child.getNext();
n--;
return row;
```

```
if (!hashTableReady) {
    while (leftRow = left.getNext()) != EOF:
        ht.add(leftKey, leftRow);
        hashTableReady = true;
}
while (rightRow = right.getNext()) != EOF {
    leftRow = ht.lookup(rightRow);
    return join(leftRow, rightRow);
}
return EOF;
```

```
row = readRowAtCursor();
if (row == EndOfTable):
    return EOF;
advanceCursor();
return row;
```



```
while (row = child.getNext()) != EOF {
    if check(cond, row):
        return row;
}
return EOF;
```

```
row = readRowAtCursor();
if (row == EndOfTable):
    return EOF;
advanceCursor();
return row;
```

# Advantages

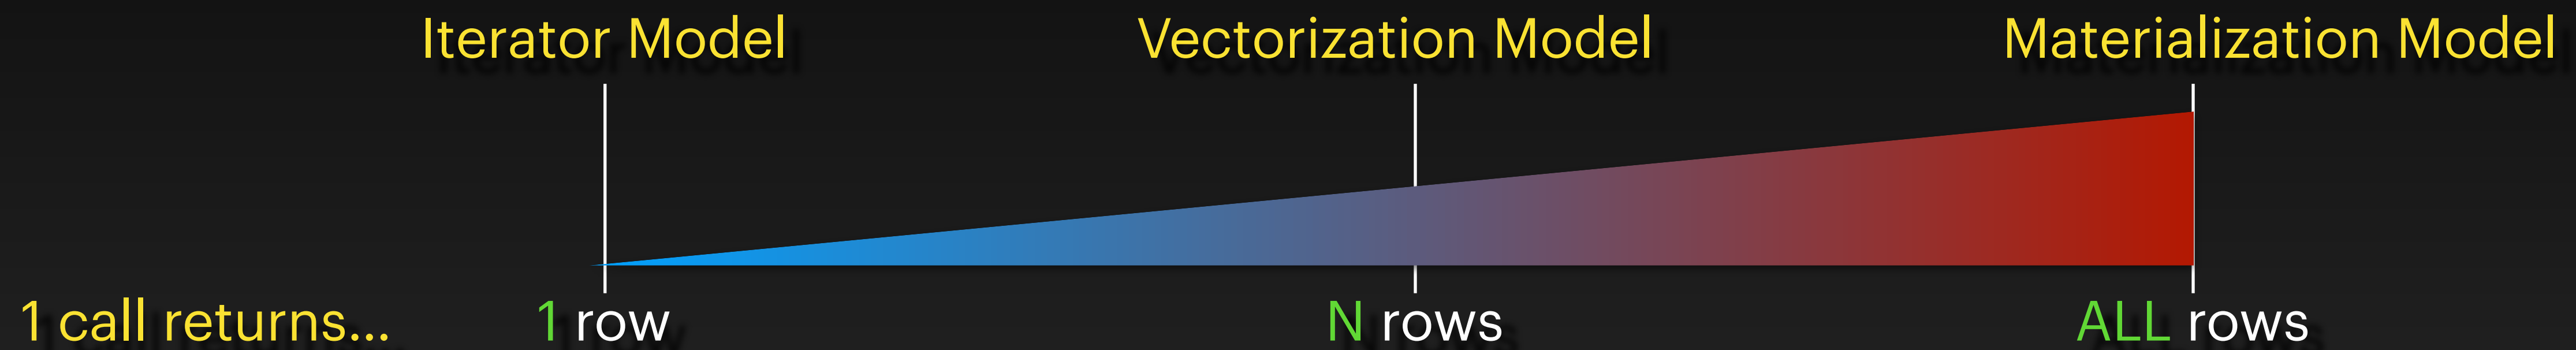
- Allows for **pipelining**
  - Several operators can be “active” at the same time
  - Start producing outputs before reading all the inputs
- Avoid unnecessary data reading (e.g. in the case of **LIMIT**)
- Does not require keeping all intermediate results in memory
- Most (or all) DBMSs support this model



# Disadvantages

- Slightly more complex way to implement
- Lots of context switching between operators, to return one row at a time
- Lots of function calls - can be expensive
- Some operators can be **blocking**
  - Need to consume all the input before they can start producing output
  - Examples: hash join, aggregation, sort

# Vectorization Model



- Each operator implements **getNextBatch()** function
  - Each call returns **next batch of rows** in result, or **EOF**
- Batch size can depend on available memory

```

class TableScan : Node {
    Table table;

    ResultSet getResults() {
        ResultSet out;
        for row in table;
            out.add(row);
        return out;
    }
};

```

## Materialization Model

```

class TableScan : Node {
    open() {
        initializeCursor();
    }

    Row getNext() {
        row = readRowAtCursor();
        if (row == EndOfTable):
            return EOF;
        advanceCursor();
        return row;
    }
};

```

## Iterator Model

```

class TableScan : Node {
    open() {
        initializeCursor();
    }

    ResultSet getNextBatch() {
        ResultSet out;
        for (i = 0; i < batchSize; ++i) {
            row = readRowAtCursor;
            if (row == EndOfTable):
                break;
            advanceCursor();
            out.add(row);
        }
        if (out.empty()):
            return EOF;
        else
            return out;
    }
};

```

## Vectorization Model



# Advantages

- Doesn't need to produce ALL outputs of each operation before executing the next one
  - Will start producing results before reading entire inputs
- Each operation produces a batch that can “usually” fit in memory
- Can work with large datasets (unlike materialization model)
- Can work on complex queries with many operations operators don't have to wait for too long
- Fewer function calls than iterator model

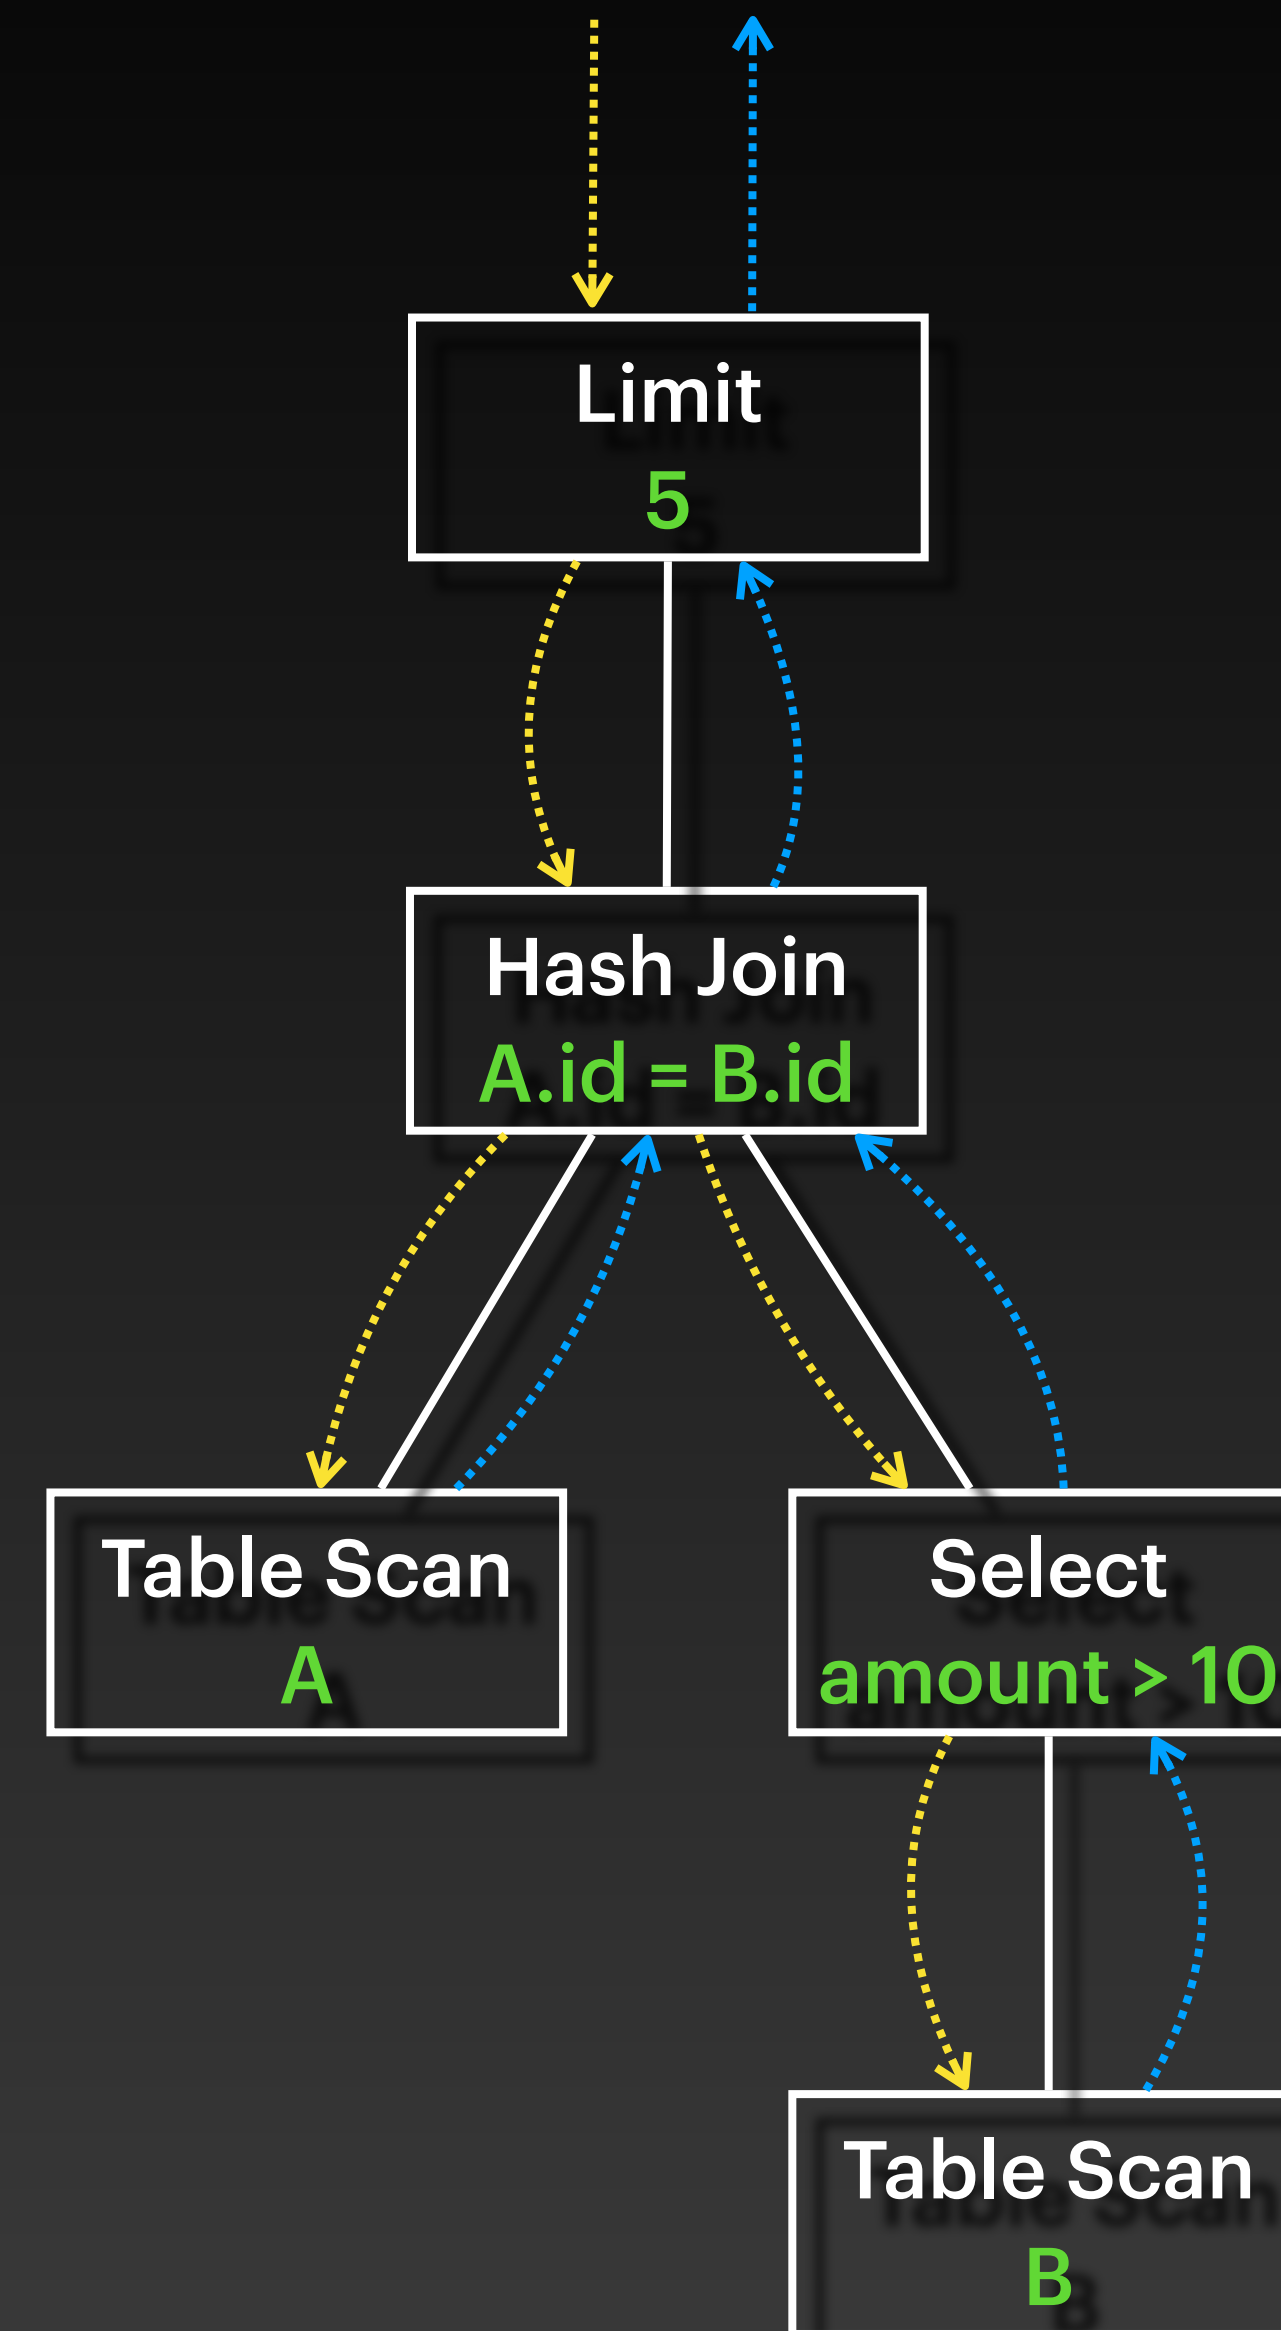
# Disadvantages

- More complex implementation
  - We have to think about batches in every operation
- Longer time (than iterator model) until the first output is produced
- “Some” wasted work potentially, but not as much as with materialization

# Pull-based vs Push-based

# Pull-based Processing

- Top-down
- Each node “pulls” data from its children
- While a node is executing, everything else is waiting (single thread)
- (+) simple to implement & debug
- (-) no parallelism





# Push-based Processing

- Bottom-up
- Buffers are added between nodes
- Each node “pushes” data to its parent
- (+) Multiple nodes can work in parallel
- (-) More complex

